

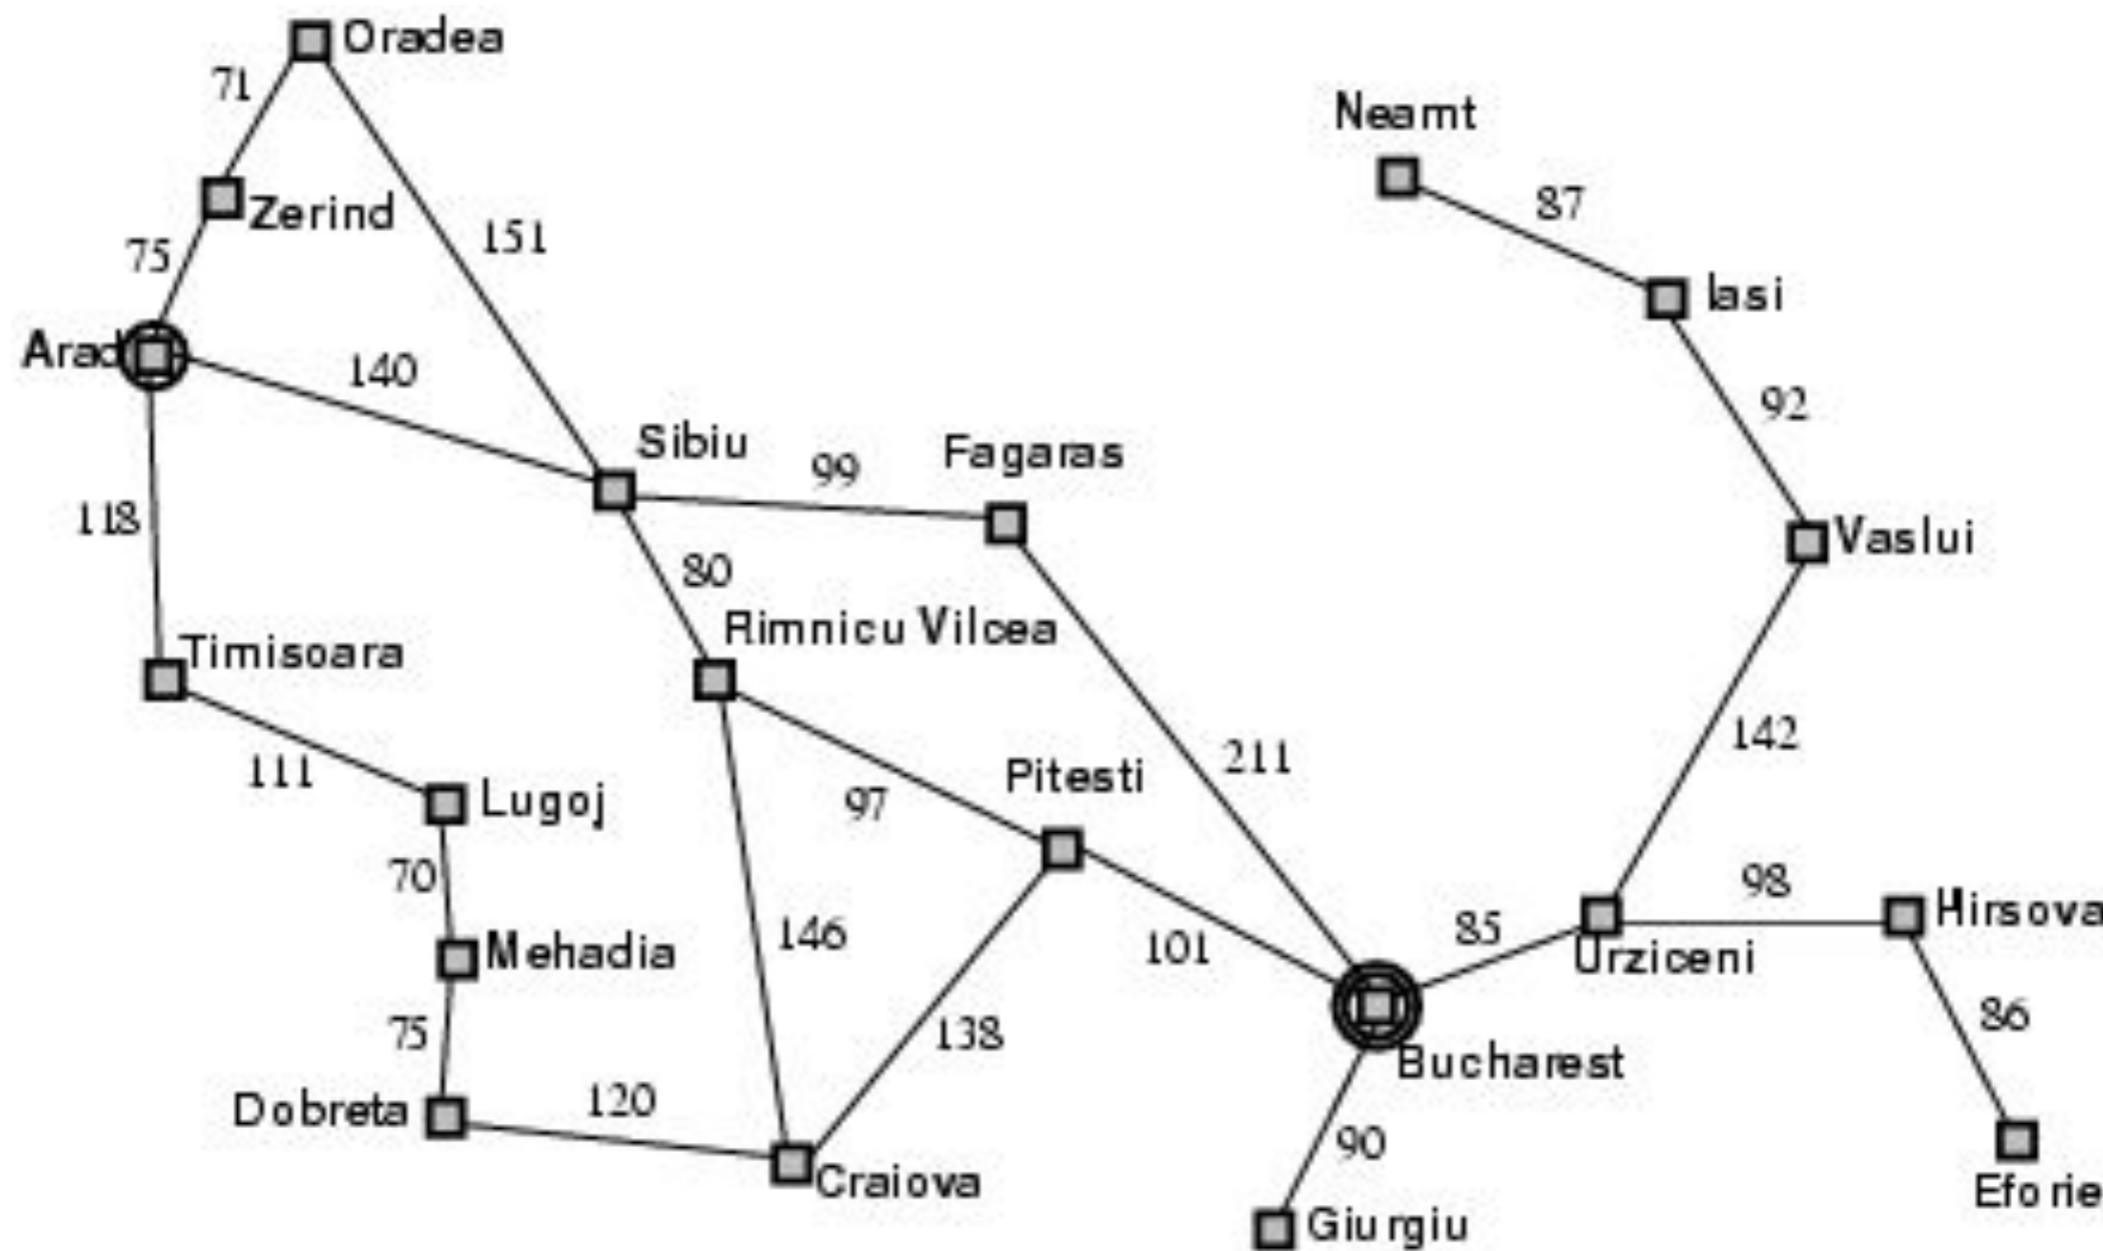
# Lecture 3: Informed search and optimization

Artificial Intelligence  
CS-GY-6613  
Julian Togelius  
[julian.togelius@nyu.edu](mailto:julian.togelius@nyu.edu)

# The chef recommends:

- Best-first search
- Greedy best-first search
- A\*
- Heuristics
- Optimization versus tree search
- Hill-climbing
- Simulated annealing
- Evolutionary algorithms
- Assignment 1

# Remember Romania



# Tree search

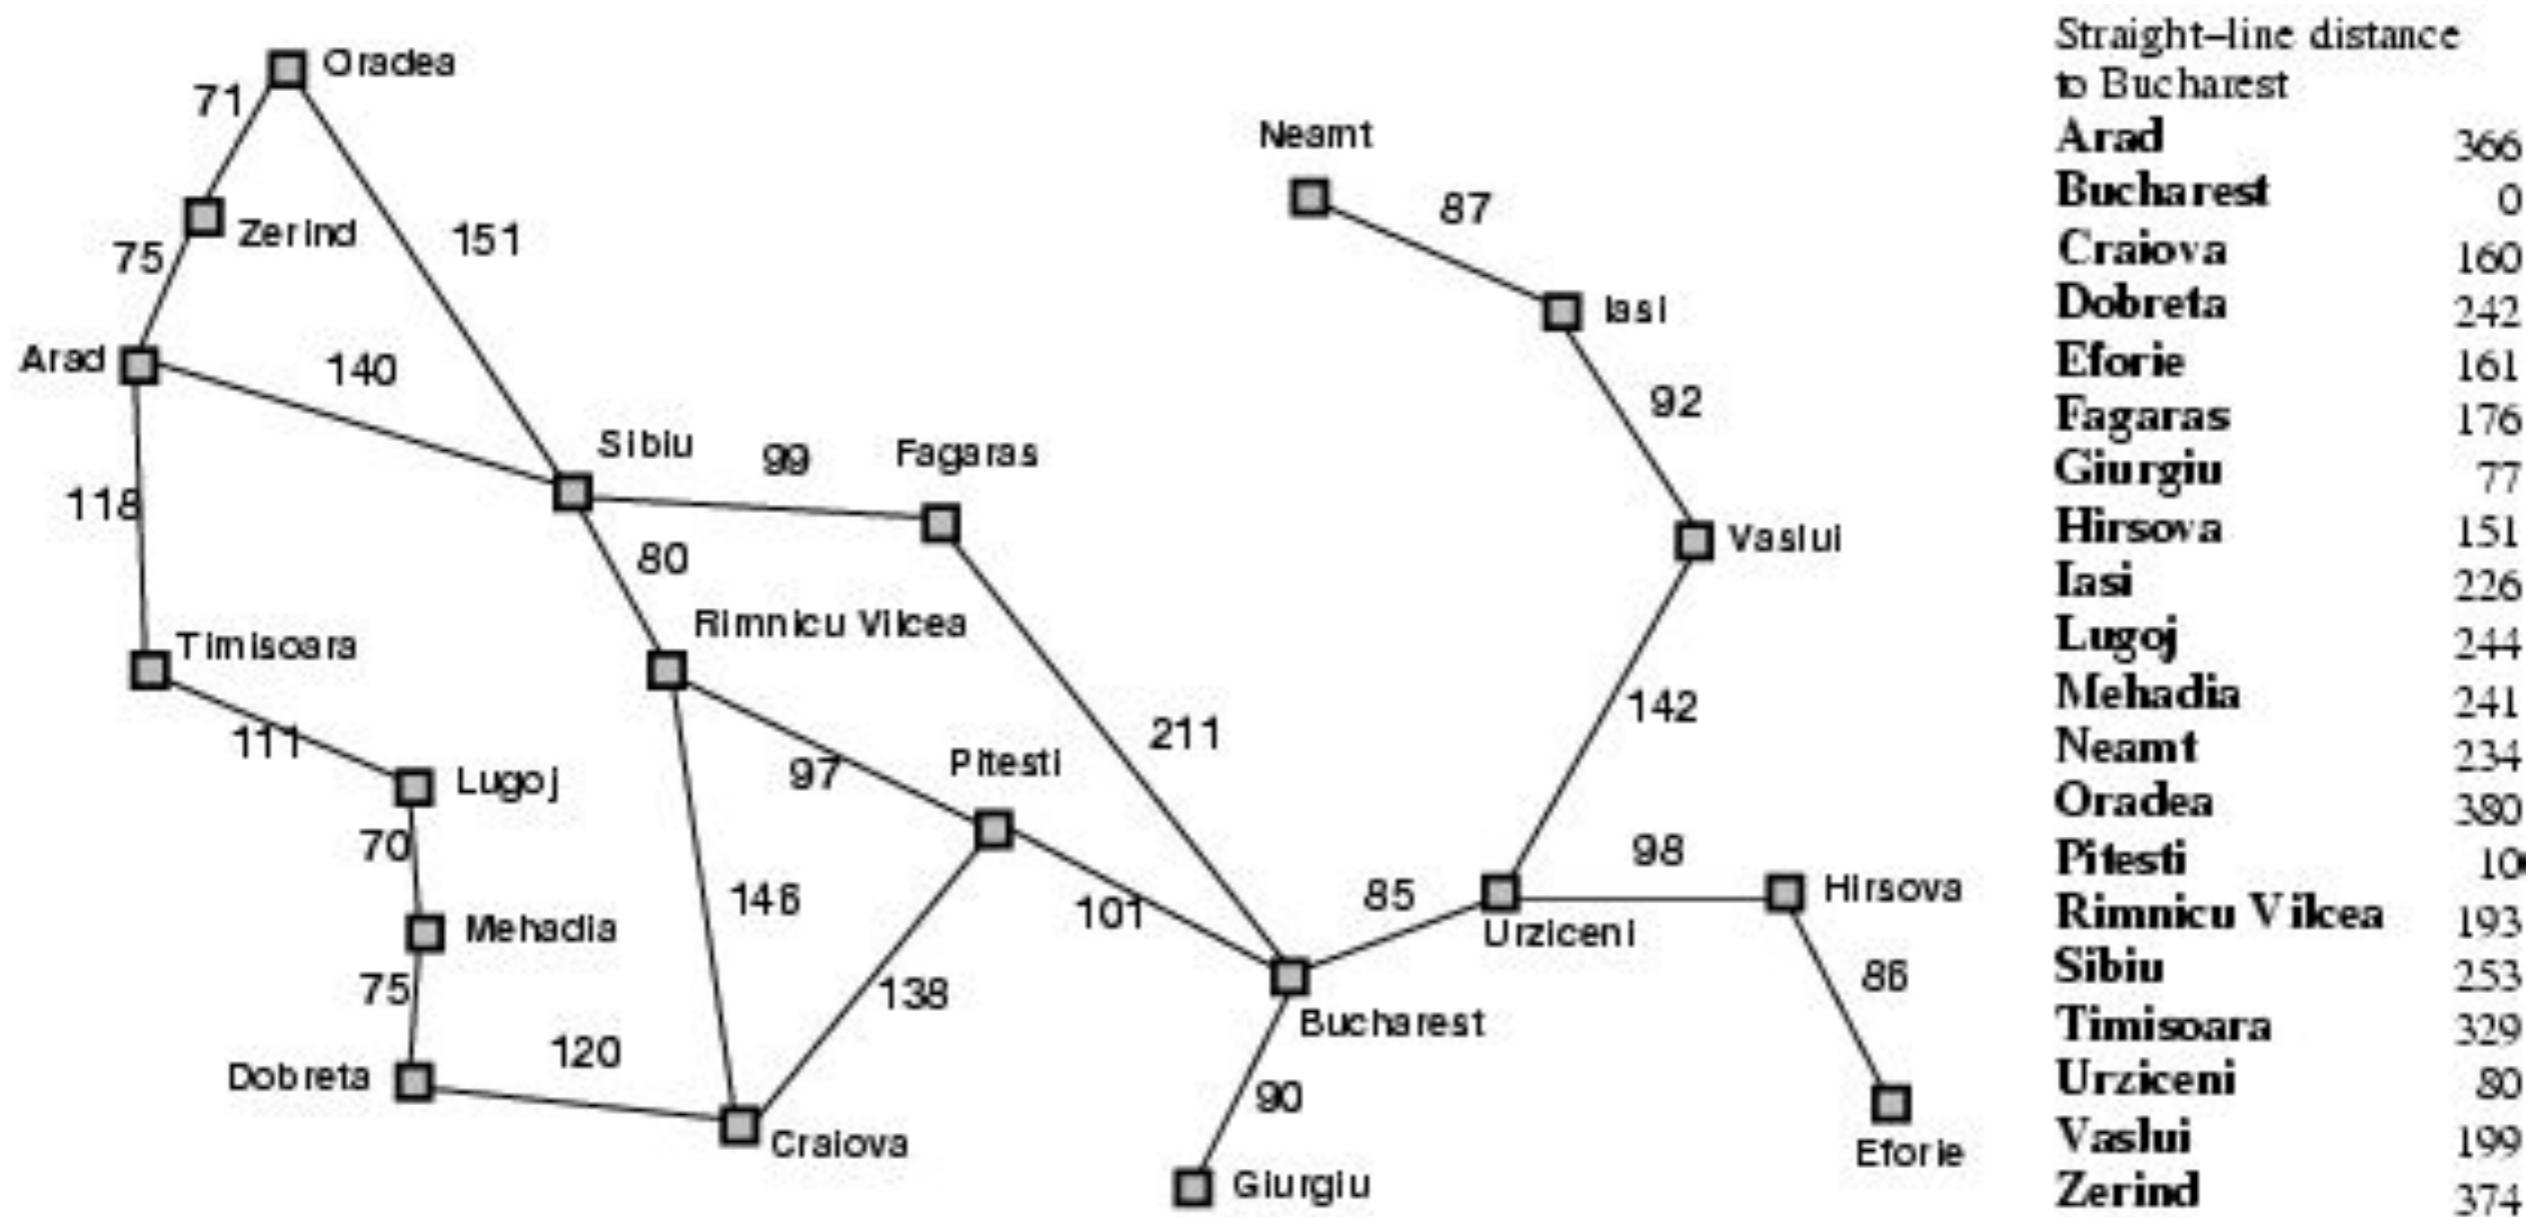
- offline, simulated exploration of state space by generating successors of already-explored states (a.k.a.~expanding states)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
```

# Uninformed search

- Uninformed search strategies use only the information available in the problem definition
- Breadth-first search
- Uniform cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search

# Now with straight line distances!



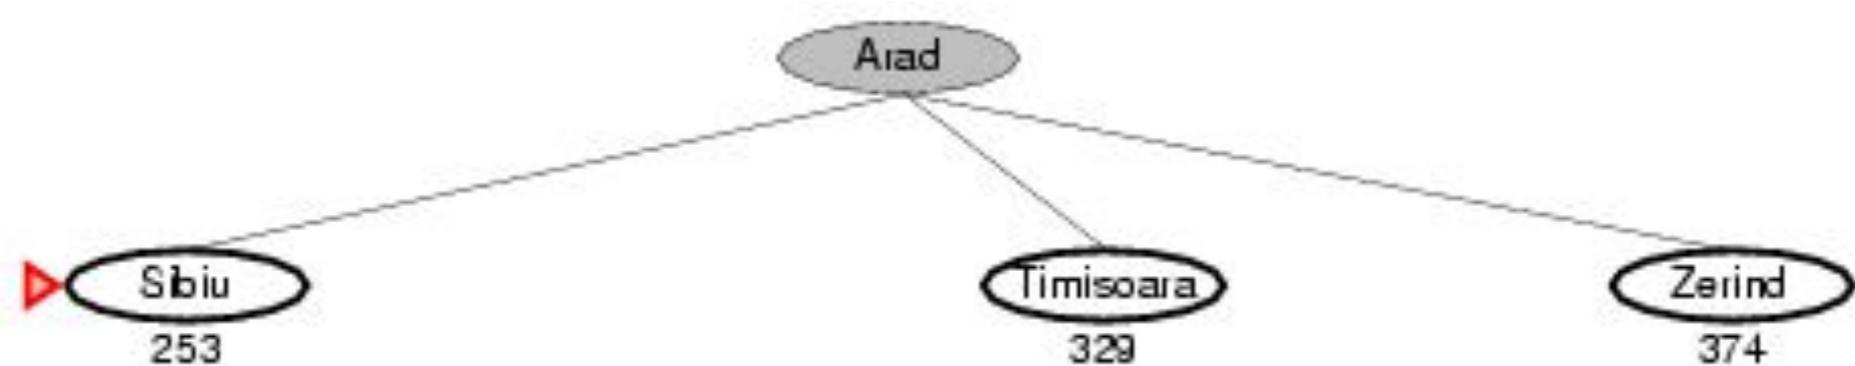
# Greedy best-first search

- Evaluation function  $f(n) = h(n)$  (heuristic)  
= estimate of cost from  $n$  to goal
- e.g.,  $h_{SLD}(n)$  = straight-line distance from  $n$  to Bucharest
- Greedy best-first search expands the node that *appears* to be closest to goal

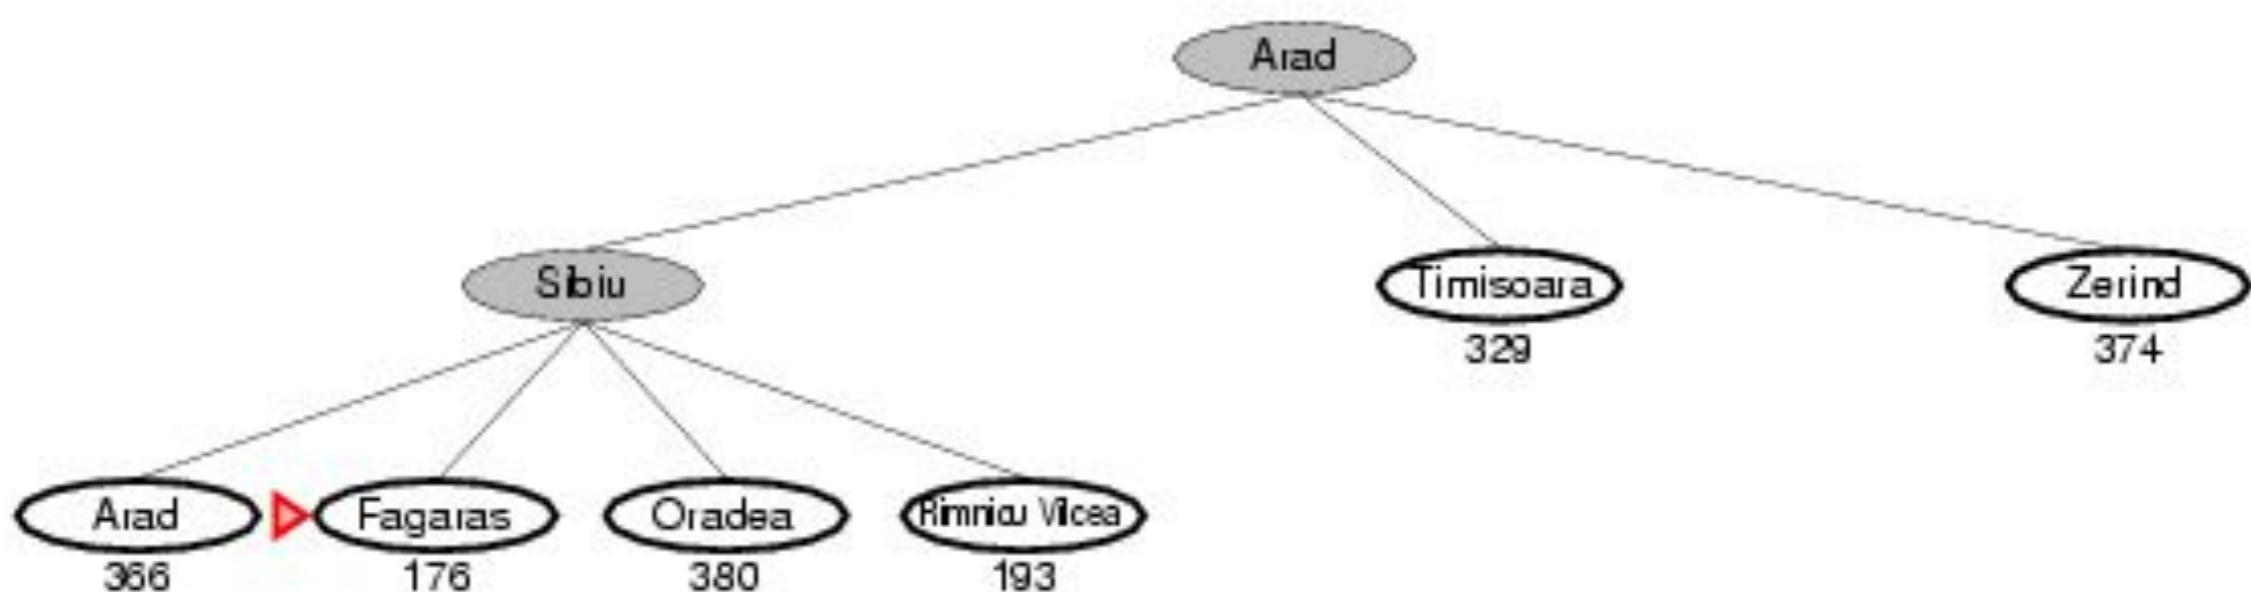
# Greedy best-first example



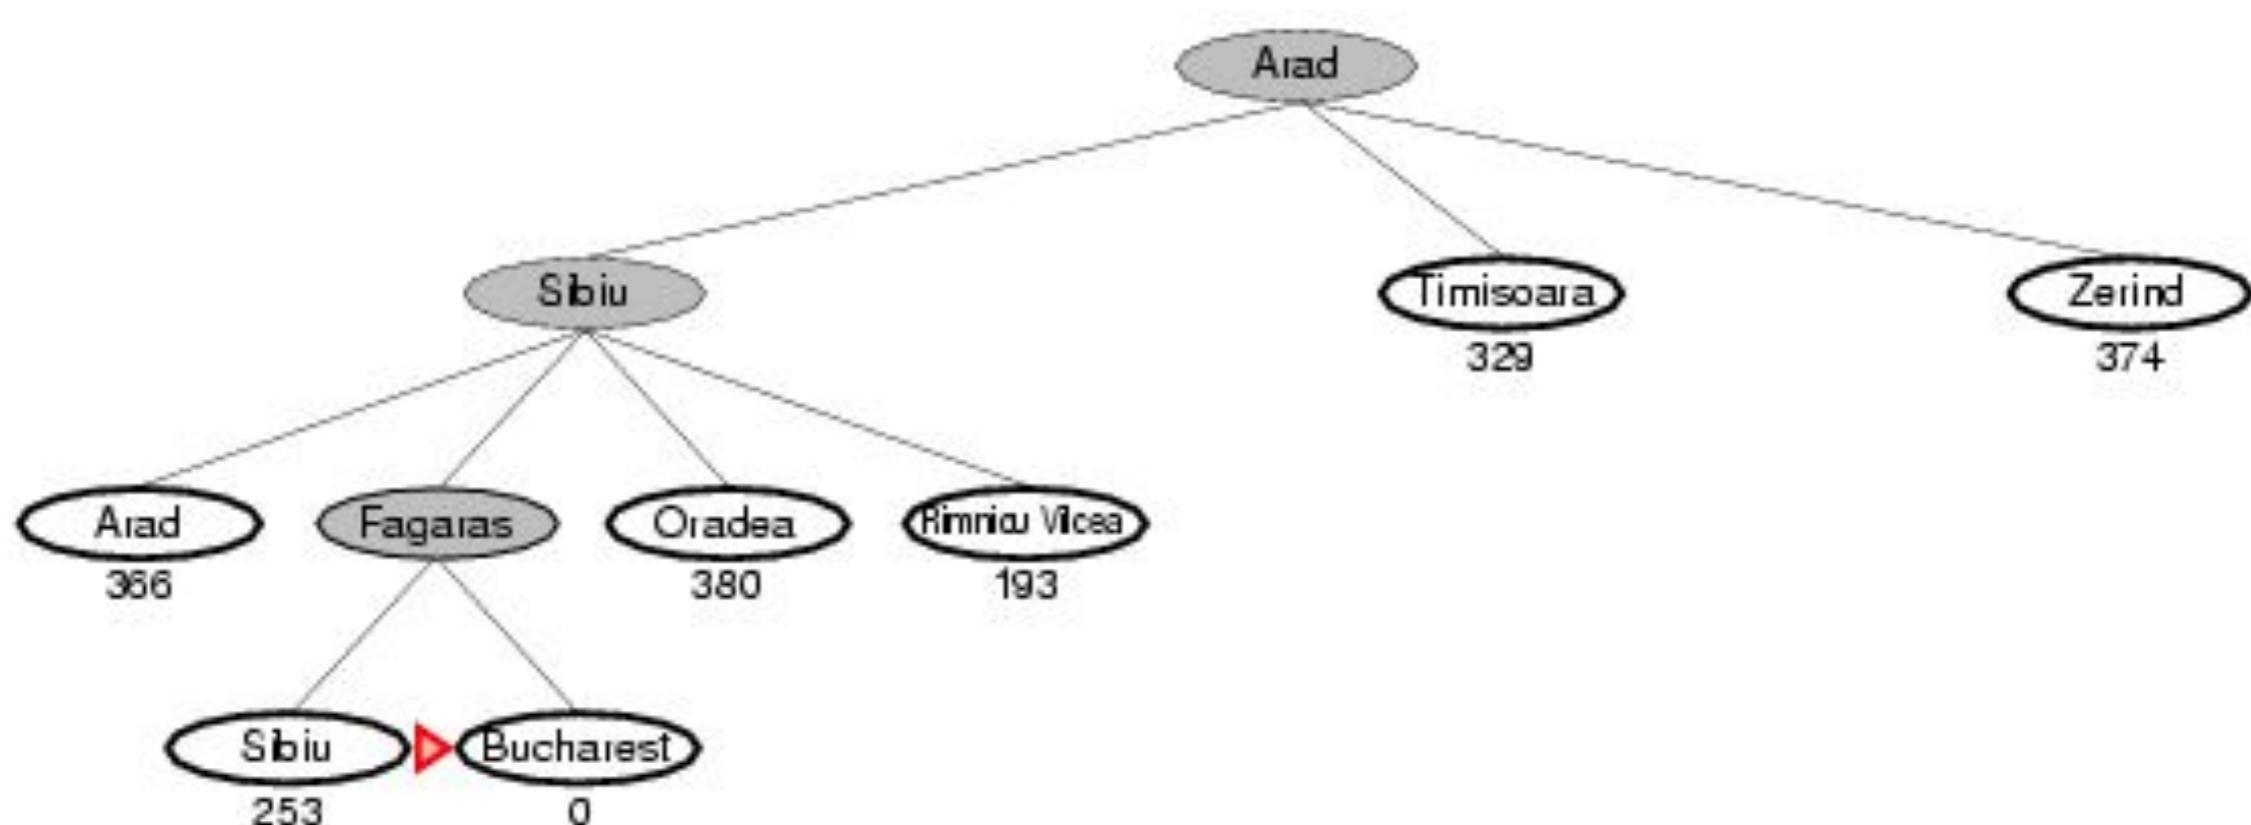
# Greedy best-first example



# Greedy best-first example



# Greedy best-first example



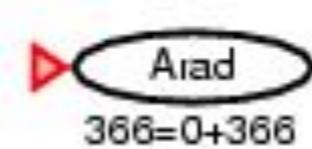
# Greedy best-first

- *Complete?* No – can get stuck in loops, e.g., Lugoj > Mehadia > Lugoj > Mehadia >...
- *Time?*  $O(b^m)$ , but a good heuristic can give dramatic improvement
- *Space?*  $O(b^m)$  -- keeps all nodes in memory
- *Optimal?* No

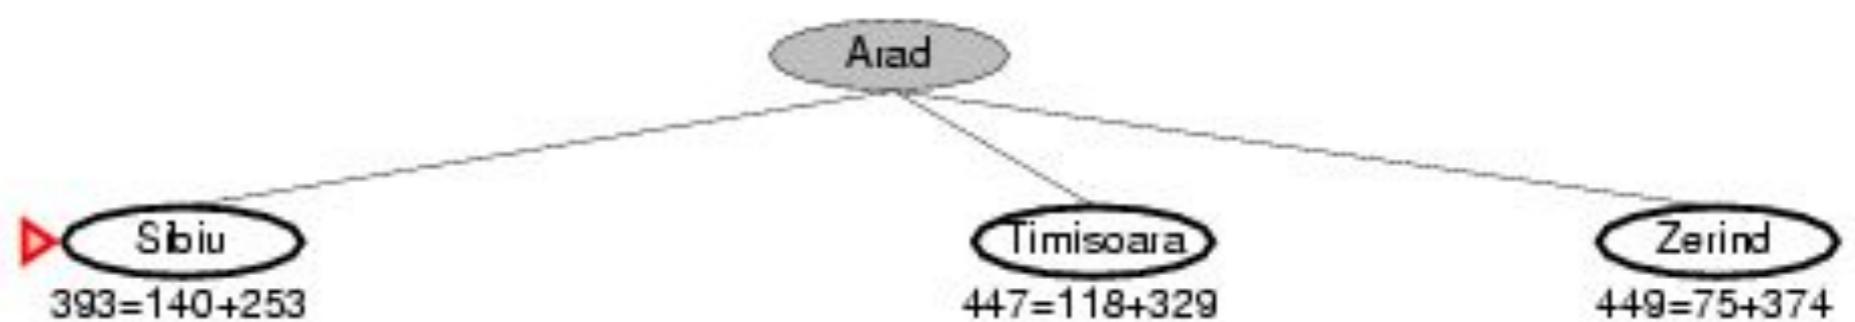
# A\* search

- Idea: avoid expanding paths that are already expensive
- Evaluation function  $f(n) = g(n) + h(n)$
- $g(n)$  = cost so far to reach n
- $h(n)$  = estimated cost from n to goal
- $f(n)$  = estimated total cost of path through n to goal

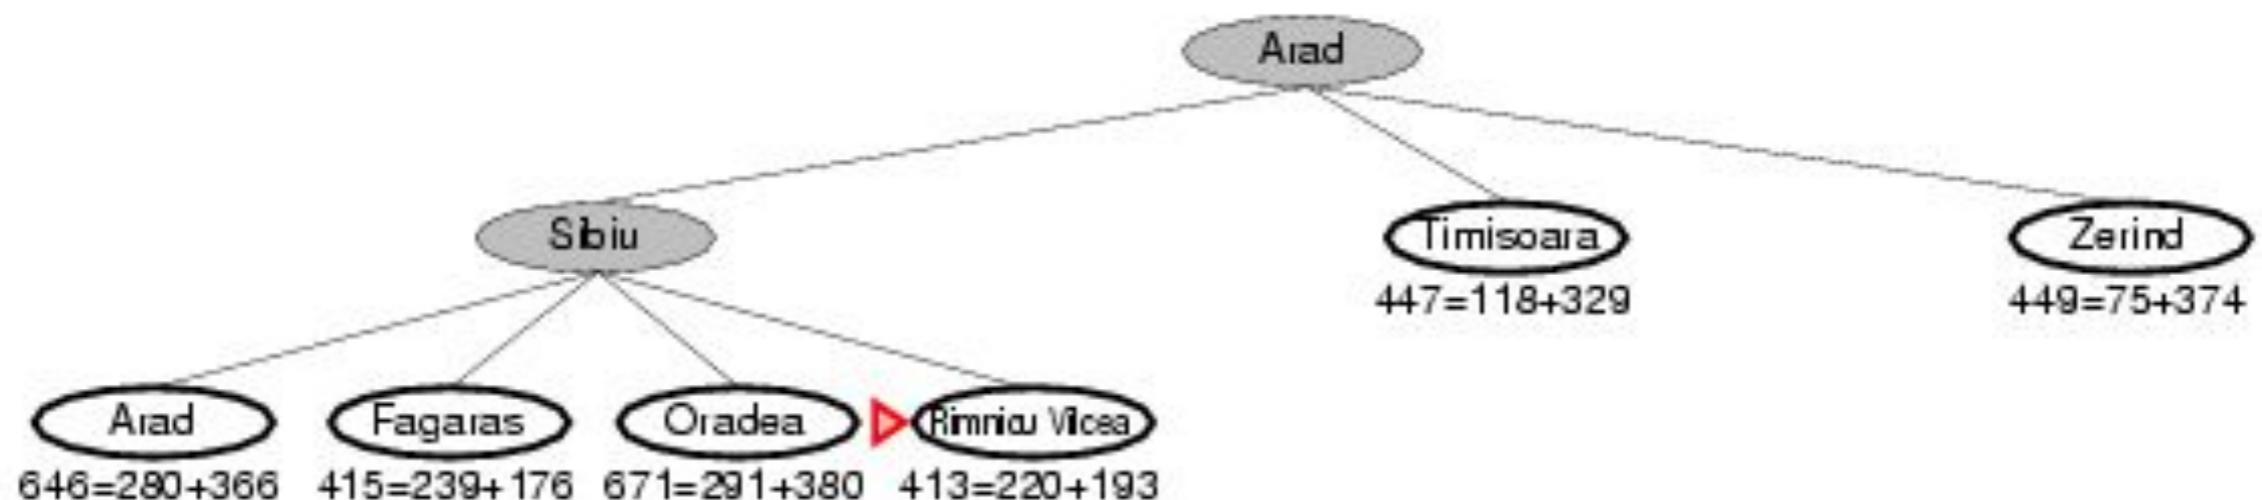
# A\* search example



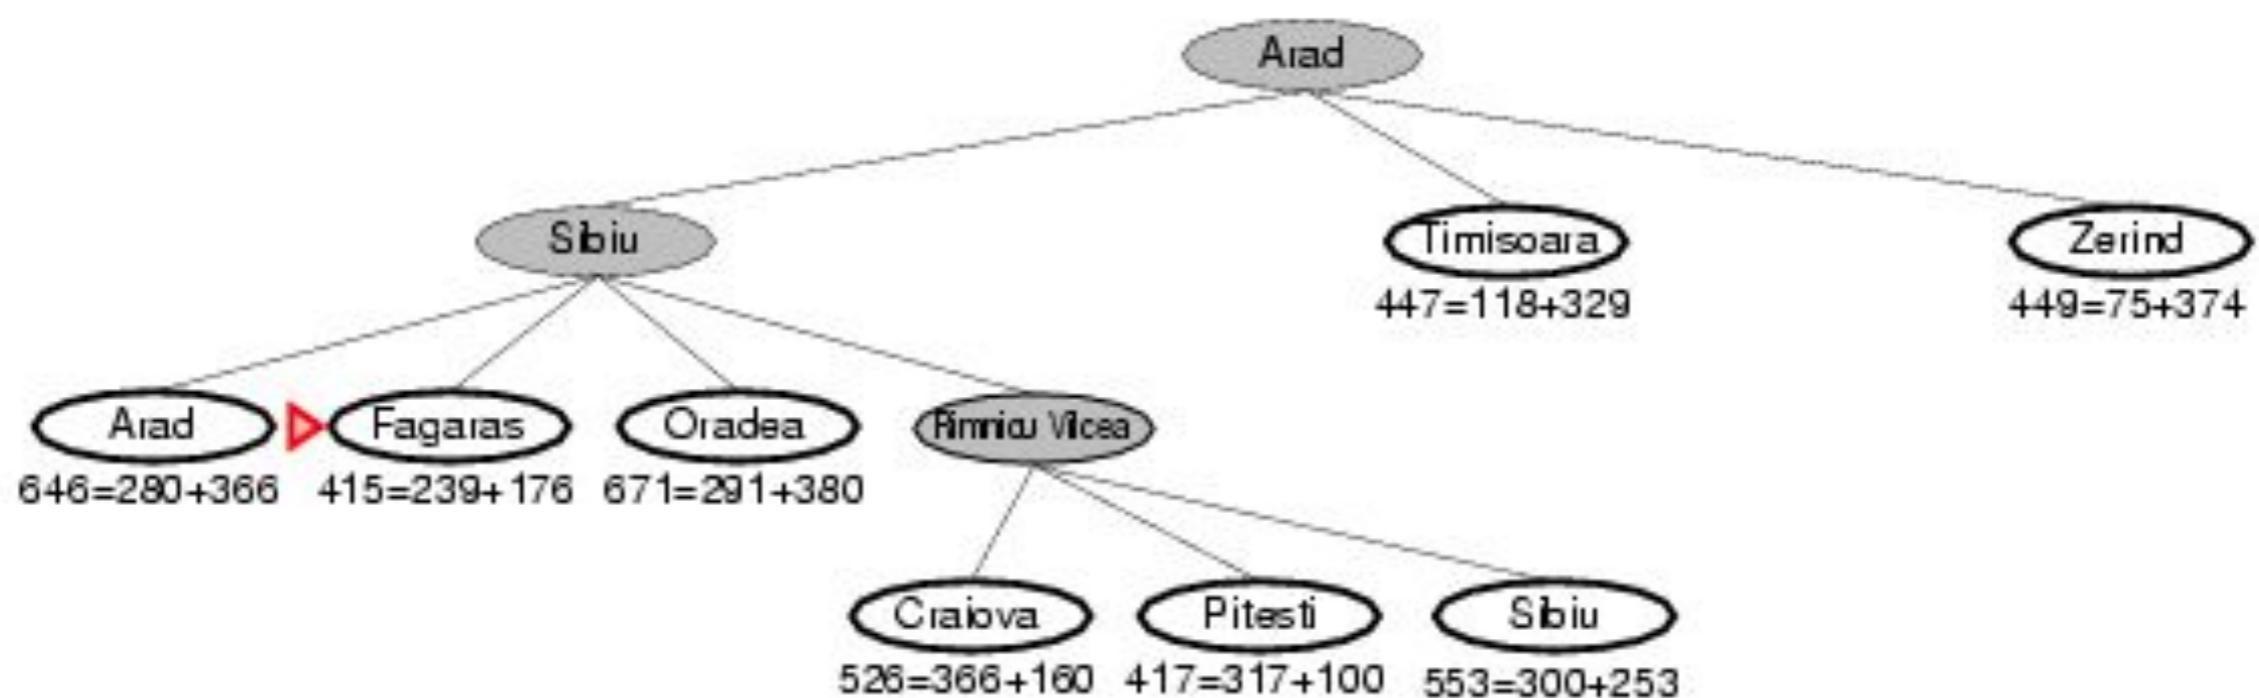
# A\* search example



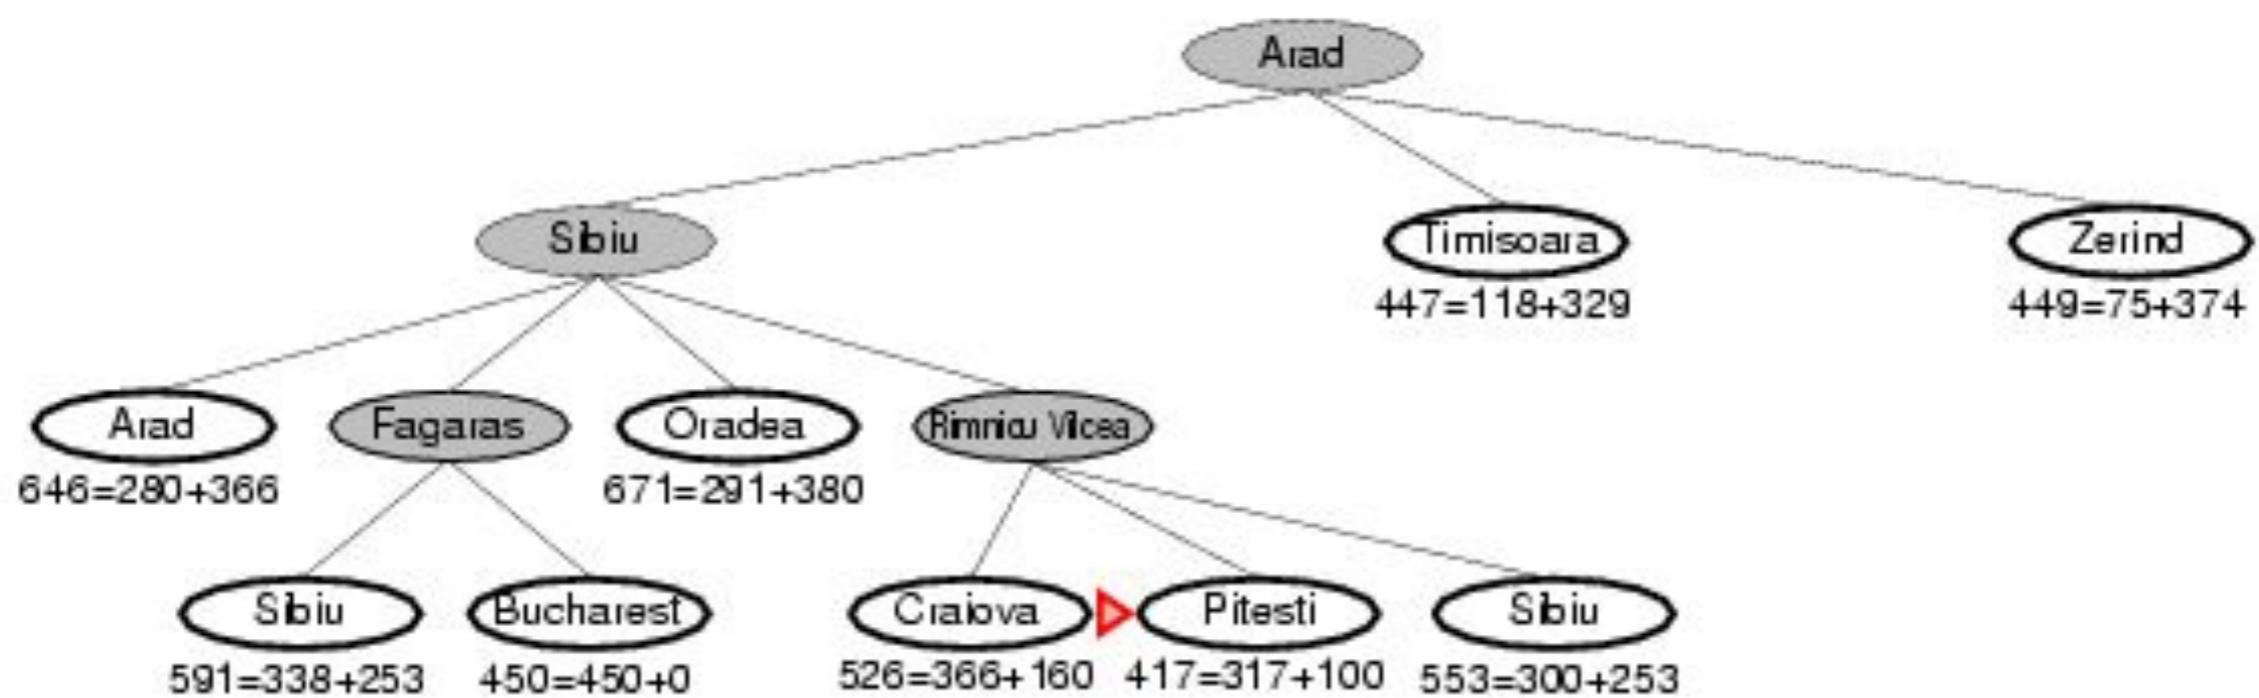
# A\* search example



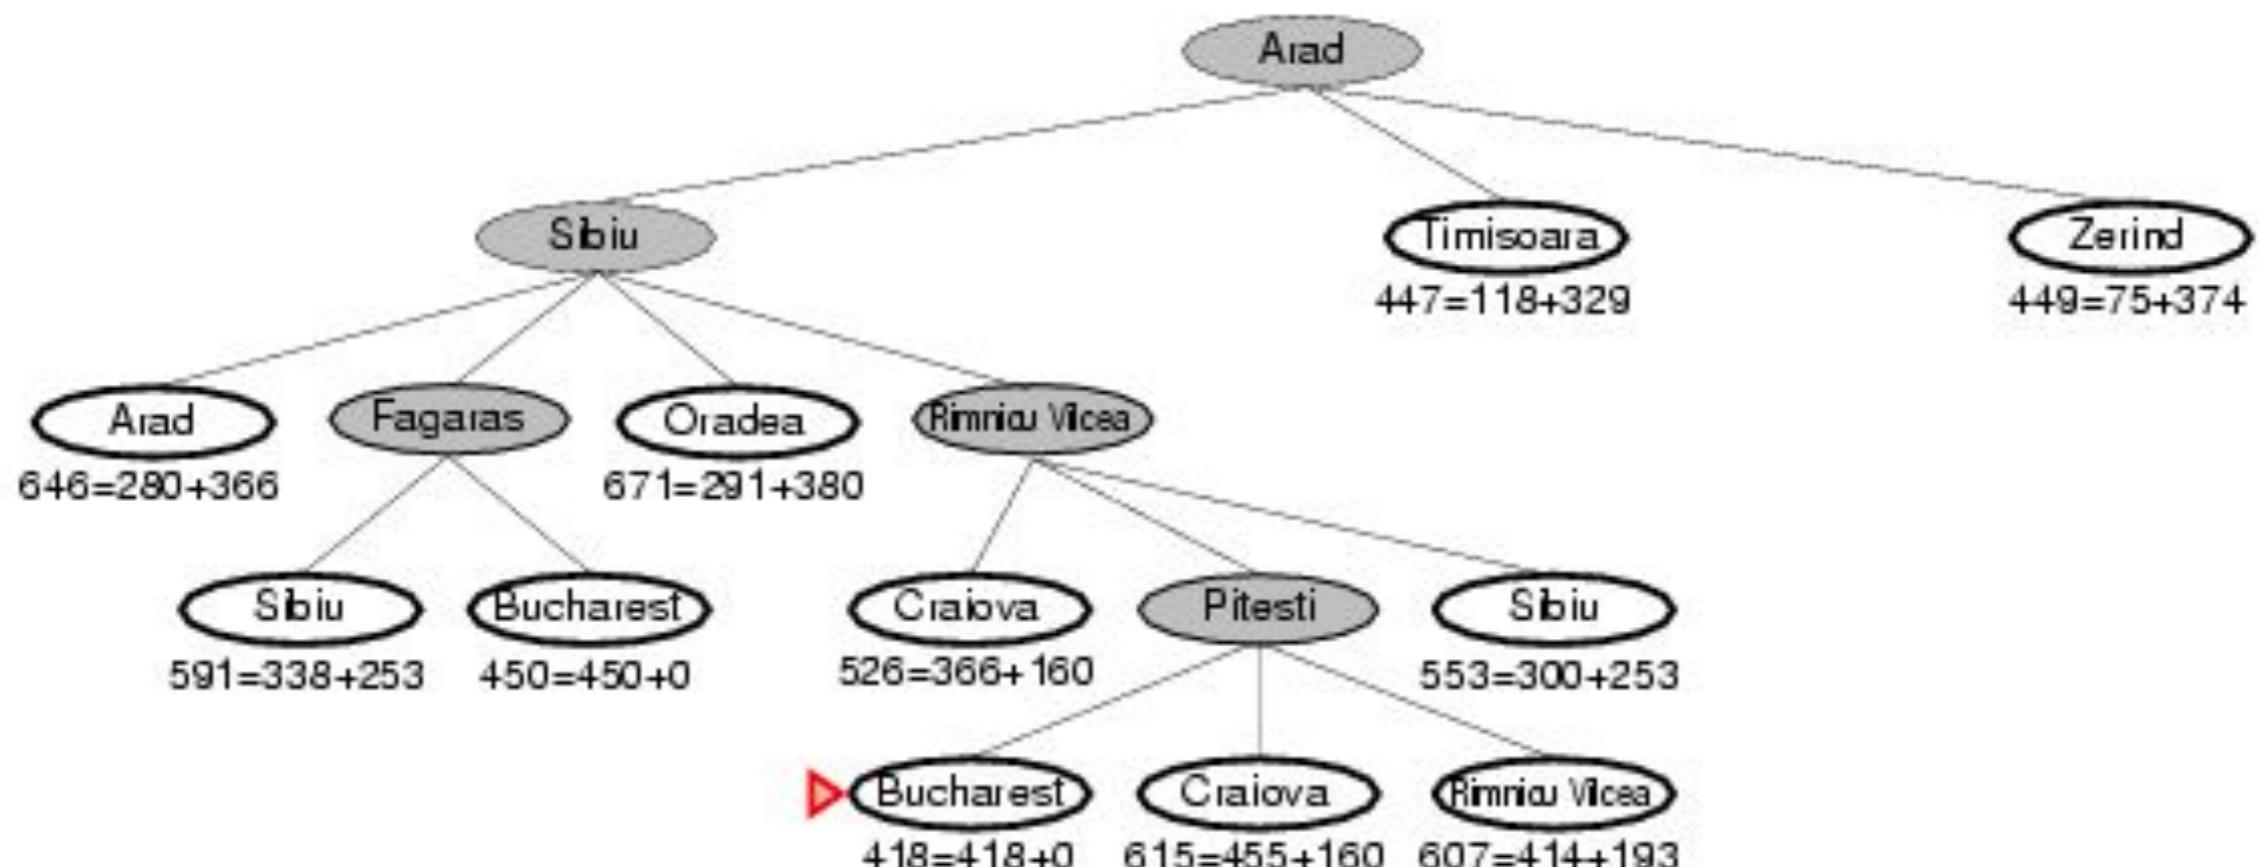
# A\* search example



# A\* search example



# A\* search example



1. Create a search graph  $G$ , consisting solely of the start node,  $sn$ . Put  $sn$  on a list called OPEN.
2. Create a list called CLOSED that is initially empty.
3. If OPEN is empty, exit with failure.
4. Select the first node on OPEN, remove it from OPEN, and put it on CLOSED. Call this node  $n$ .
5. If  $n$  is a goal node, exit successfully with the solution obtained by tracing a path along the pointers from  $n$  to  $sn$  in  $G$ . (The pointers define a search tree and are established in Step 7.)
6. Expand node  $n$ , generating the set  $M$ , of its successors that are not already ancestors of  $n$  in  $G$ . Install these members of  $M$  as successors of  $n$  in  $G$ .
7. Establish a pointer to  $n$  from each of those members of  $M$  that were not already in  $G$  (i.e., not already on either OPEN or CLOSED). Add these members of  $M$  to OPEN. For each member,  $m$ , of  $M$  that was already on OPEN or CLOSED, redirect its pointer to  $n$  if the best path to  $m$  found so far is through  $n$ . For each member of  $M$  already on CLOSED, redirect the pointers of each of its descendants in  $G$  so that they point backward along the best paths found so far to these descendants.
8. Reorder the list OPEN in order of increasing  $f$  values. (Ties among minimal  $f$  values are resolved in favor of the deepest node in the search tree.)
9. Go to Step 3.

# In-class exercise

# Admissible heuristics

- A heuristic  $h(n)$  is admissible if for every node  $n$ ,  $h(n) \leq h^*(n)$ , where  $h^*(n)$  is the true cost to reach the goal state from  $n$ .
- An admissible heuristic never overestimates the cost to reach the goal, i.e., it is optimistic
- Example:  $hSLD(n)$  (never overestimates the actual road distance)
- Theorem: If  $h(n)$  is admissible,  $A^*$  is optimal

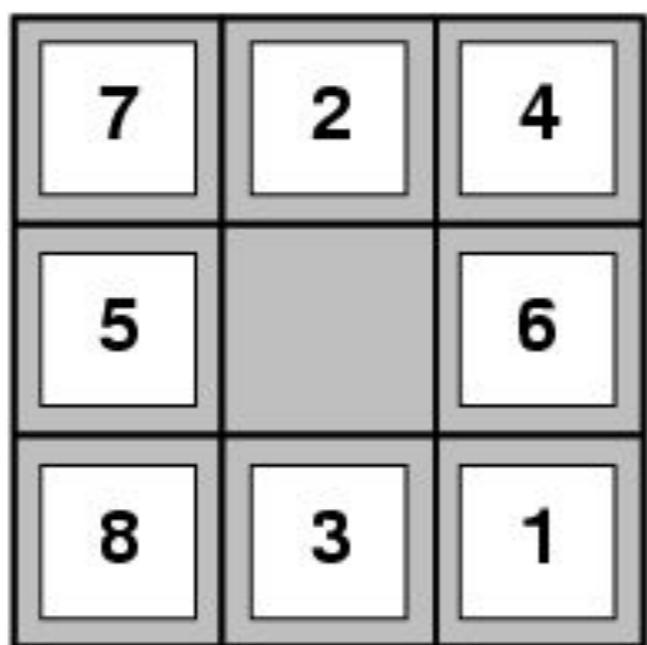
# Admissible heuristics

- E.g., for the 8-puzzle:

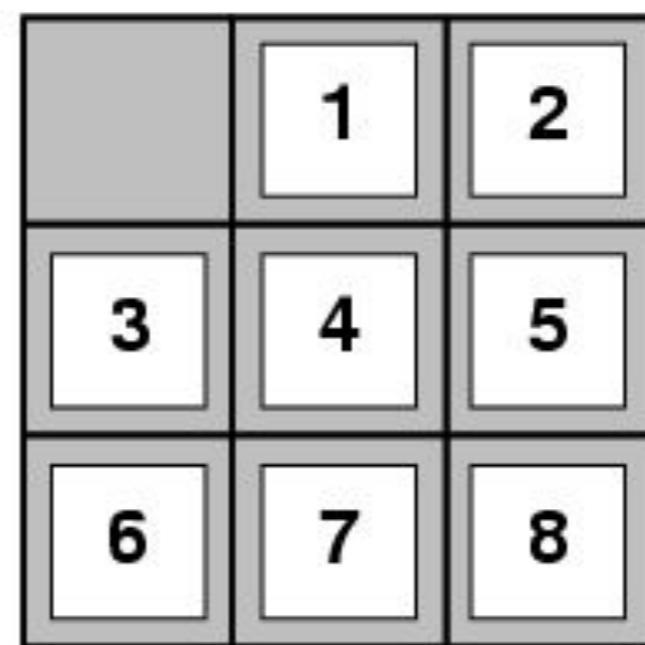
$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total Manhattan distance

(i.e., no. of squares from desired location of each tile)



Start State



Goal State

# Dominance

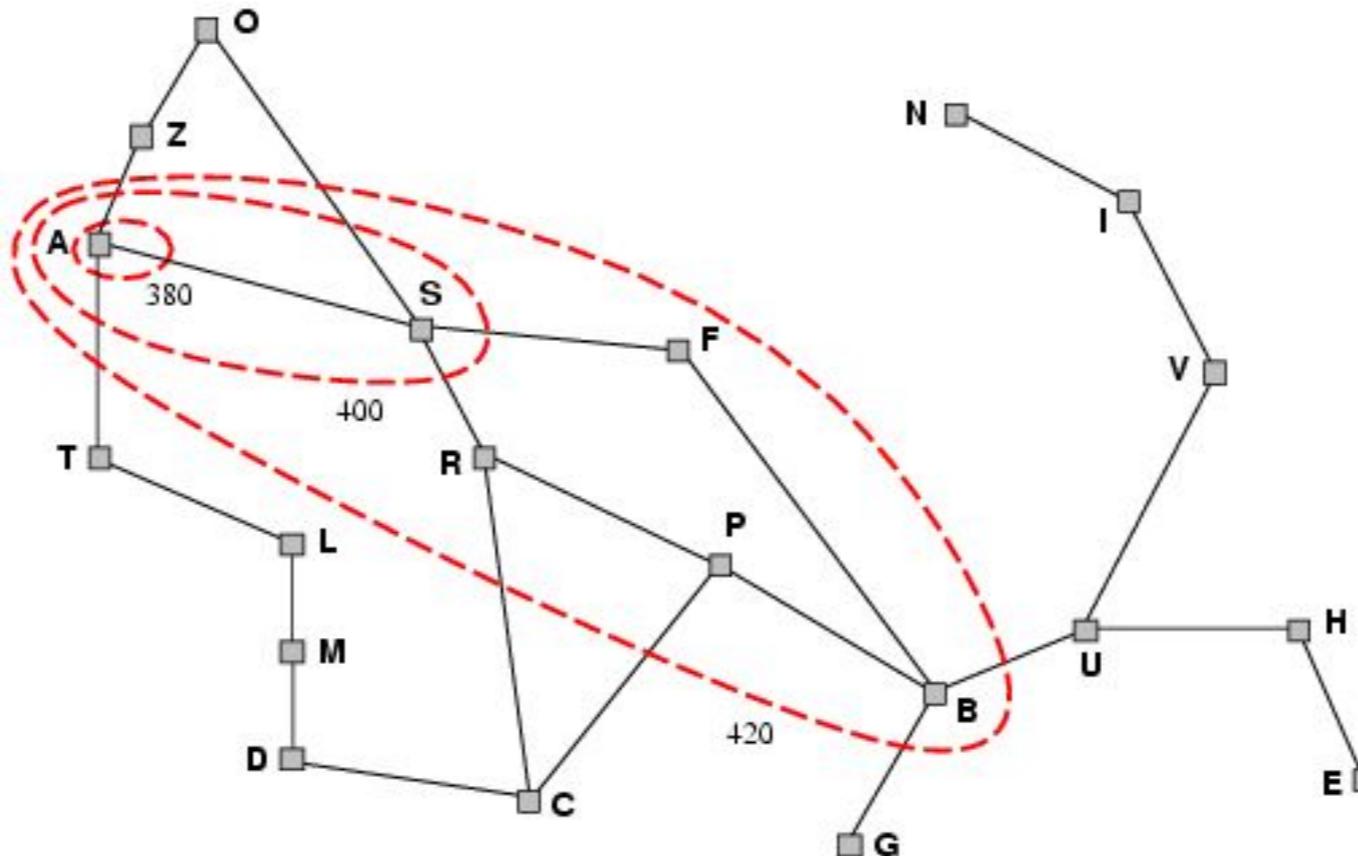
- If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible) then  $h_2$  dominates  $h_1$
- $h_2$  is better for search

# Relaxed problems

- A problem with fewer restrictions on the actions is called a relaxed problem
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move anywhere, then  $h_1(n)$  gives the shortest solution
- If the rules are relaxed so that a tile can move to any adjacent square, then  $h_2(n)$  gives the shortest solution

# Optimality of A\*

- A\* expands nodes in order of increasing f value
- Gradually adds "f-contours" of nodes
- Contour  $i$  has all nodes with  $f=f_i$ , where  $f_i < f_{i+1}$



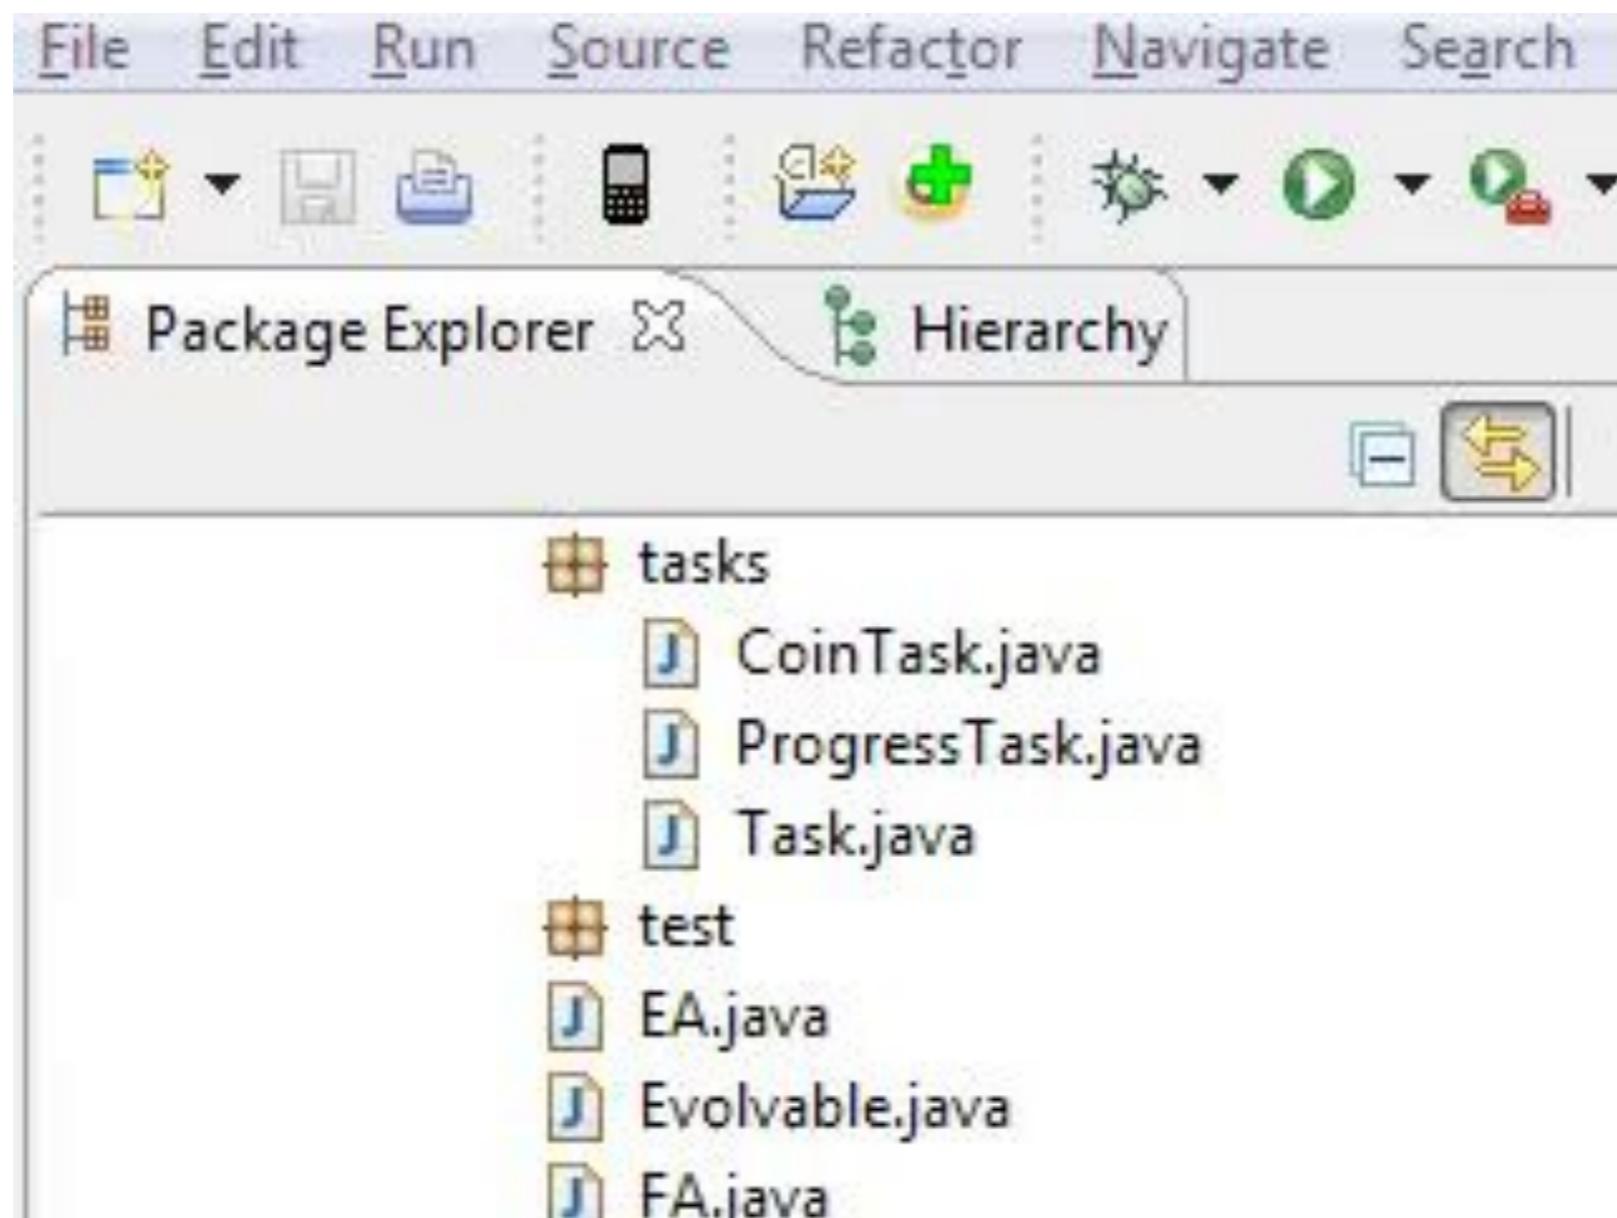
# Properties of A\*

- Complete? Yes (unless there are infinitely many nodes with  $f \leq f(G)$ )
- Time? Exponential in the worst case
- Space? Keeps all nodes in memory
- Optimal? Yes

# Admissible heuristics for...

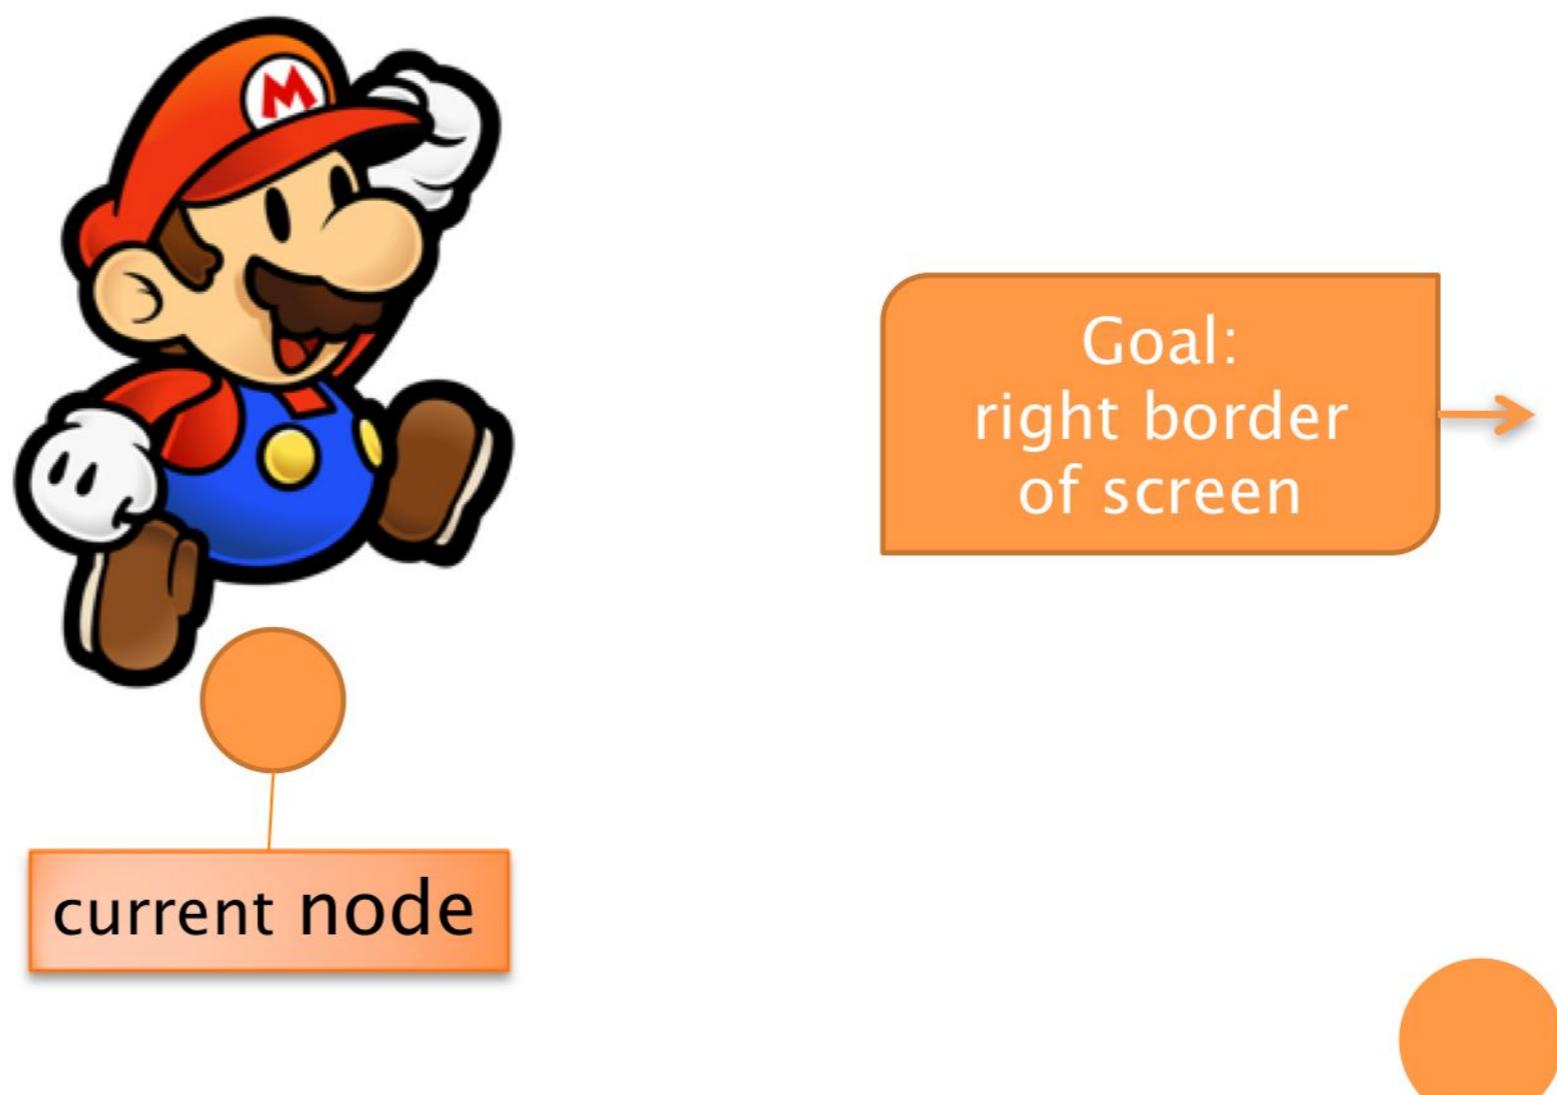


# A\* in Mario

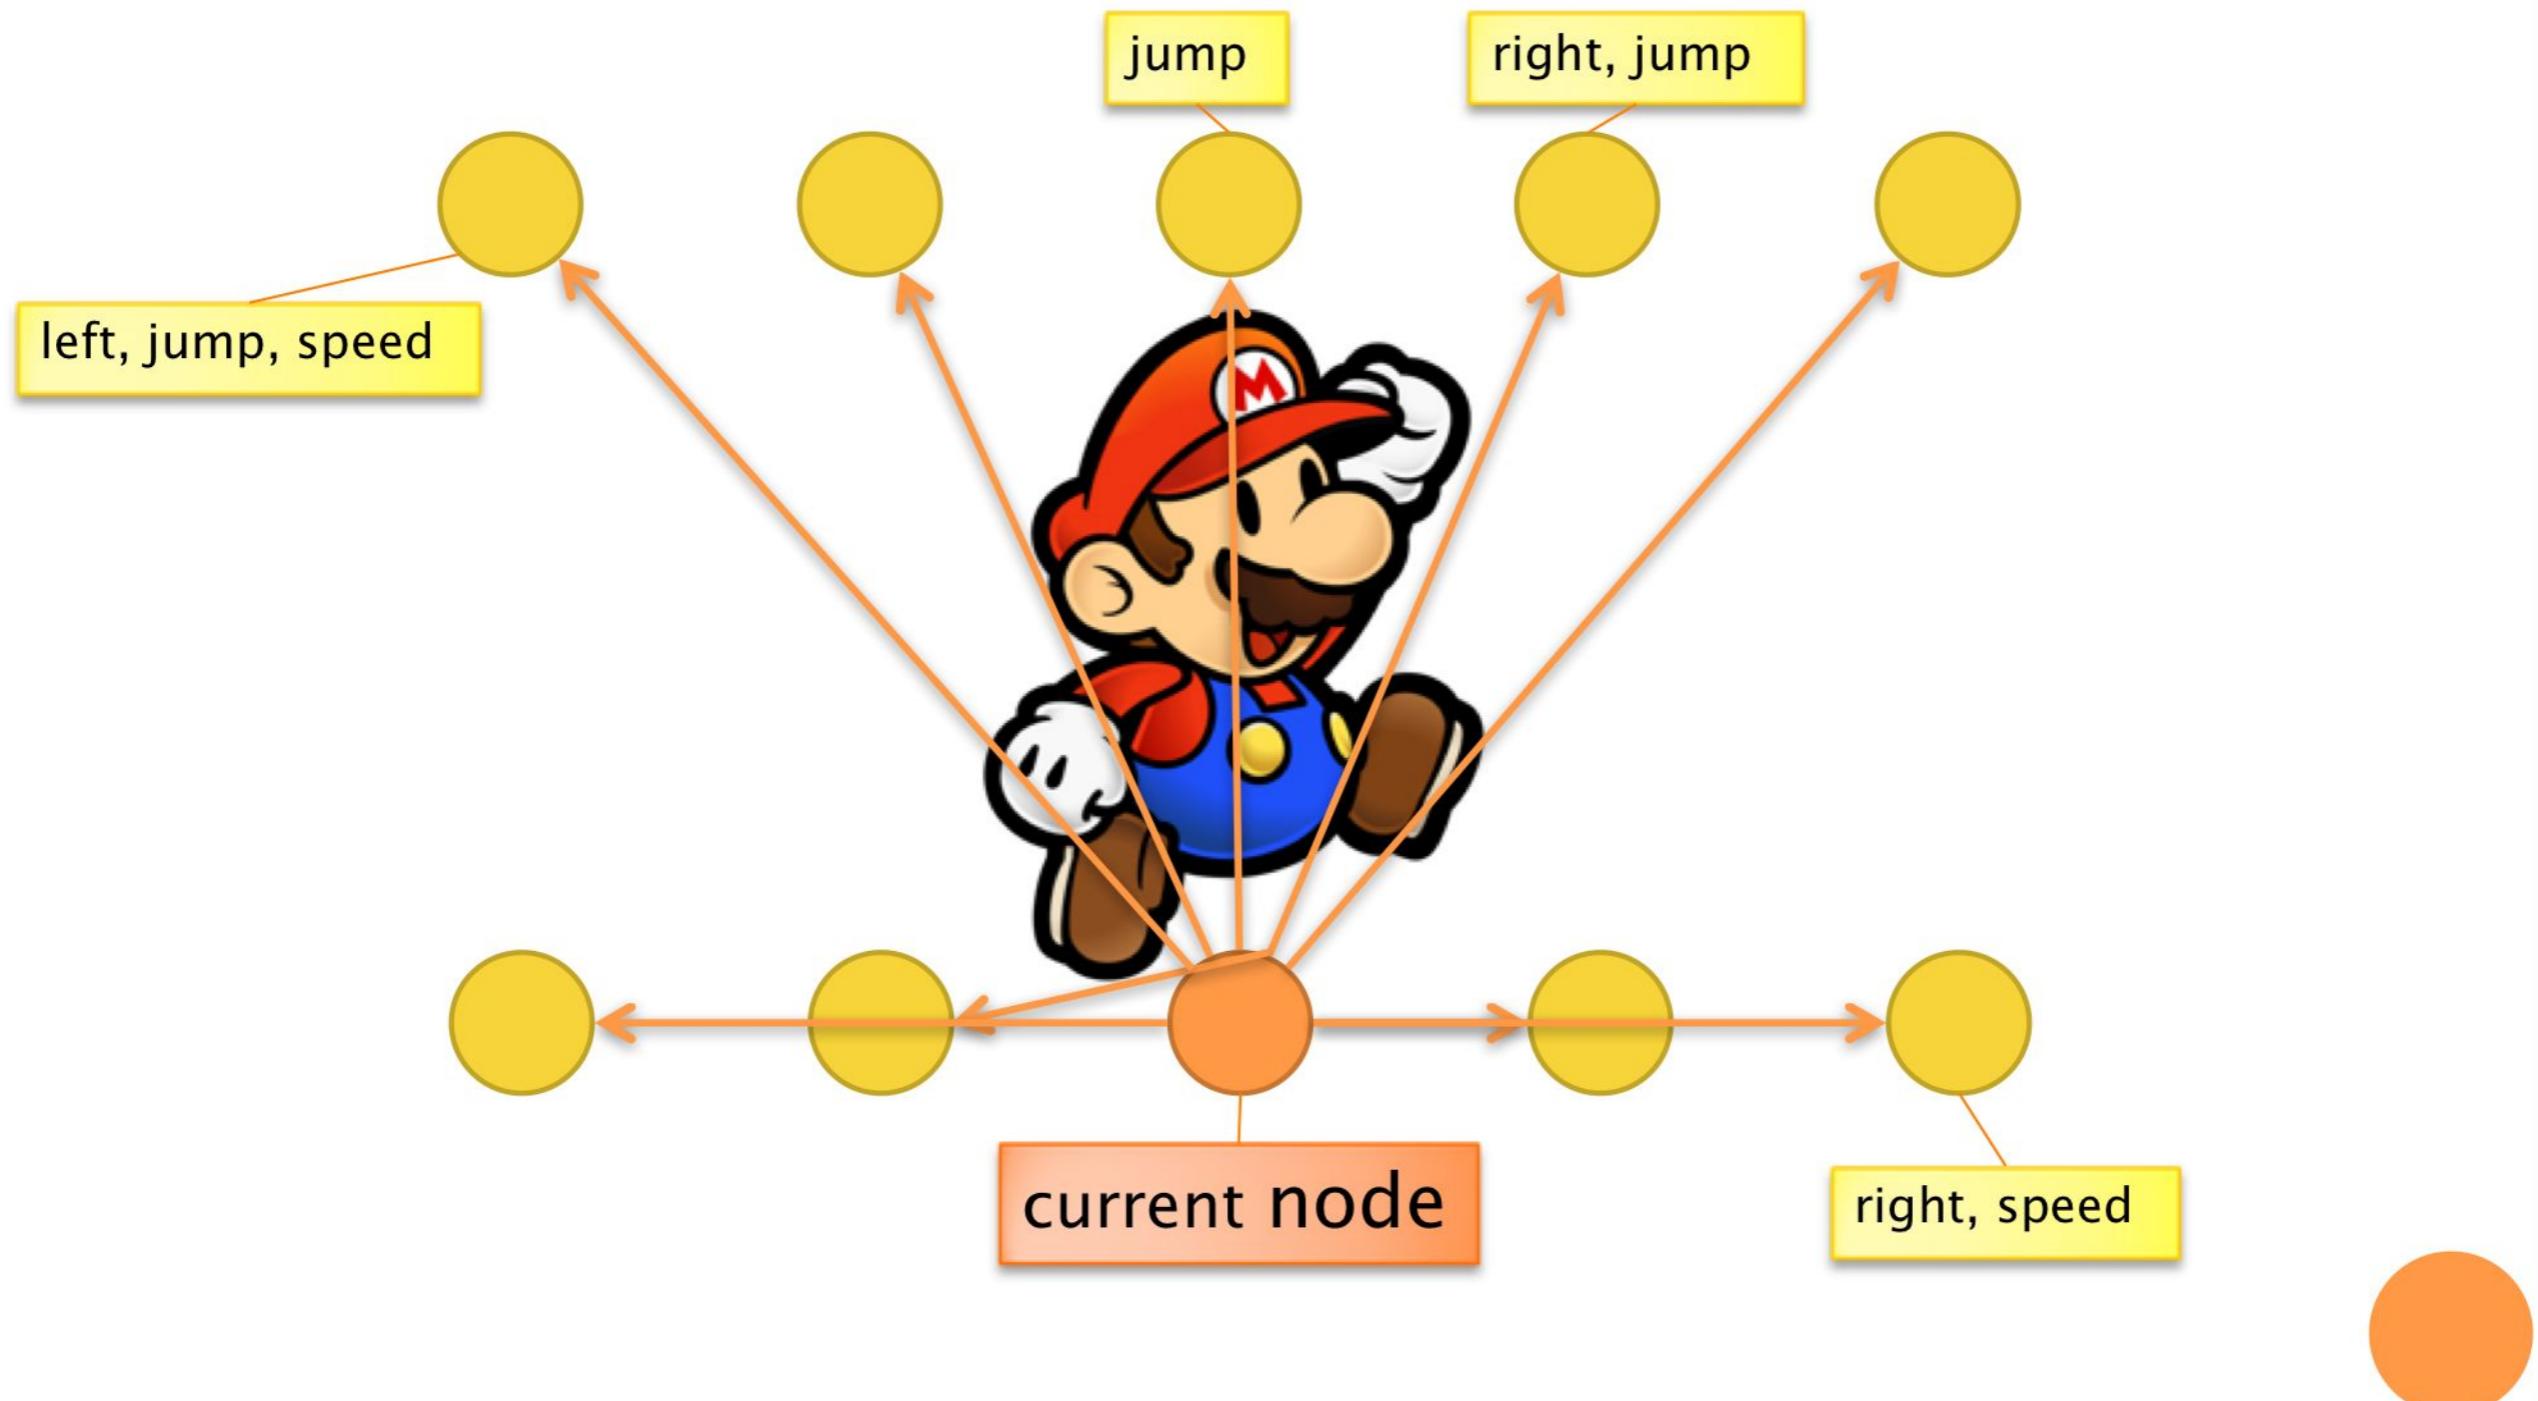


<https://www.youtube.com/watch?v=DlkMs4ZHr8>

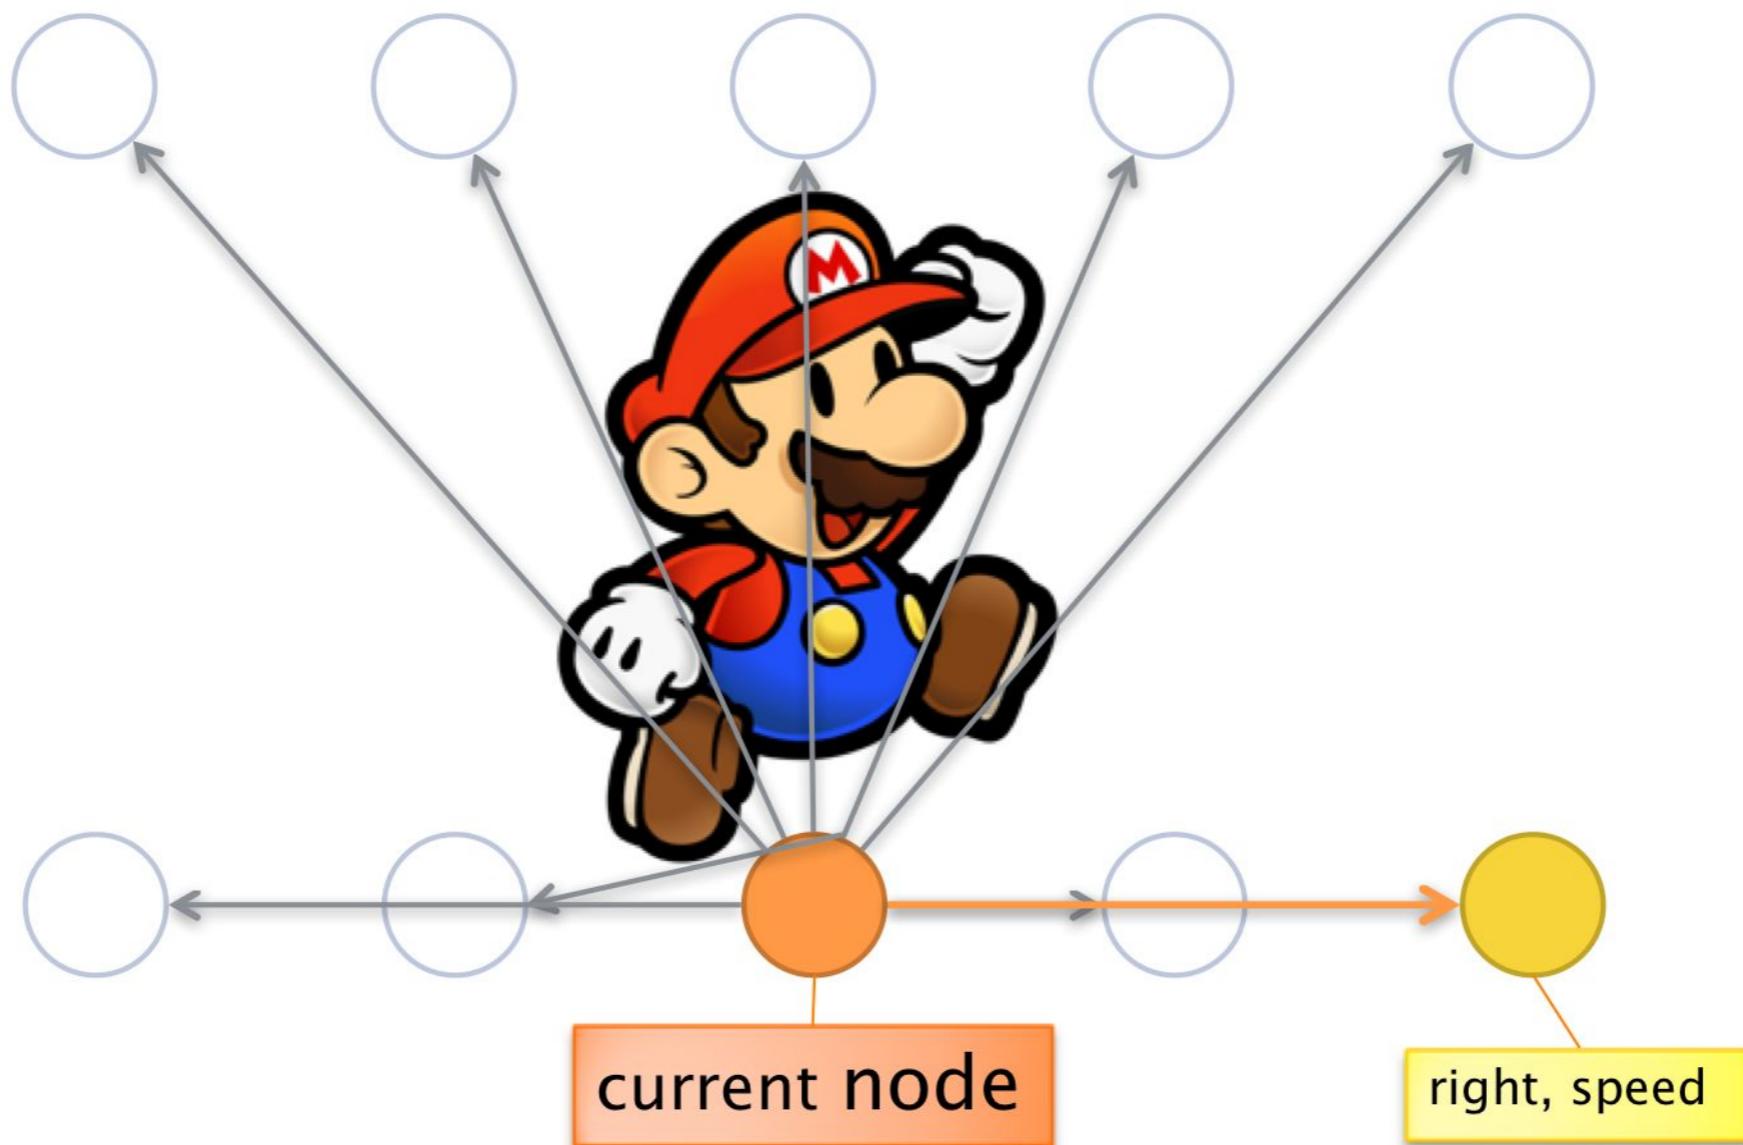
# A\* IN MARIO: CURRENT POSITION



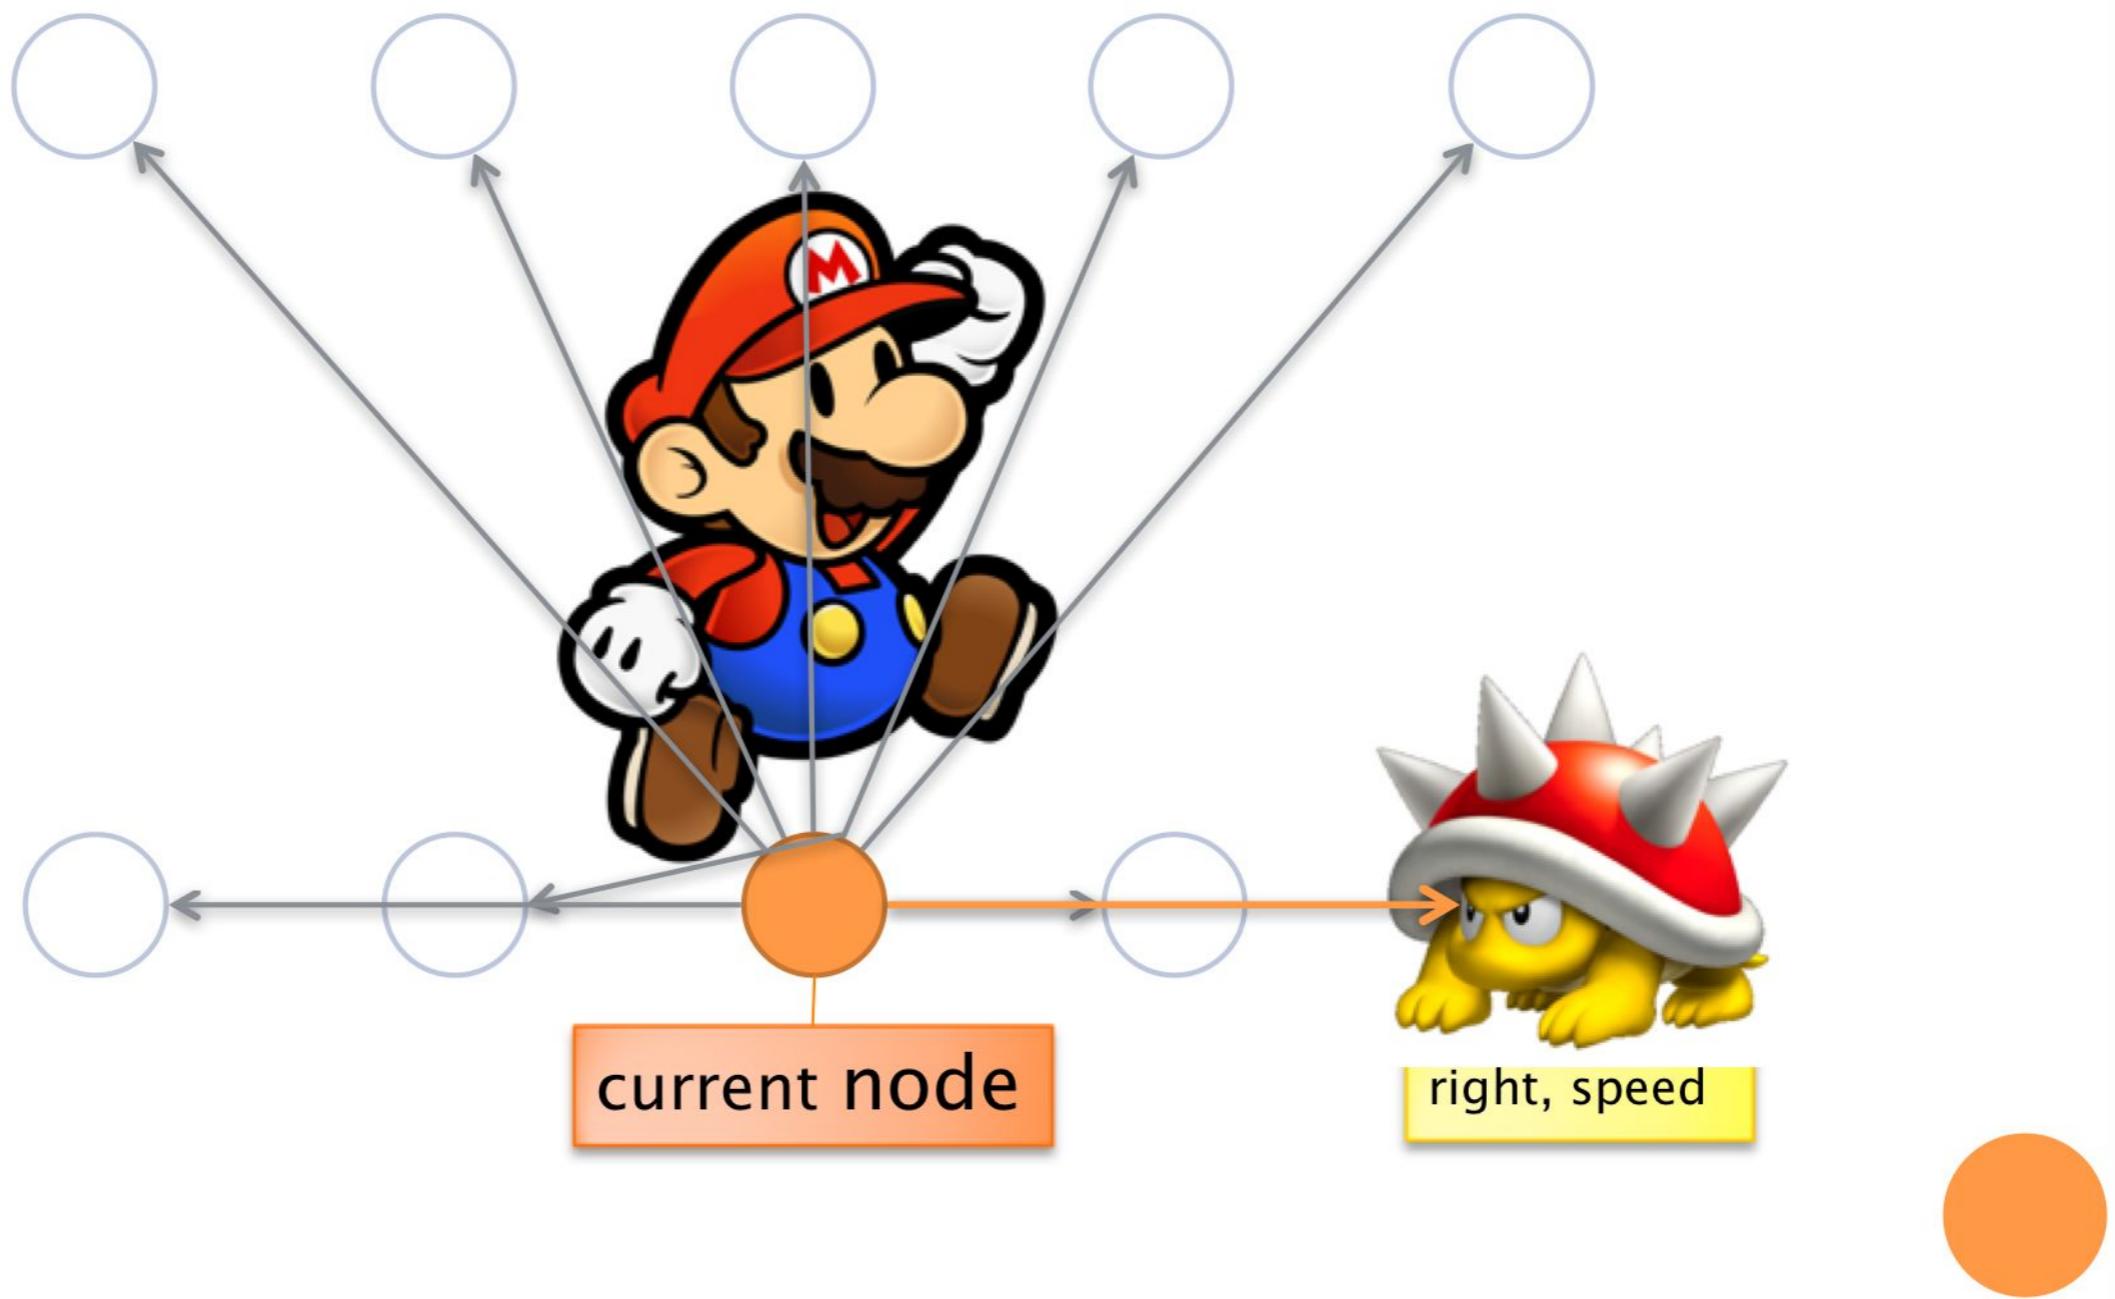
# A\* IN MARIO: CHILD NODES



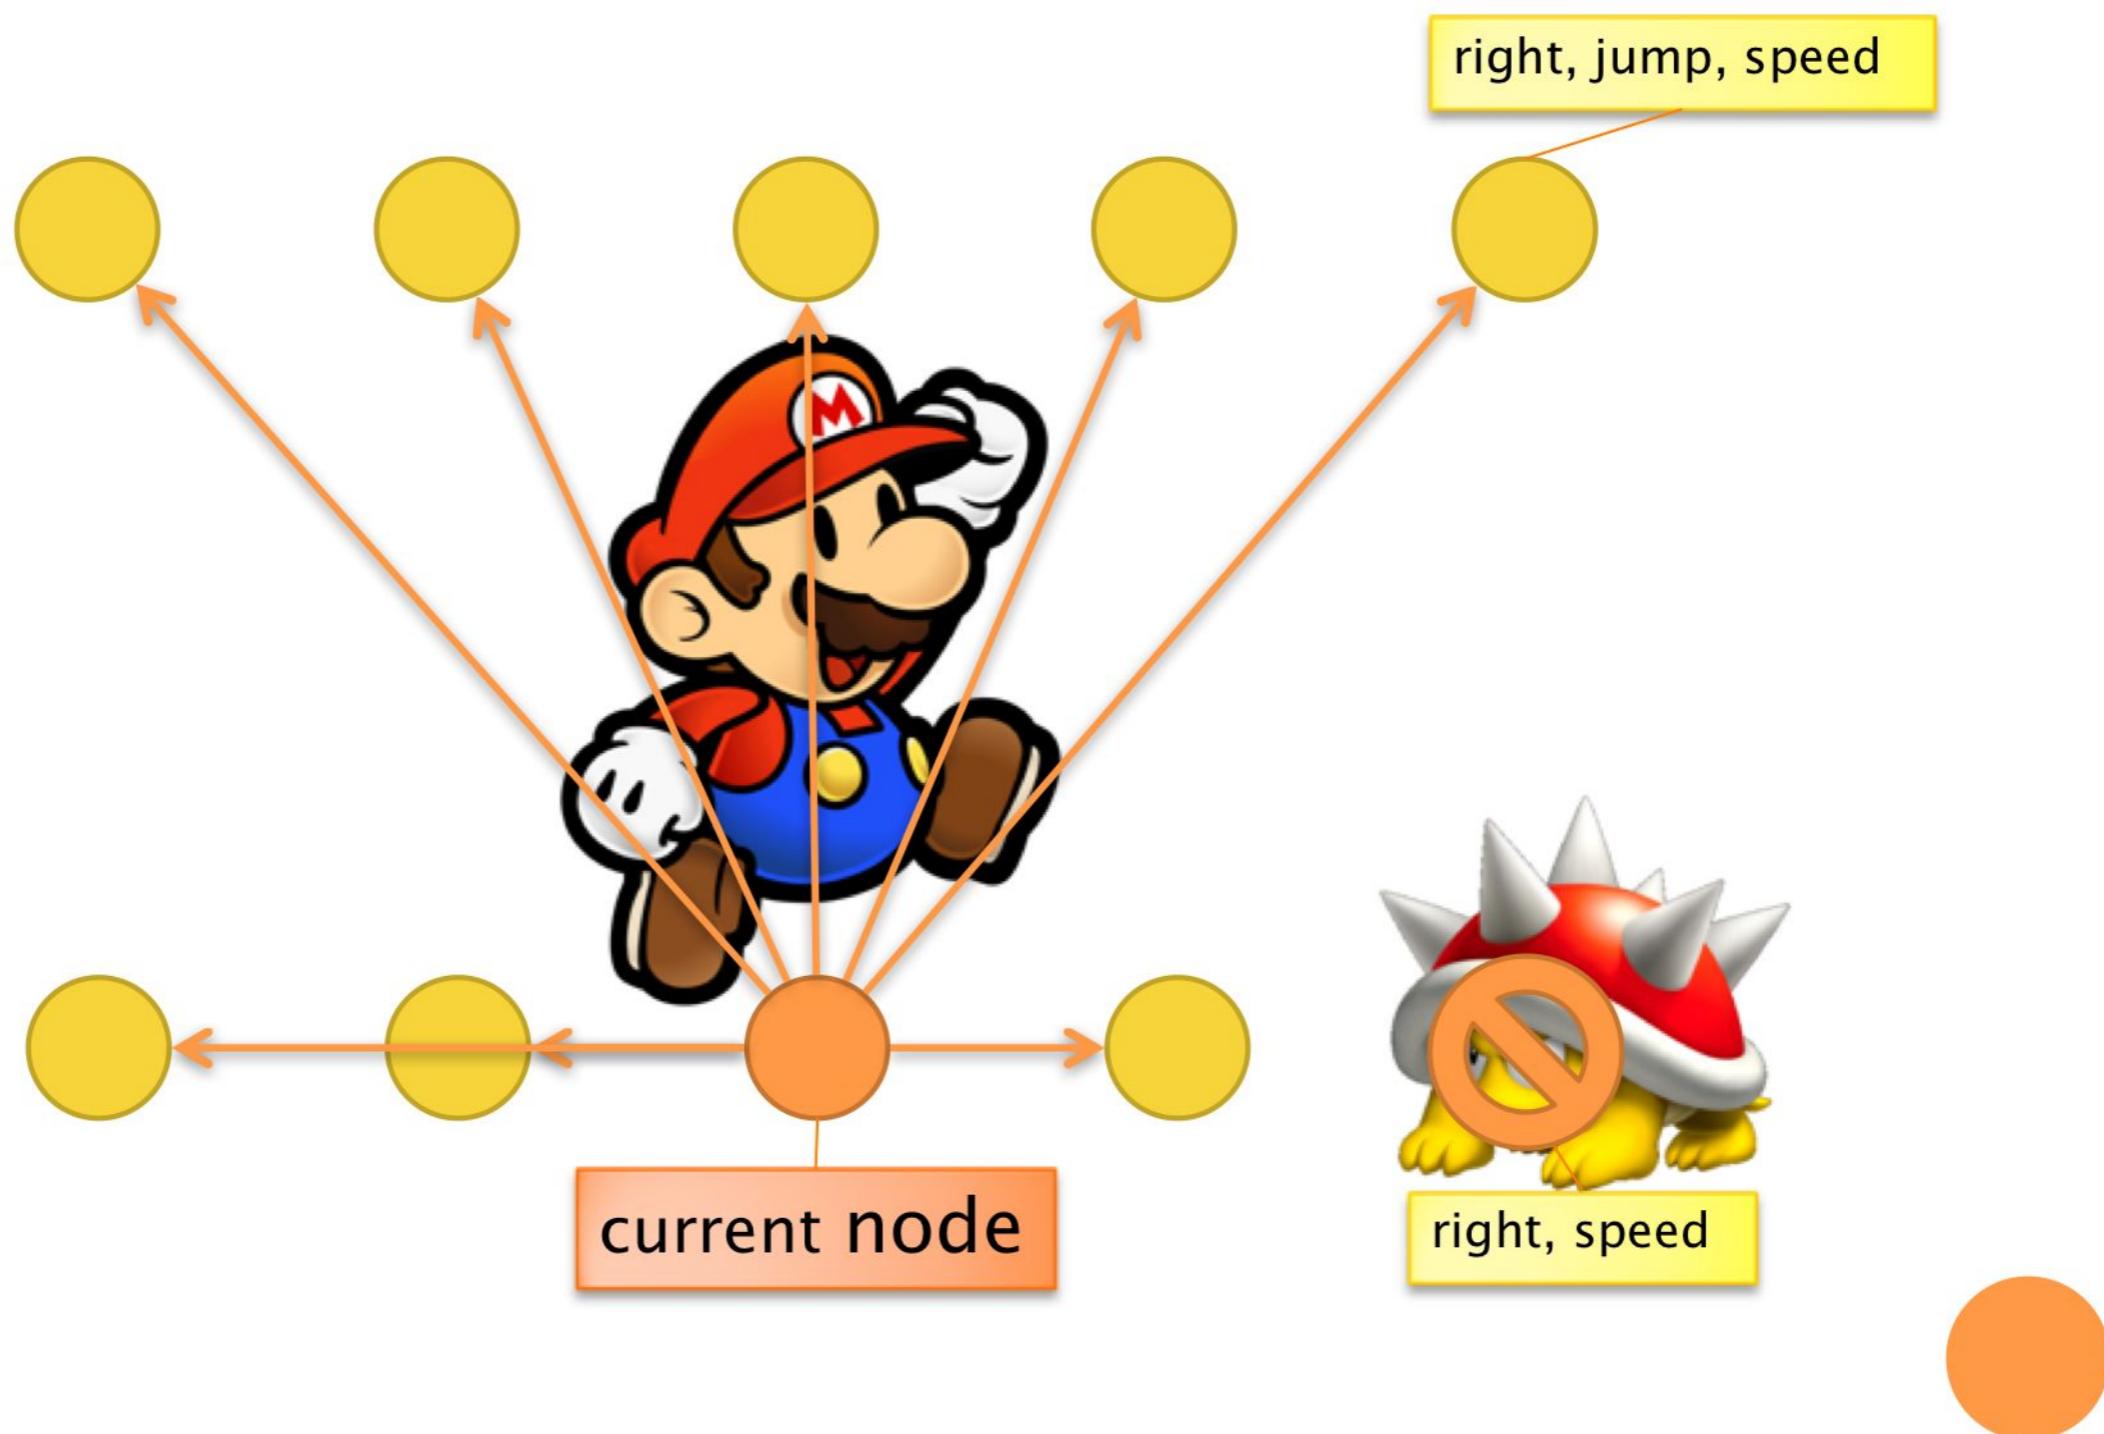
# A\* IN MARIO: BEST FIRST



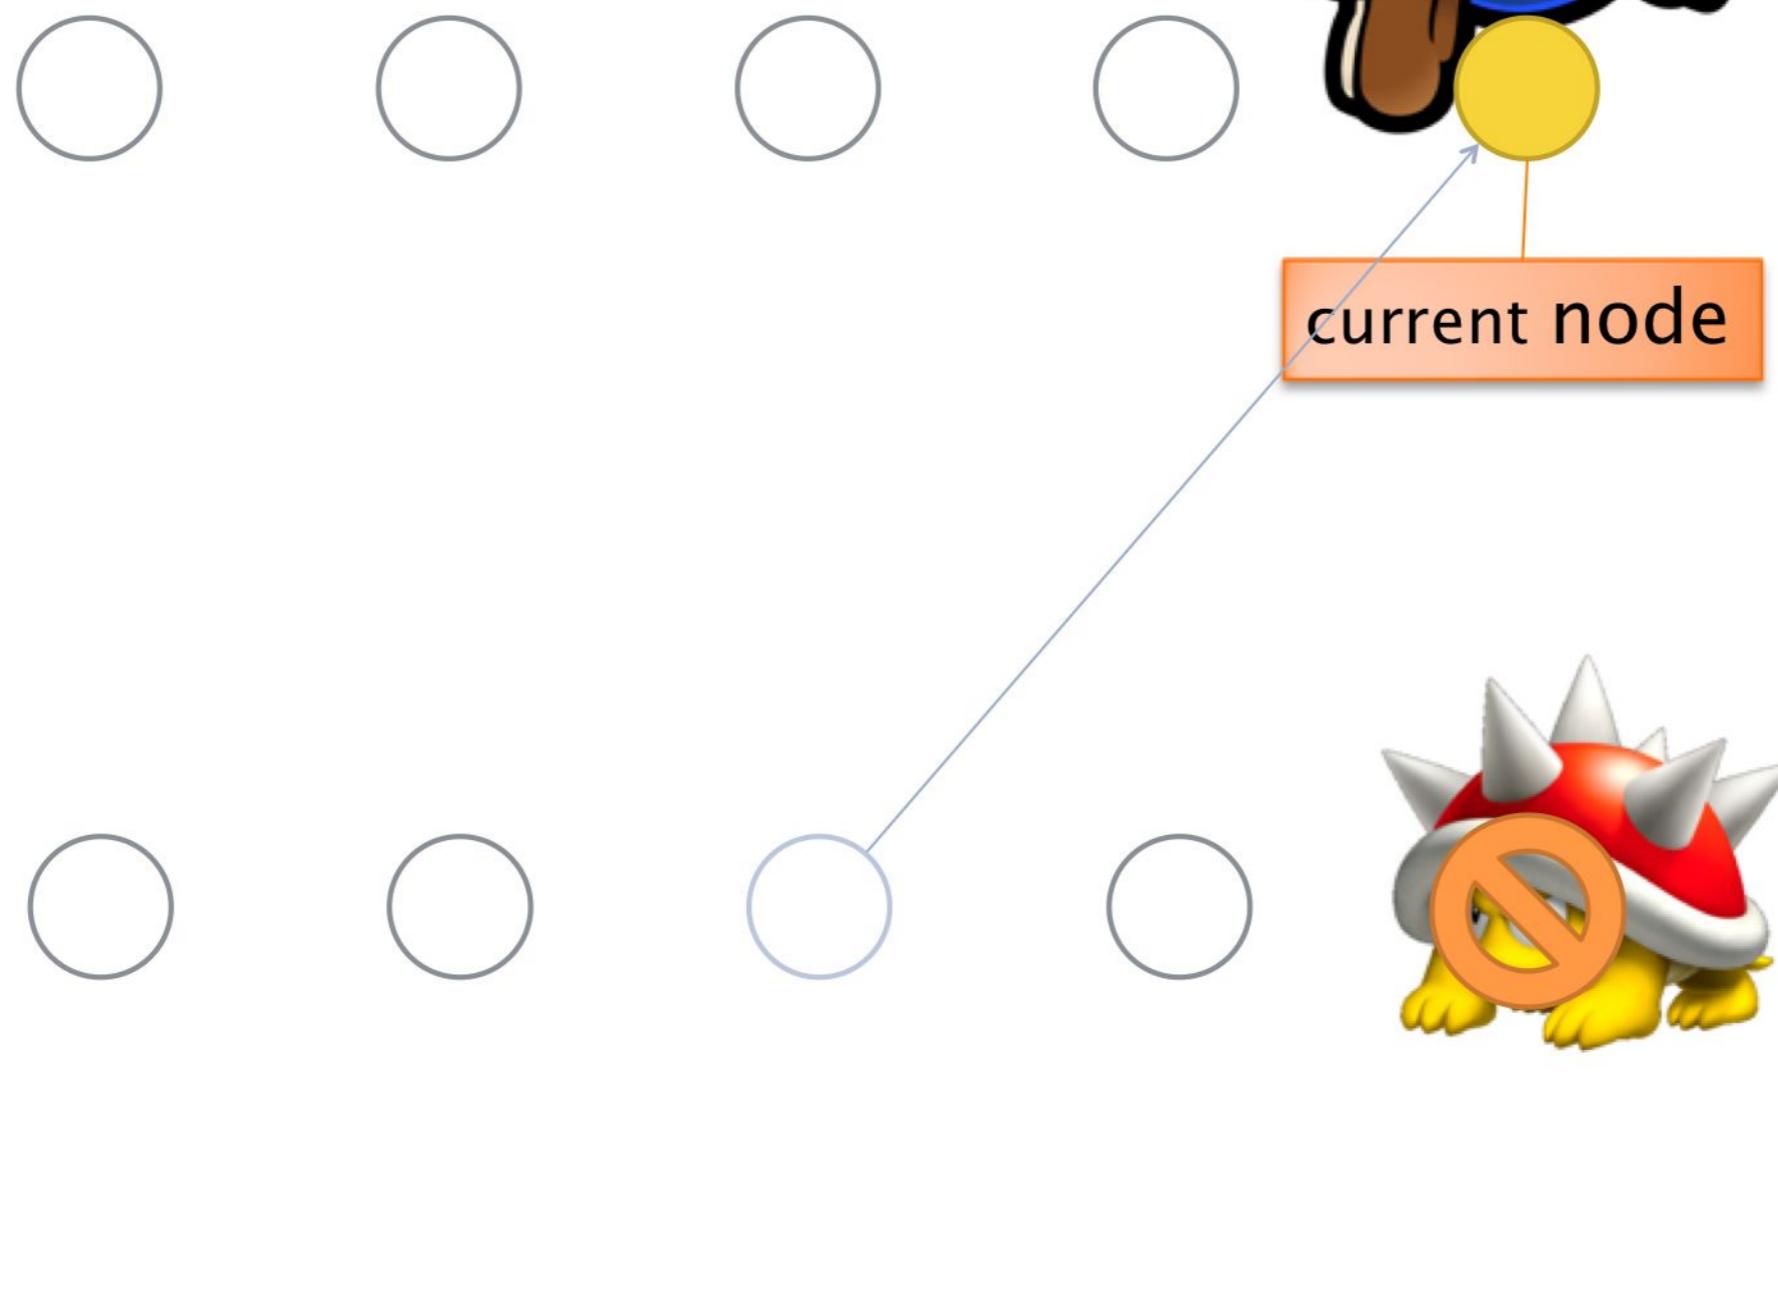
# A\* IN MARIO: EVALUATE NODE



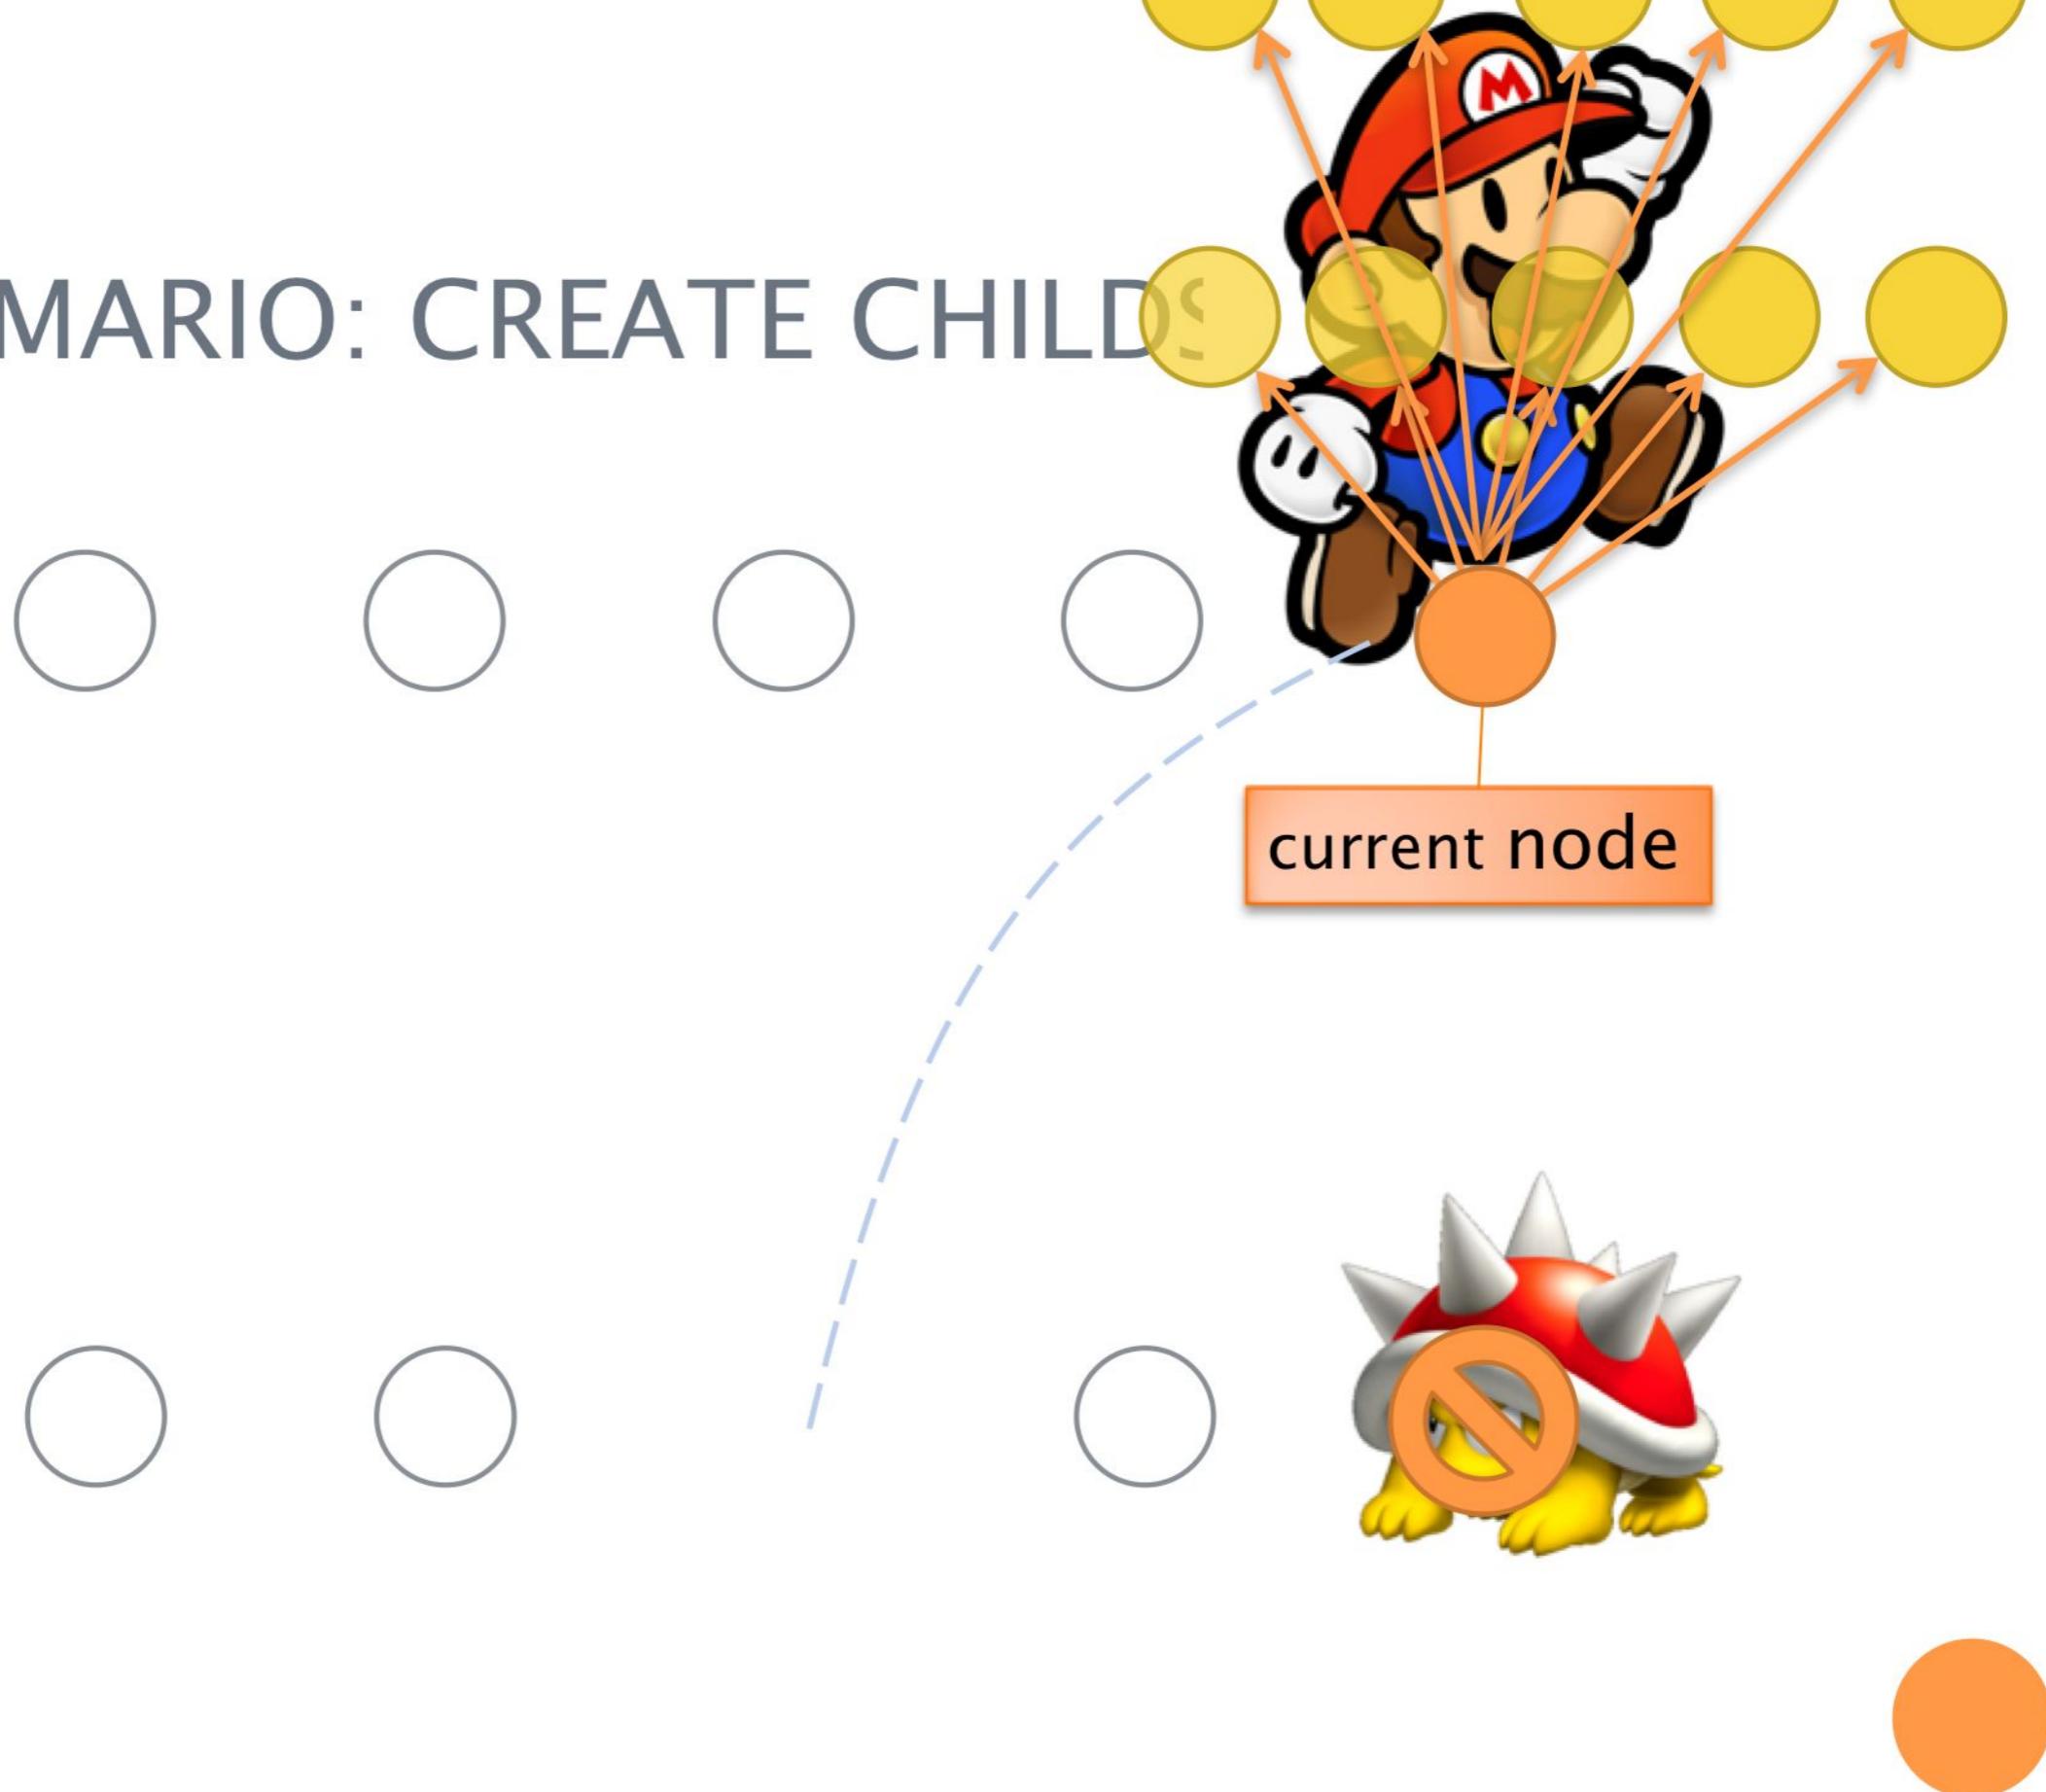
# A\* IN MARIO: BACKTRACK



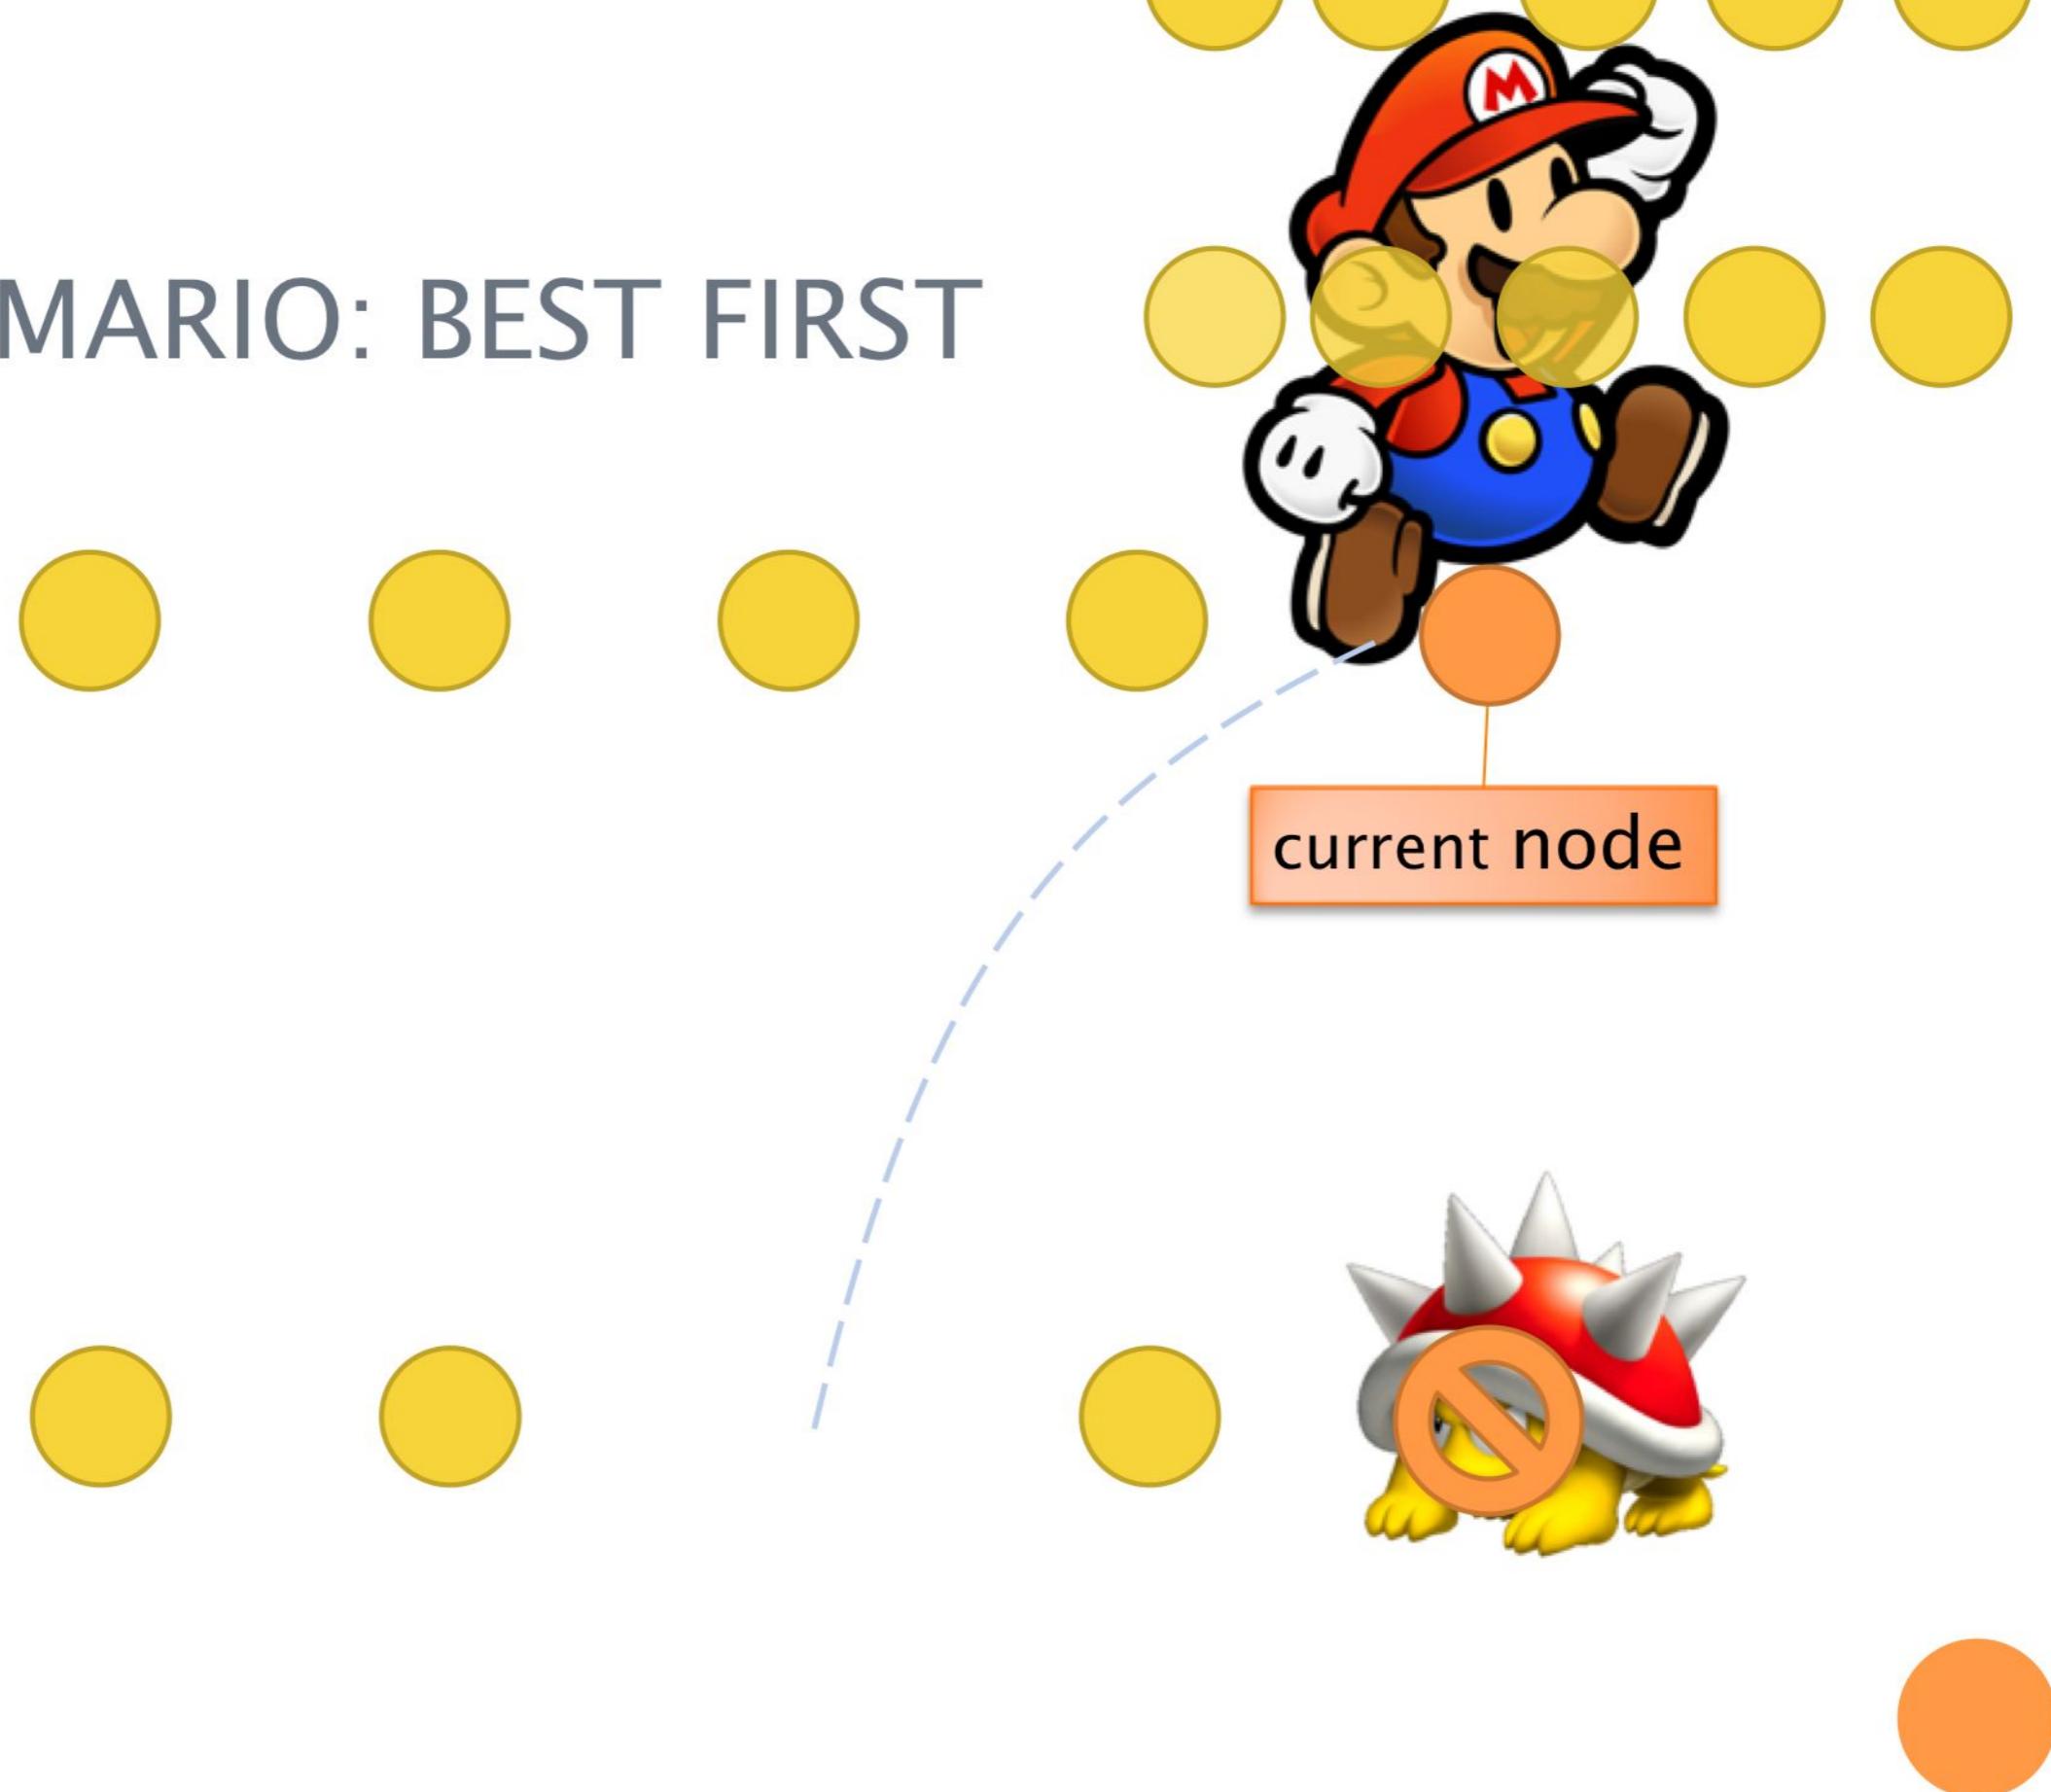
# A\* IN MARIO: EVALUATE



# A\* IN MARIO: CREATE CHILDREN



# A\* IN MARIO: BEST FIRST



# Running into a wall



Tree search versus  
optimization

# Local search algorithms

- In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use local search algorithms: keep a single "current" state, try to improve it

# n-queens

- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal

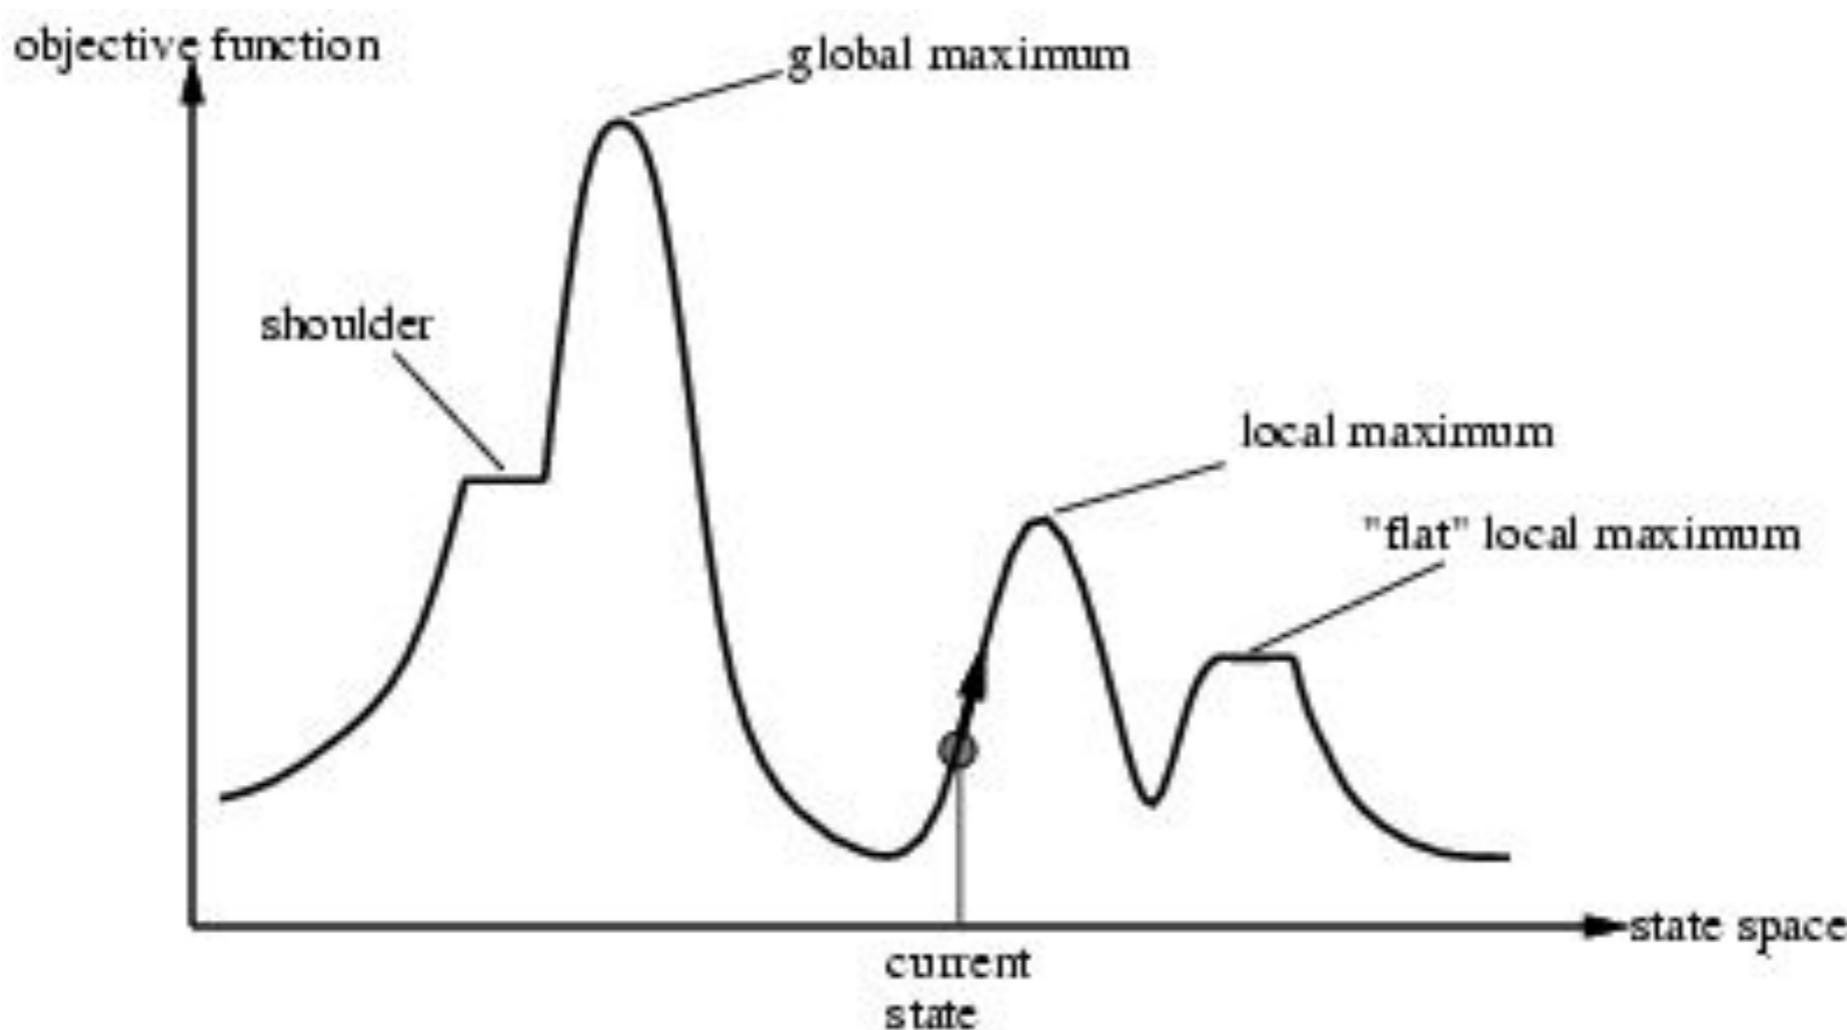


# Hill-climbing

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node
  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor  $\leftarrow$  a (highest-valued) successor of current
    if VALUE[neighbor]  $\leq$  VALUE[current] then return STATE[current]
    current  $\leftarrow$  neighbor
```

# Hill-climbing

- Can get stuck in local maxima/minima

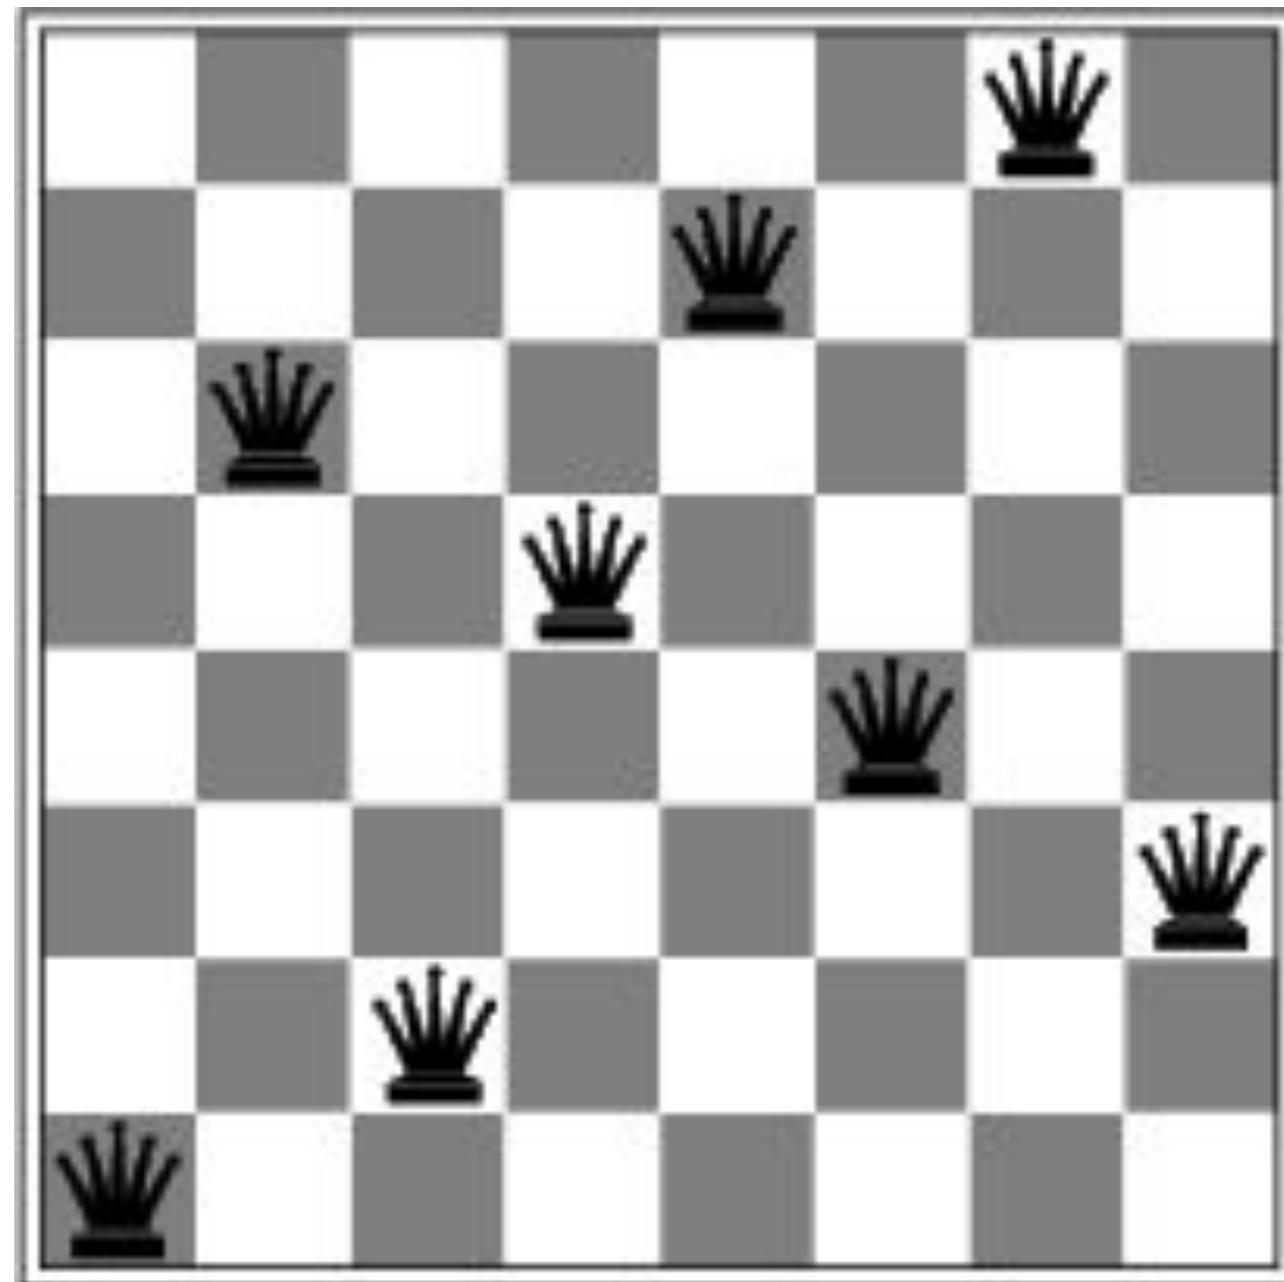


# Hill-climbing 8-queens

- $h$  = number of pairs of queens that are attacking each other, either directly or indirectly. Here,  $h=17$ .

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	14	13	16	13	16
14	14	17	15	14	16	16	16
17	14	16	18	15	14	15	16
18	14	14	15	15	14	14	16
14	14	13	17	12	14	12	18

# Local minimum



# Hill-climbing variations

- Simple hill climber: generates one successor, accepts it if better than the current solution
- “Permissive” hill climber: accepts the successor, accepts it if equal to or better than the current solution
- Steepest ascent hill climber: generates all successors, chooses the best one

# In-class exercise

# Optimization in general

- Optimal: no
- Complete: no
- Space: reasonable
- Time: who knows

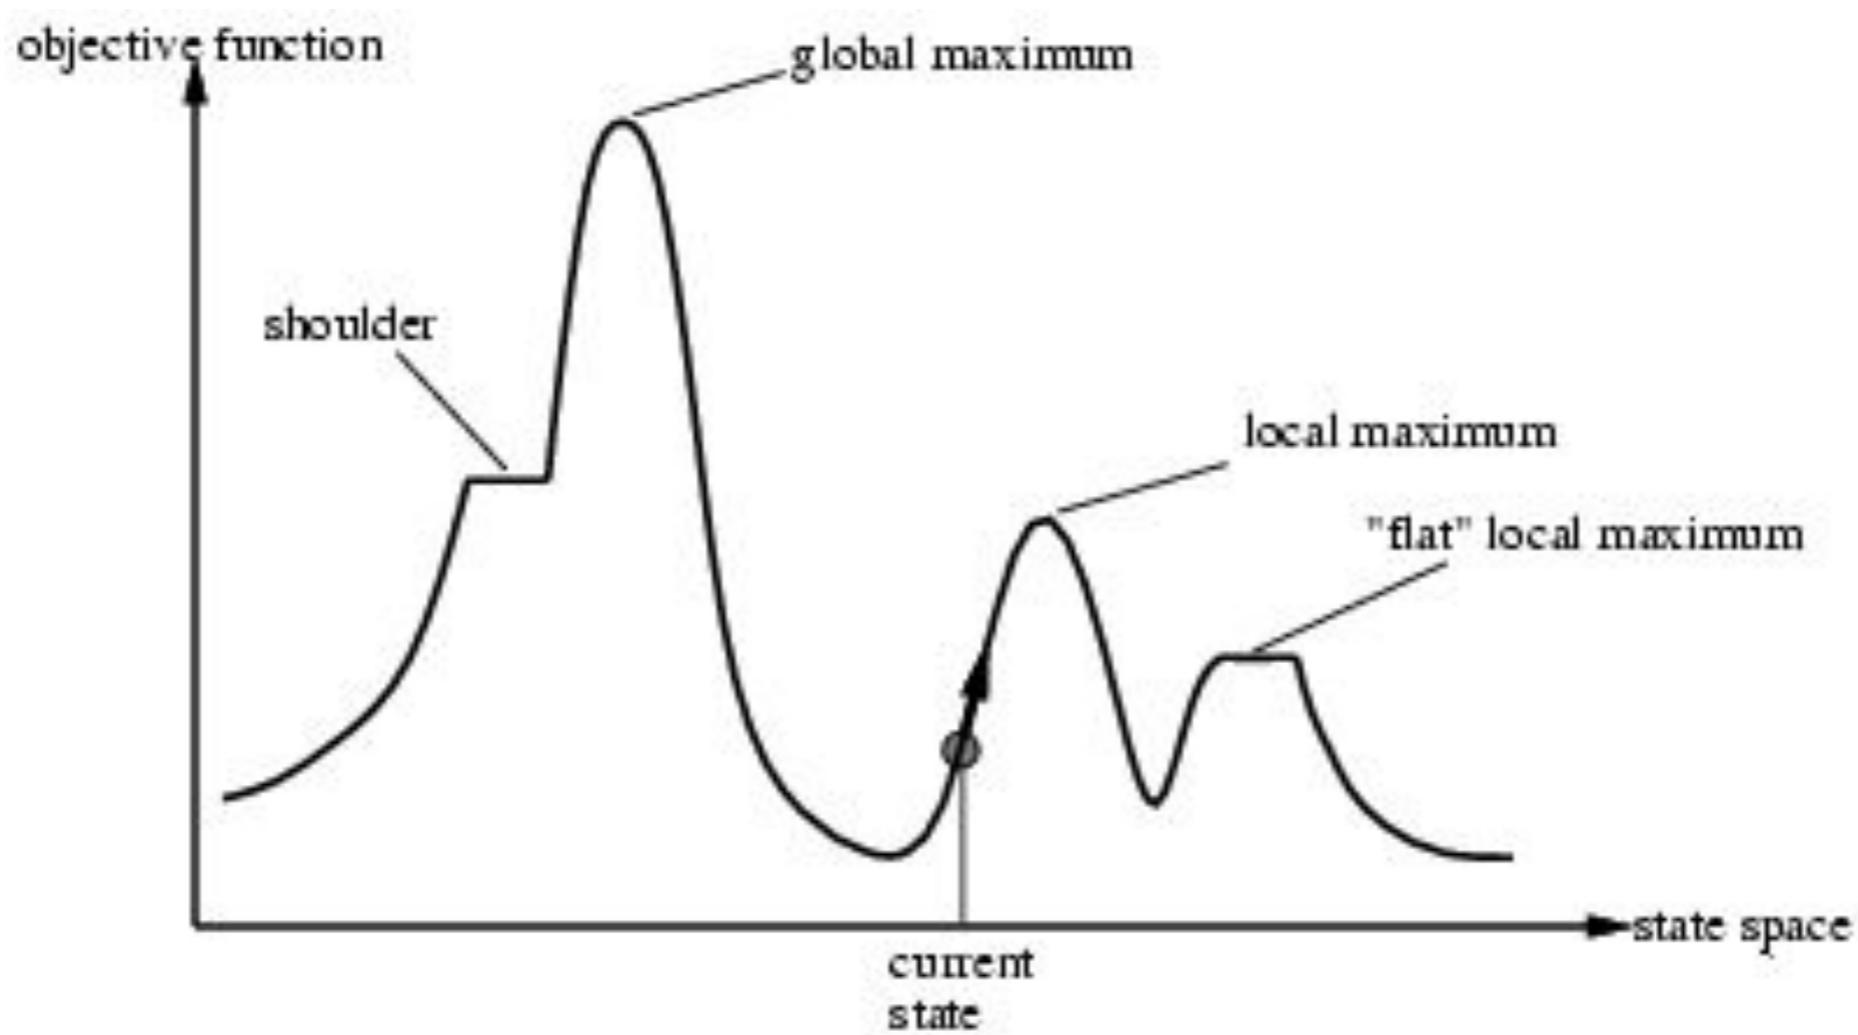
# Other ideas...

- Exhaustive search: try all configurations, one after another
- Random search: generate random configurations

# Random restarts

- When you stop making progress, start another hillclimber somewhere else
- Keep the best solution you found so far

# Hill-climbing



# Simulated annealing

- Do bad moves with decreasing probability

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
          schedule, a mapping from time to “temperature”
  local variables: current, a node
                    next, a node
                    T, a “temperature” controlling prob. of downward steps

  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  for t  $\leftarrow$  1 to  $\infty$  do
    T  $\leftarrow$  schedule[t]
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE[next] - VALUE[current]
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

# Simulated annealing

- One can prove: If  $T$  decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1

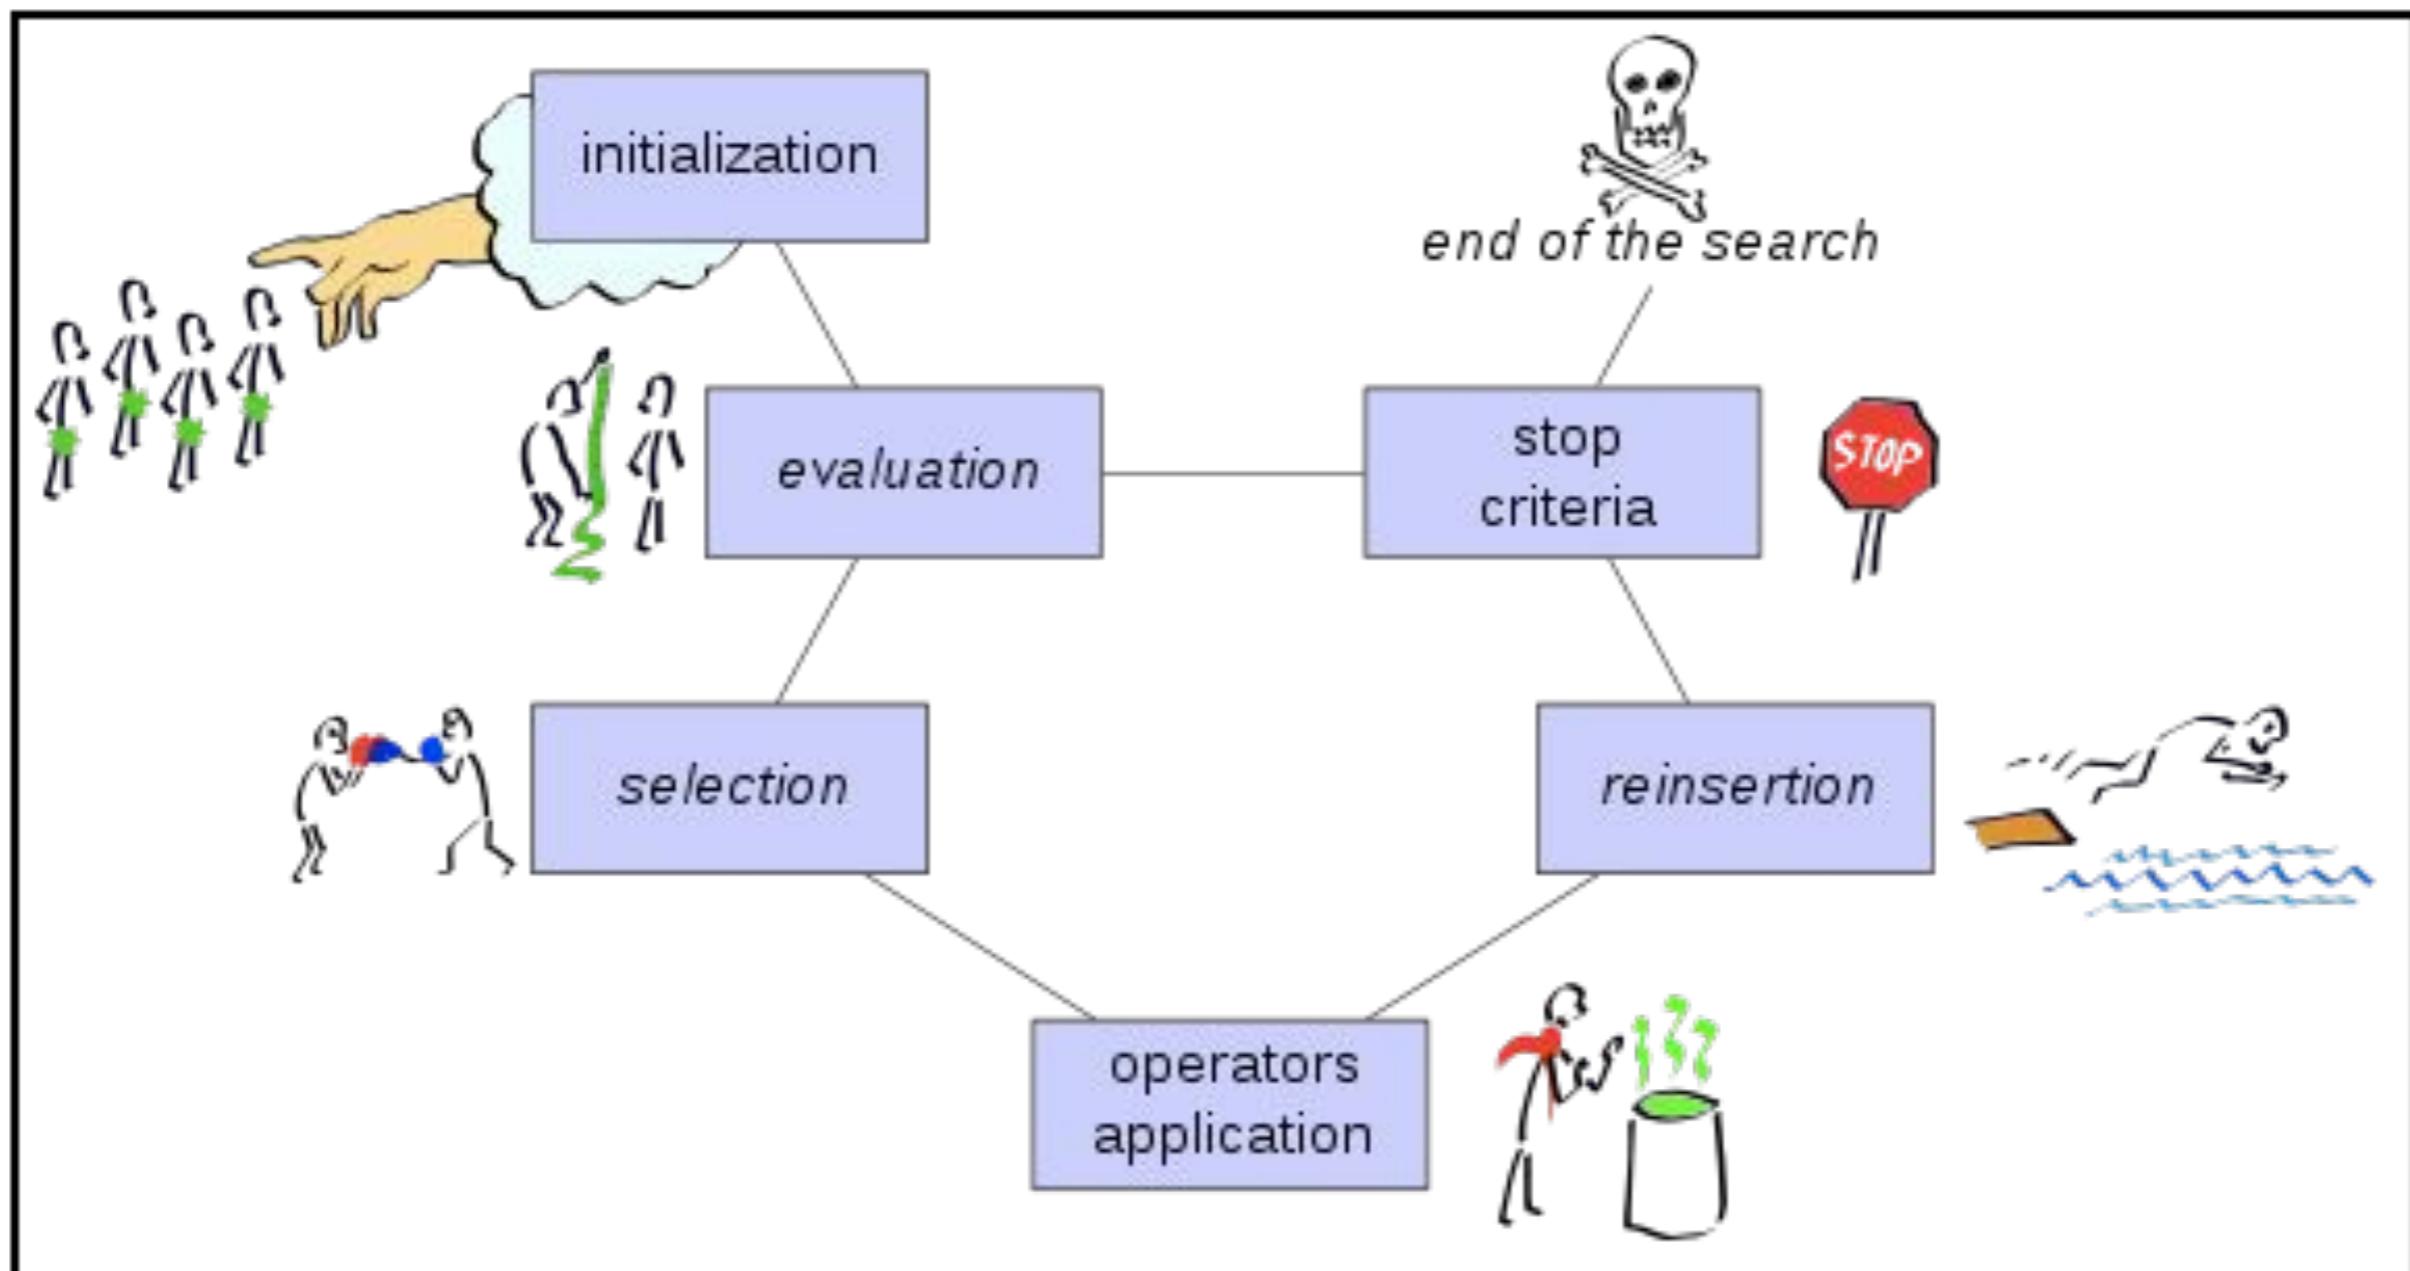
# Limitations of local search

- Susceptible to local optima
- No sharing of information between parallel hill-climbing runs
- No possibility to learn from previous search experience

# Evolutionary algorithms

- Stochastic global optimisation algorithms
- Inspired by Darwinian natural evolution
- Extremely domain-general, widely used in practice

# Evolutionary computation



*General schema of an Evolutionary Algorithm (EA)*