

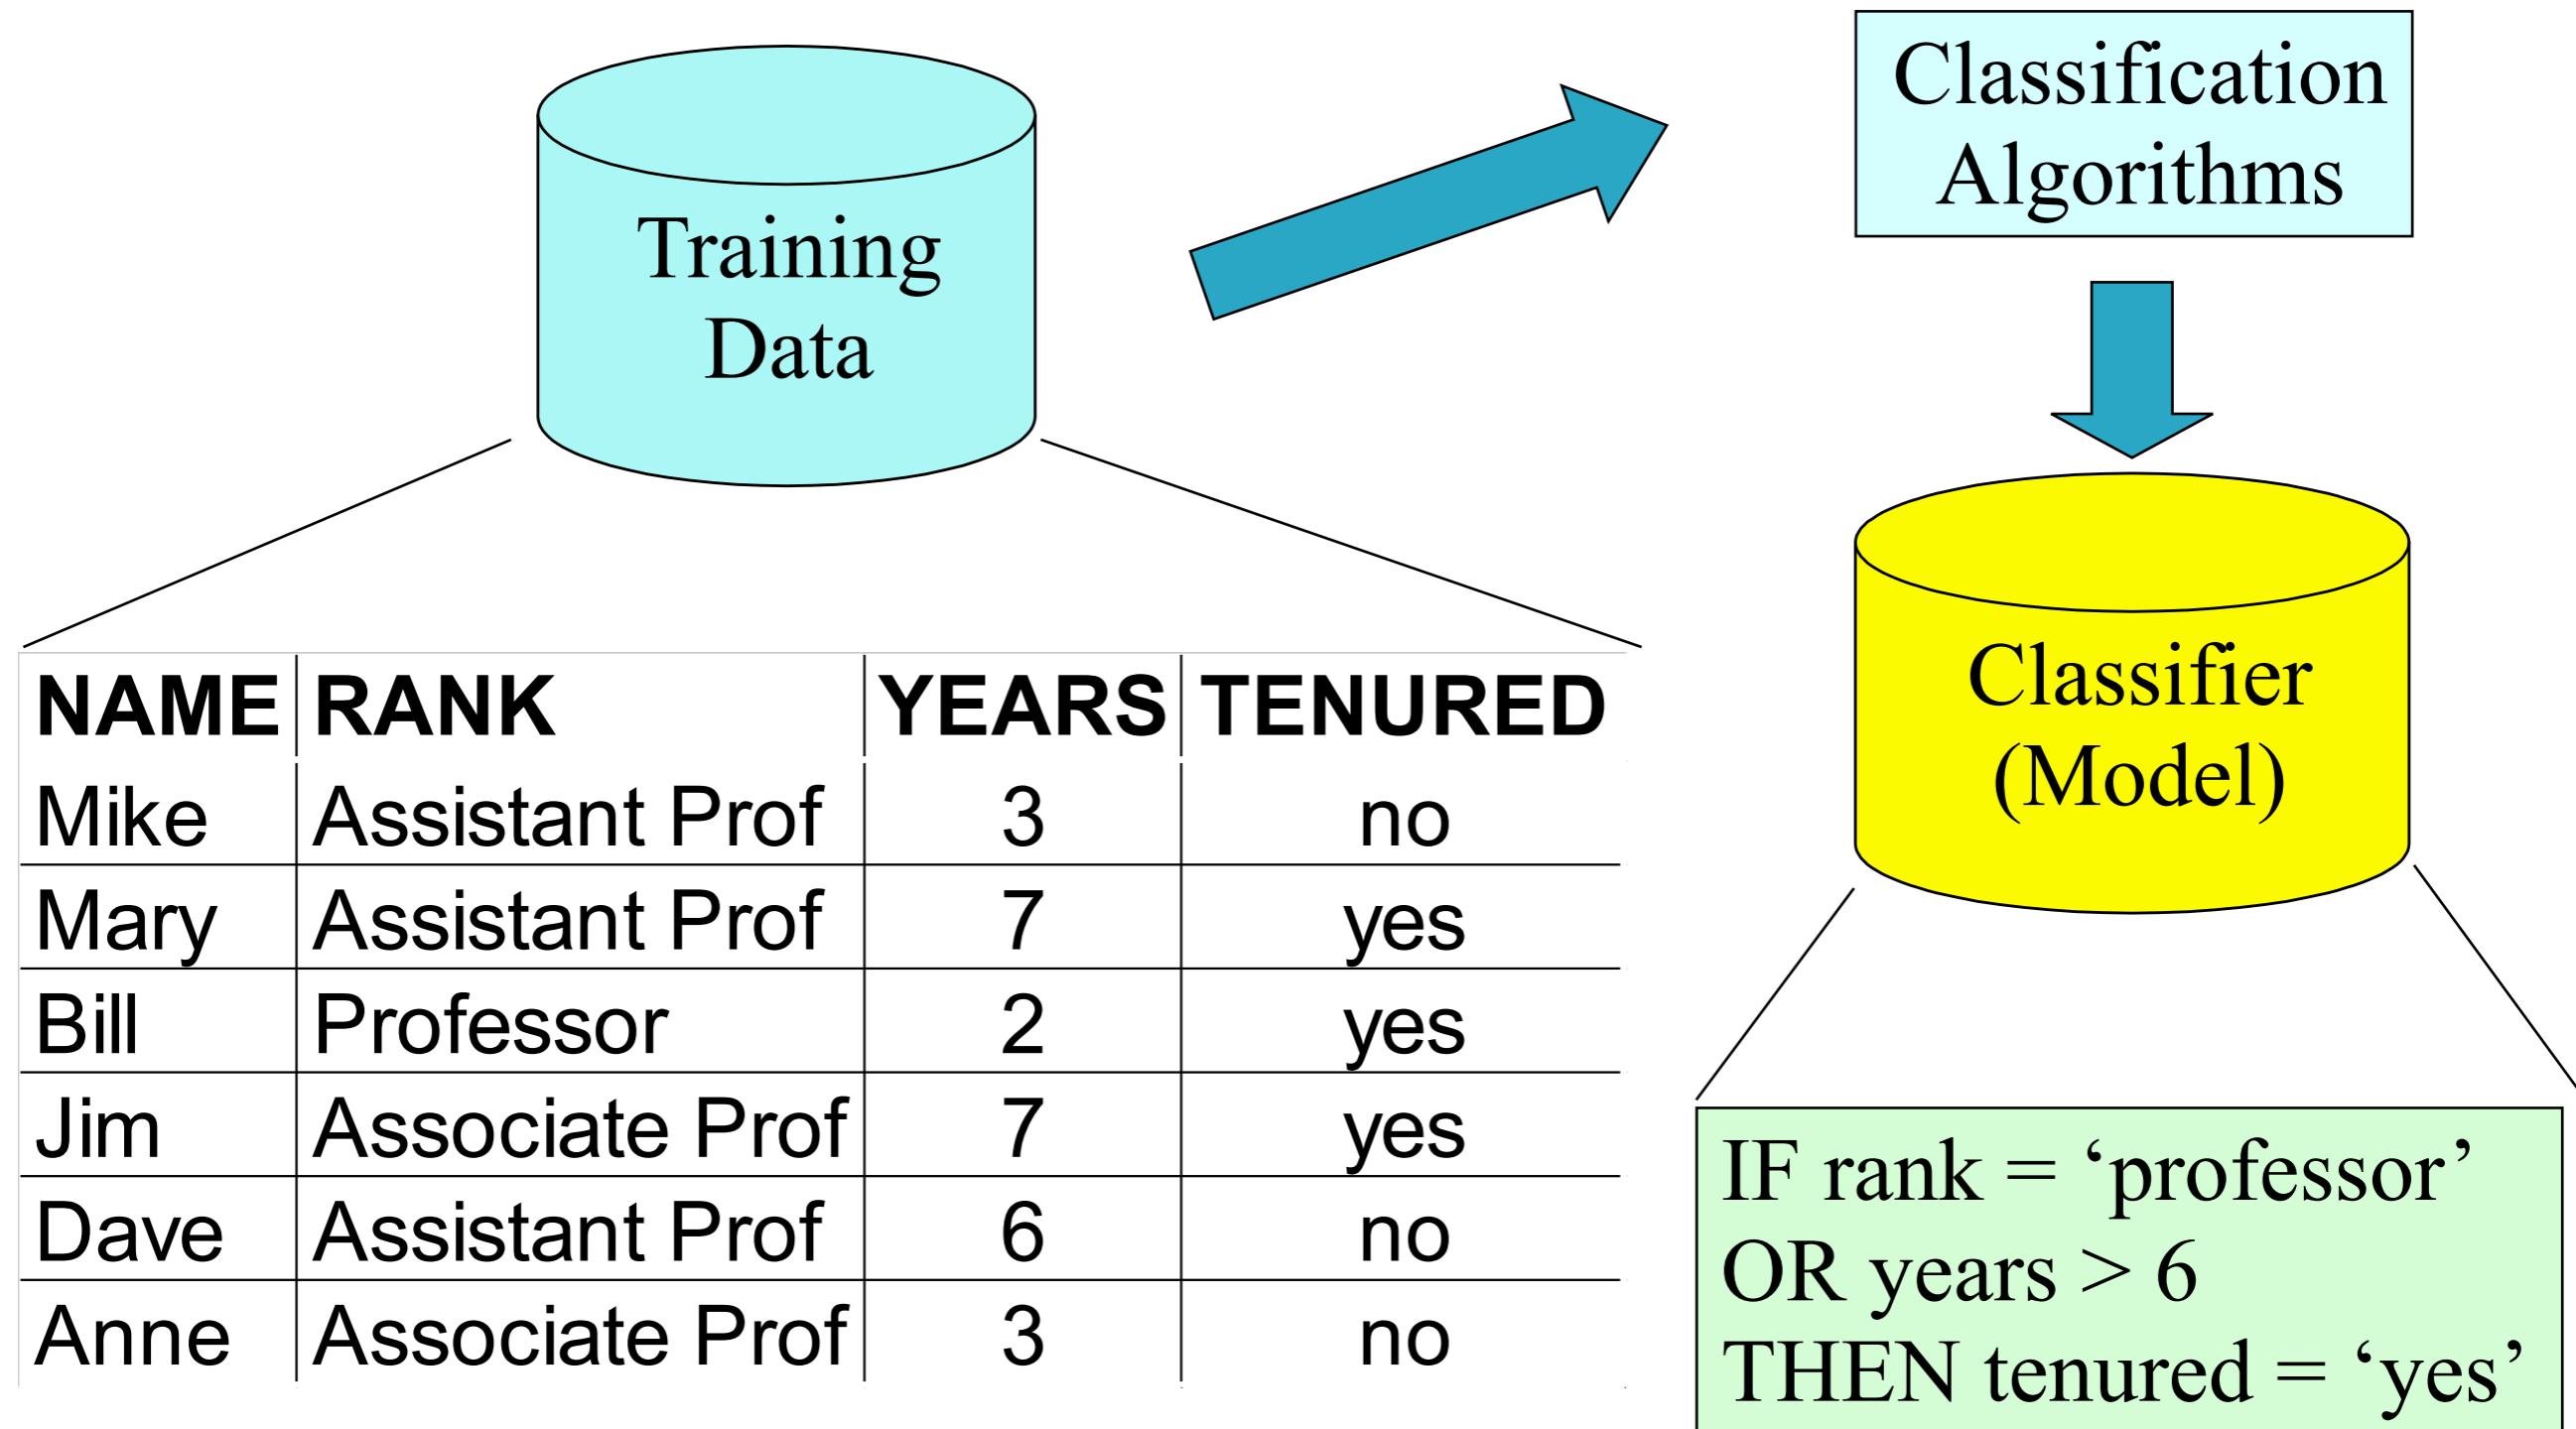
Lecture 8: Decision Trees

Artificial Intelligence
CS-GY-6613-I
Julian Togelius
julian.togelius@nyu.edu

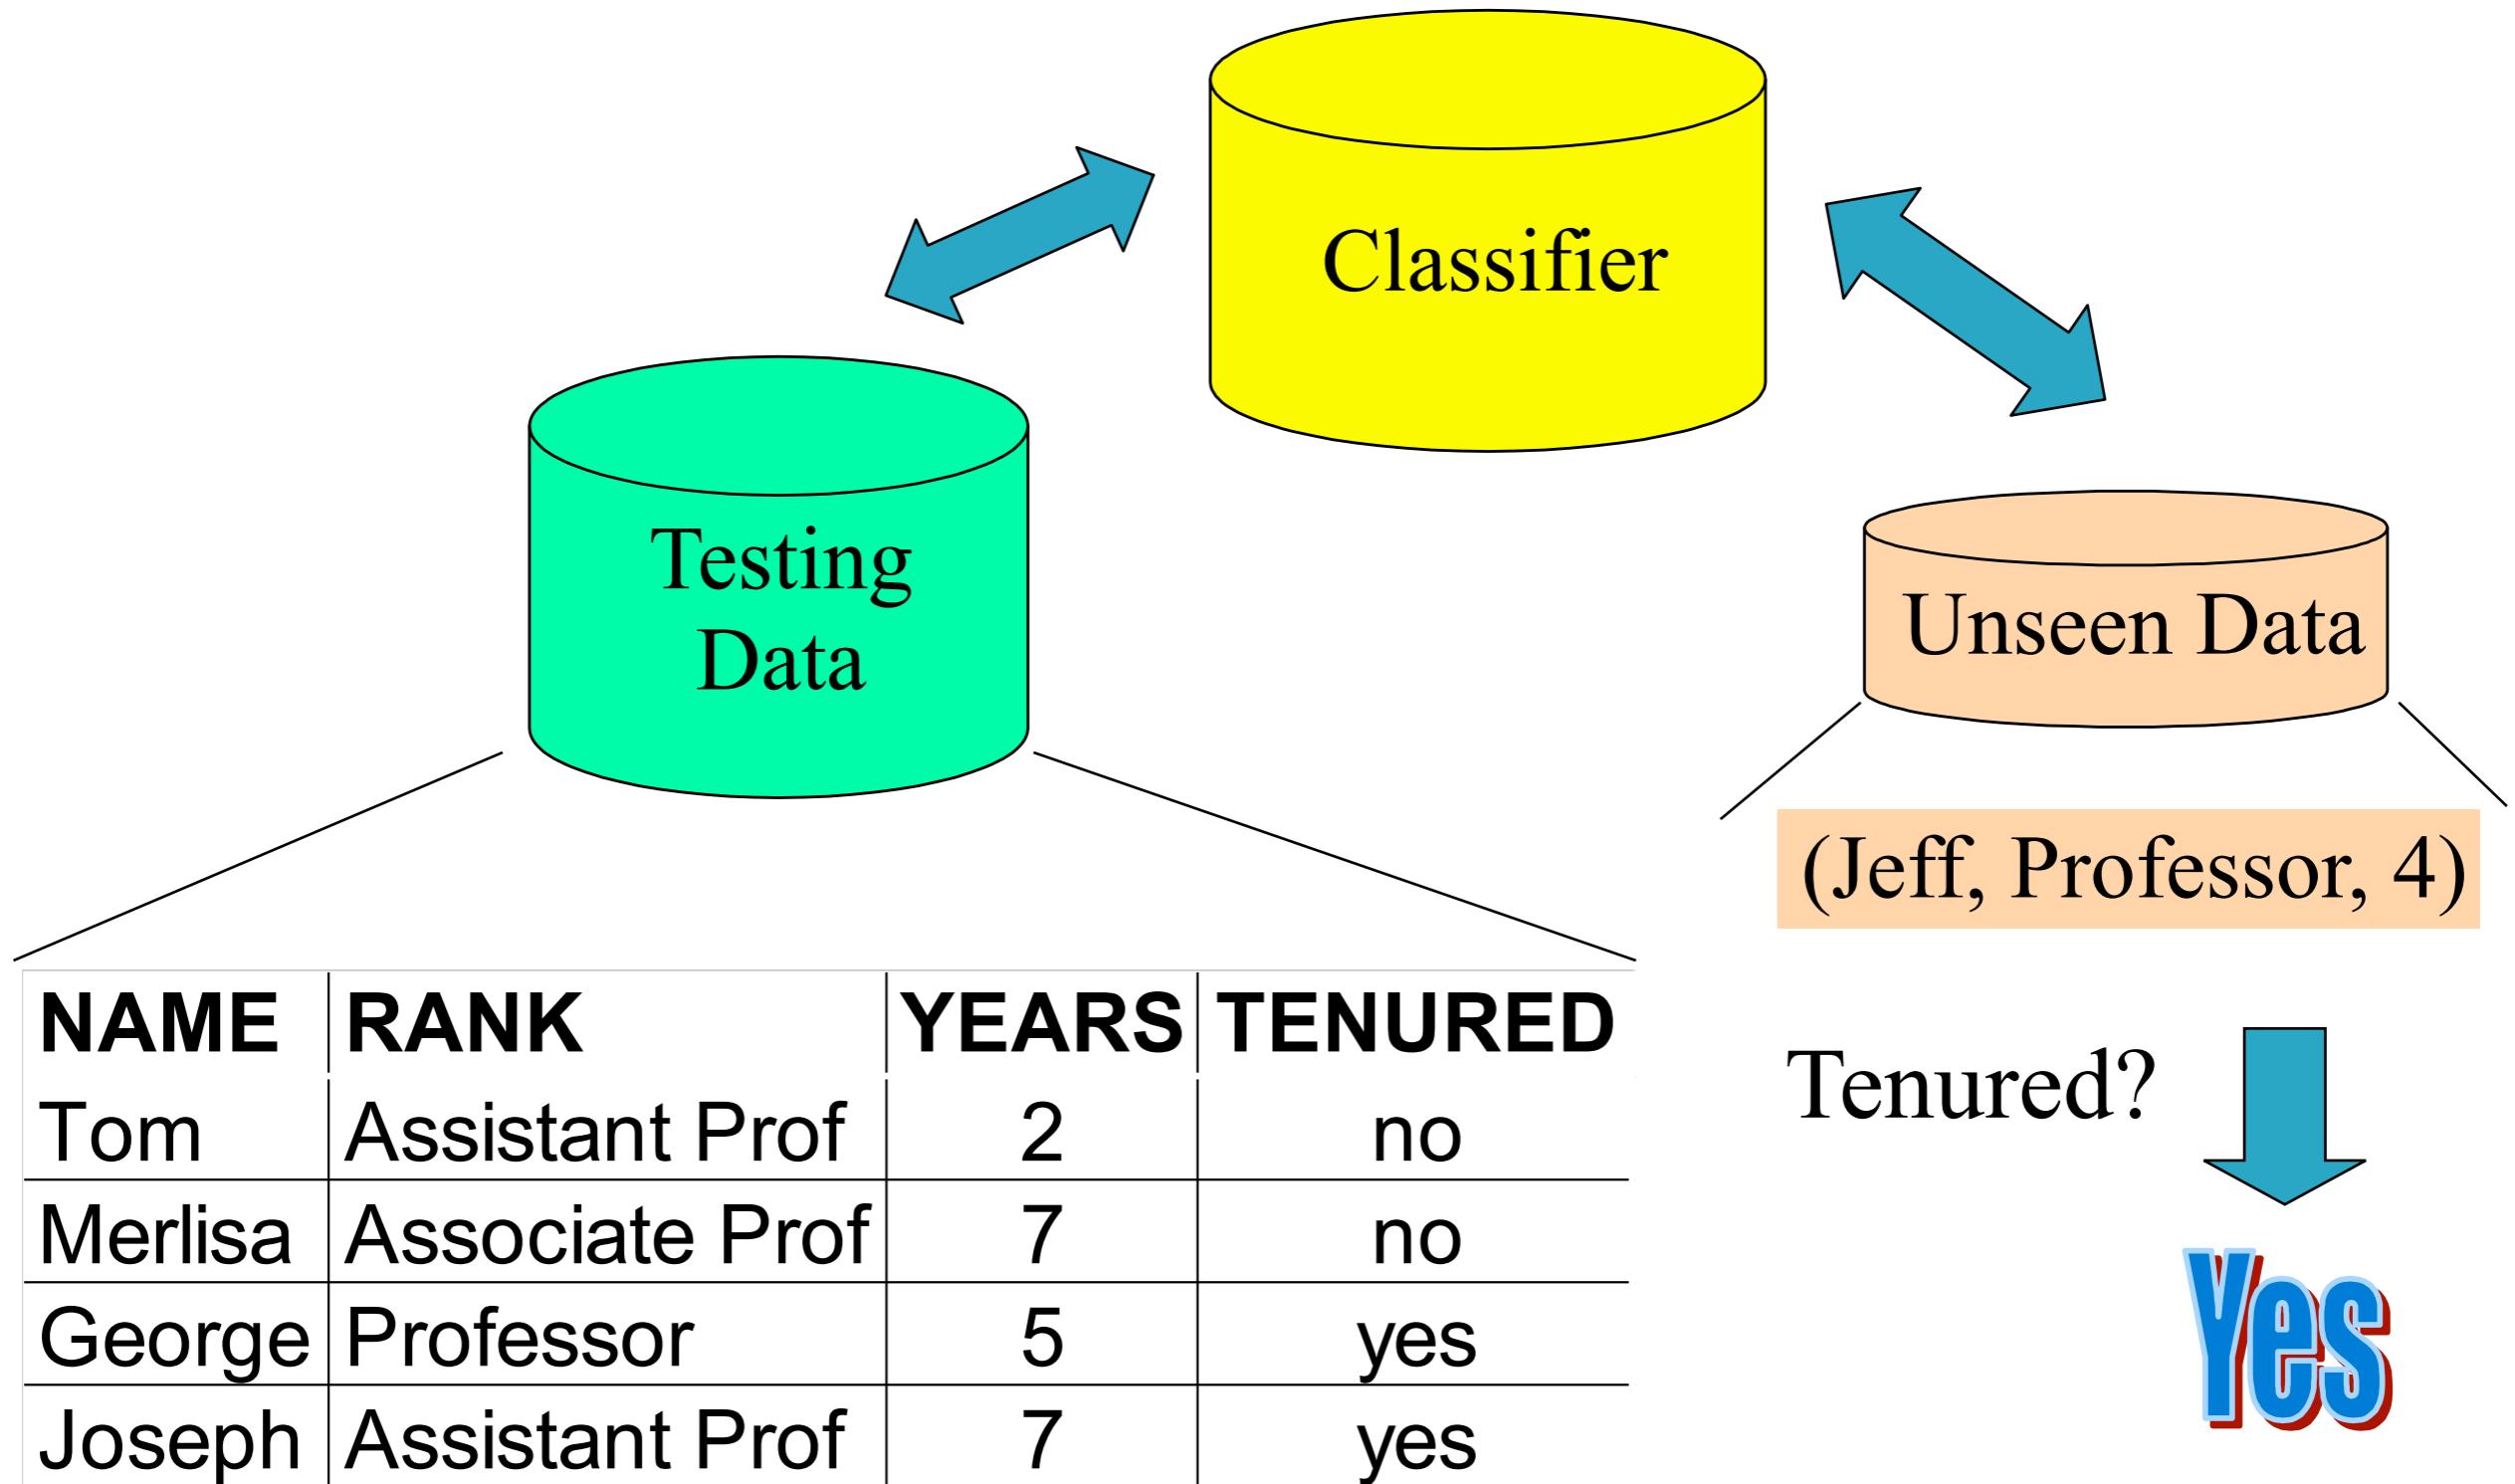
Types of learning

- **Supervised learning**
Learning to predict or classify labels based on labeled input data
- **Unsupervised learning**
Finding patterns in unlabeled data
- **Reinforcement learning**
Learning well-performing behavior from state observations and rewards

Model construction



Using the model



Classification vs prediction

- Classification: binary or nominal labels
 - Examples: pregnant or not, from which country, which type of road sign
- Prediction: continuous labels
 - Examples: future stock price, life expectancy, distance to obstacle

Terminology (supervised learning)

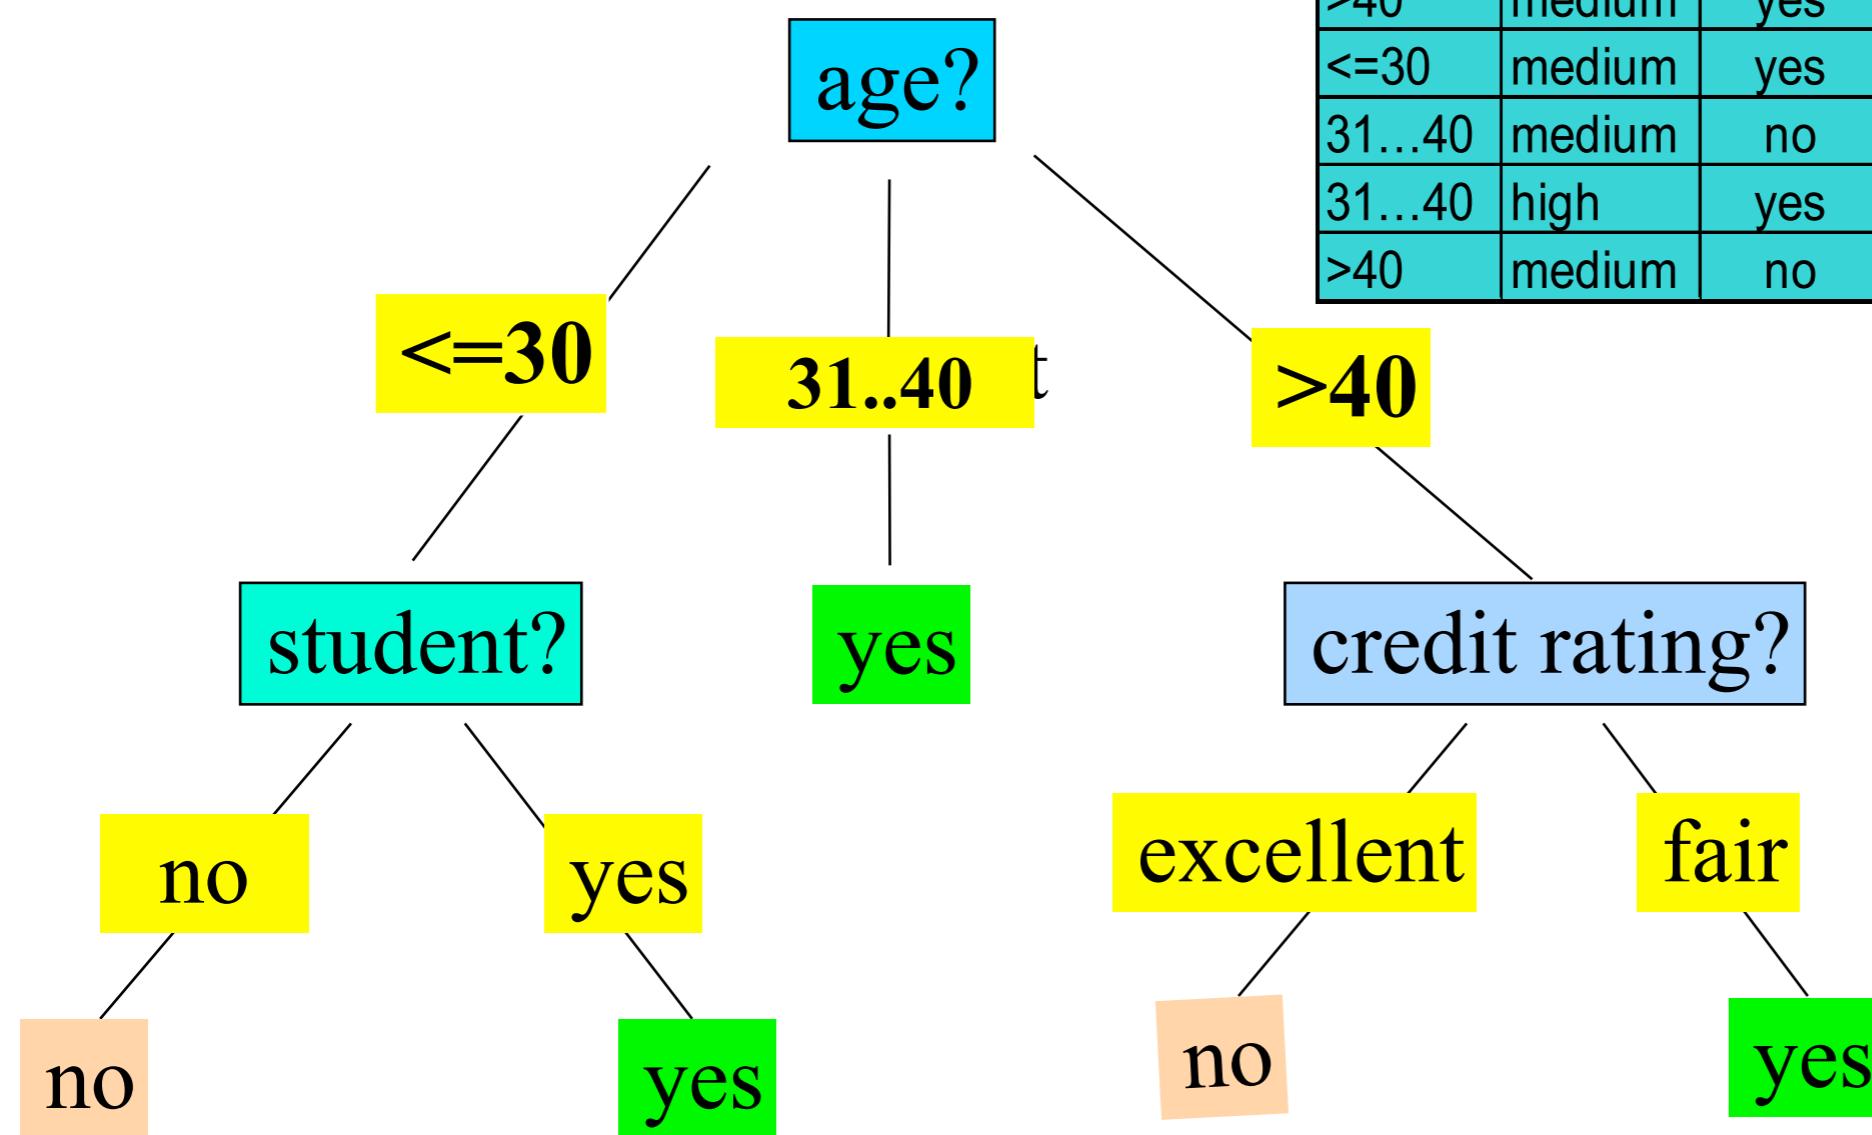
- Each line of data: instance / data point / tuple
- The features of each instance: features / attributes
- That which should be learned: labels / targets
- Each instance has features and a label
- We train on the training set...
- ...and test on the testing set

Decision trees

- A popular representation for classifiers
- Human-readable
- Can be learned (induced) efficiently using algorithms based on information theory
 - An eager learning method
- Often yields high-accuracy classifiers

An example

Classify: buys_computer



age	income	student	credit_rating	buys_computer
≤ 30	high	no	fair	no
≤ 30	high	no	excellent	no
31..40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31..40	low	yes	excellent	yes
≤ 30	medium	no	fair	no
≤ 30	low	yes	fair	yes
>40	medium	yes	fair	yes
≤ 30	medium	yes	excellent	yes
31..40	medium	no	excellent	yes
31..40	high	yes	fair	yes
>40	medium	no	excellent	no

ID3 Algorithm for Decision Tree Induction

- Tree is constructed in a top-down recursive divide-and-conquer manner
- At start, all the training examples are at the root
- Attributes are categorical (if continuous-valued, they are discretized in advance)
- Examples are partitioned recursively based on selected attributes
- Test attributes are selected on the basis of a heuristic or statistical measure (e.g., information gain)

ID3 Algorithm for Decision Tree Induction

- Stop partitioning when...
 - All samples for a given node belong to the same class, or...
 - There are no remaining attributes for further partitioning – (vote on the leaf) or...
 - There are no samples left

Attribute Selection Measure: Information Gain (ID3/C4.5)

- Select the attribute with the highest information gain
- Let p_i be the probability that an arbitrary tuple in D belongs to class C_i , estimated by $|C_i, D|/|D|$
- Expected information (entropy) needed to classify a tuple in D :

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$

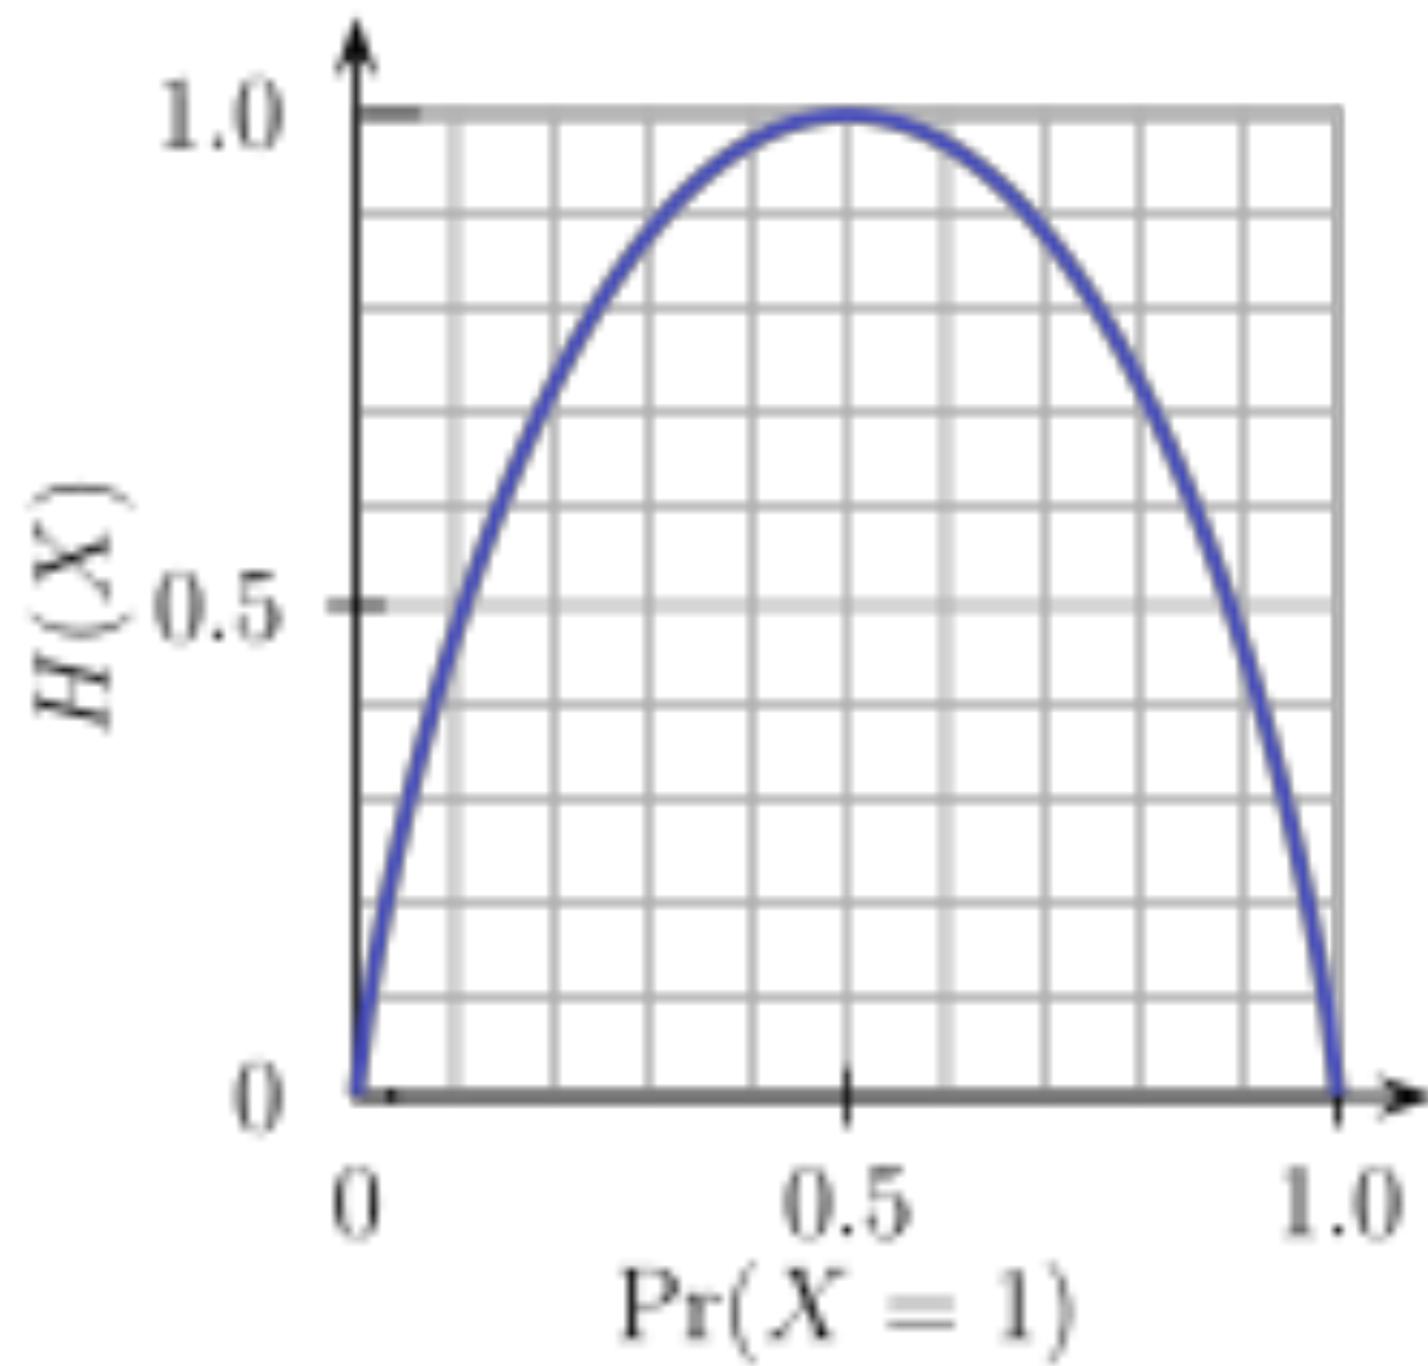
- Information needed (after using A to split D into v partitions) to classify D :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

- Information gained by branching on attribute A

$$Gain(A) = Info(D) - Info_A(D)$$

Binary entropy function



Attribute Selection: Information Gain

Class P: buys_computer = “yes”

Class N: buys_computer = “no”

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0)$$

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2(\frac{9}{14}) - \frac{5}{14} \log_2(\frac{5}{14}) = 0.940$$

$$+ \frac{5}{14} I(3,2) = 0.694$$

age	p_i	n_i	$I(p_i, n_i)$
≤ 30	2	3	0.971
31...40	4	0	0
> 40	3	2	0.971

$\frac{5}{14} I(2,3)$ means “age ≤ 30 ” has 5 out of 14 samples, with 2 yes’es and 3 no’s. Hence

age	income	student	credit_rating	buys_computer
≤ 30	high	no	fair	no
≤ 30	high	no	excellent	no
31...40	high	no	fair	yes
> 40	medium	no	fair	yes
> 40	low	yes	fair	yes
> 40	low	yes	excellent	no
31...40	low	yes	excellent	yes
≤ 30	medium	no	fair	no
≤ 30	low	yes	fair	yes
> 40	medium	yes	fair	yes
≤ 30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
> 40	medium	no	excellent	no

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly,

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit_rating) = 0.048$$

- Input: A data set, D
Output: A decision tree
 - If all the instances have the same value for the target attribute then return a decision tree that is simply this value (not really a tree - more of a stump).
- Else
 - Compute Gain values (see above) for all attributes and select an attribute with the highest value and create a node for that attribute.
 - Make a branch from this node for every value of the attribute
 - Assign all possible values of the attribute to branches.
 - Follow each branch by partitioning the dataset to be only instances whereby the value of the branch is present and then go back to 1.

function DECISION-TREE-LEARNING(*examples*, *attributes*, *parent-examples*) **returns**
tree

if *examples* is empty **then return** PLURALITY-VALUE(*parent-examples*)
else if all *examples* have the same classification **then return** the classification
else if *attributes* is empty **then return** PLURALITY-VALUE(*examples*)
else

$A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \operatorname{IMPORTANCE}(a, \text{examples})$

tree \leftarrow a new decision tree with root test *A*

for each value v_k of *A* **do**

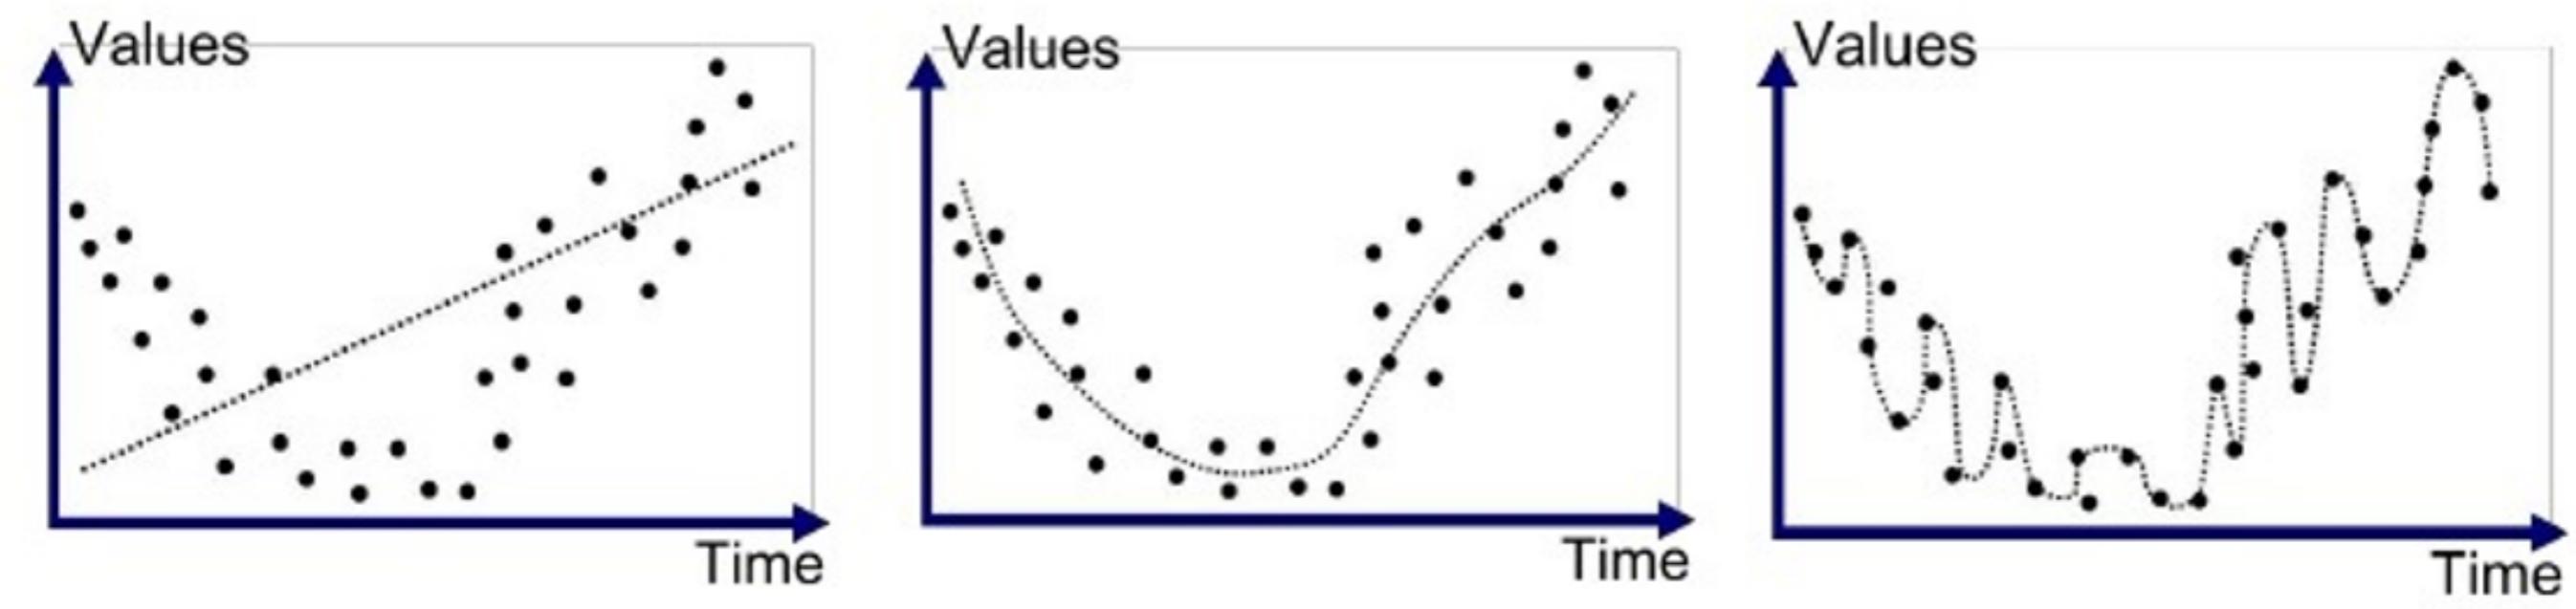
$exs \leftarrow \{e : e \in \text{examples} \text{ and } e.A = v_k\}$

subtree \leftarrow DECISION-TREE-LEARNING(*exs*, *attributes* – *A*, *examples*)

add a branch to *tree* with label (*A* = v_k) and subtree *subtree*

return *tree*

Overfitting (in supervised learning)



Underfitted

Good Fit/Robust

Overfitted

- Fitting the data too well, including the noise
- Reduces accuracy on unseen data

Overfitting and pruning

- Overfitting: An induced tree may overfit the training data (specific instance of a concept that applies to all supervised learning)
 - Too many branches, some may reflect anomalies due to noise or outliers
 - Poor accuracy for unseen samples
- Three approaches to avoid overfitting
 - Prepruning: Halt tree construction early—do not split a node if this would result in the goodness measure (e.g. accuracy on a validation set) falling below a threshold
 - Postpruning: Remove branches from a “fully grown” tree—get a sequence of progressively pruned trees
 - Use a set of data different from the training data to decide which is the “best pruned tree”
 - Limit depth or size

What to do with numerical attributes?

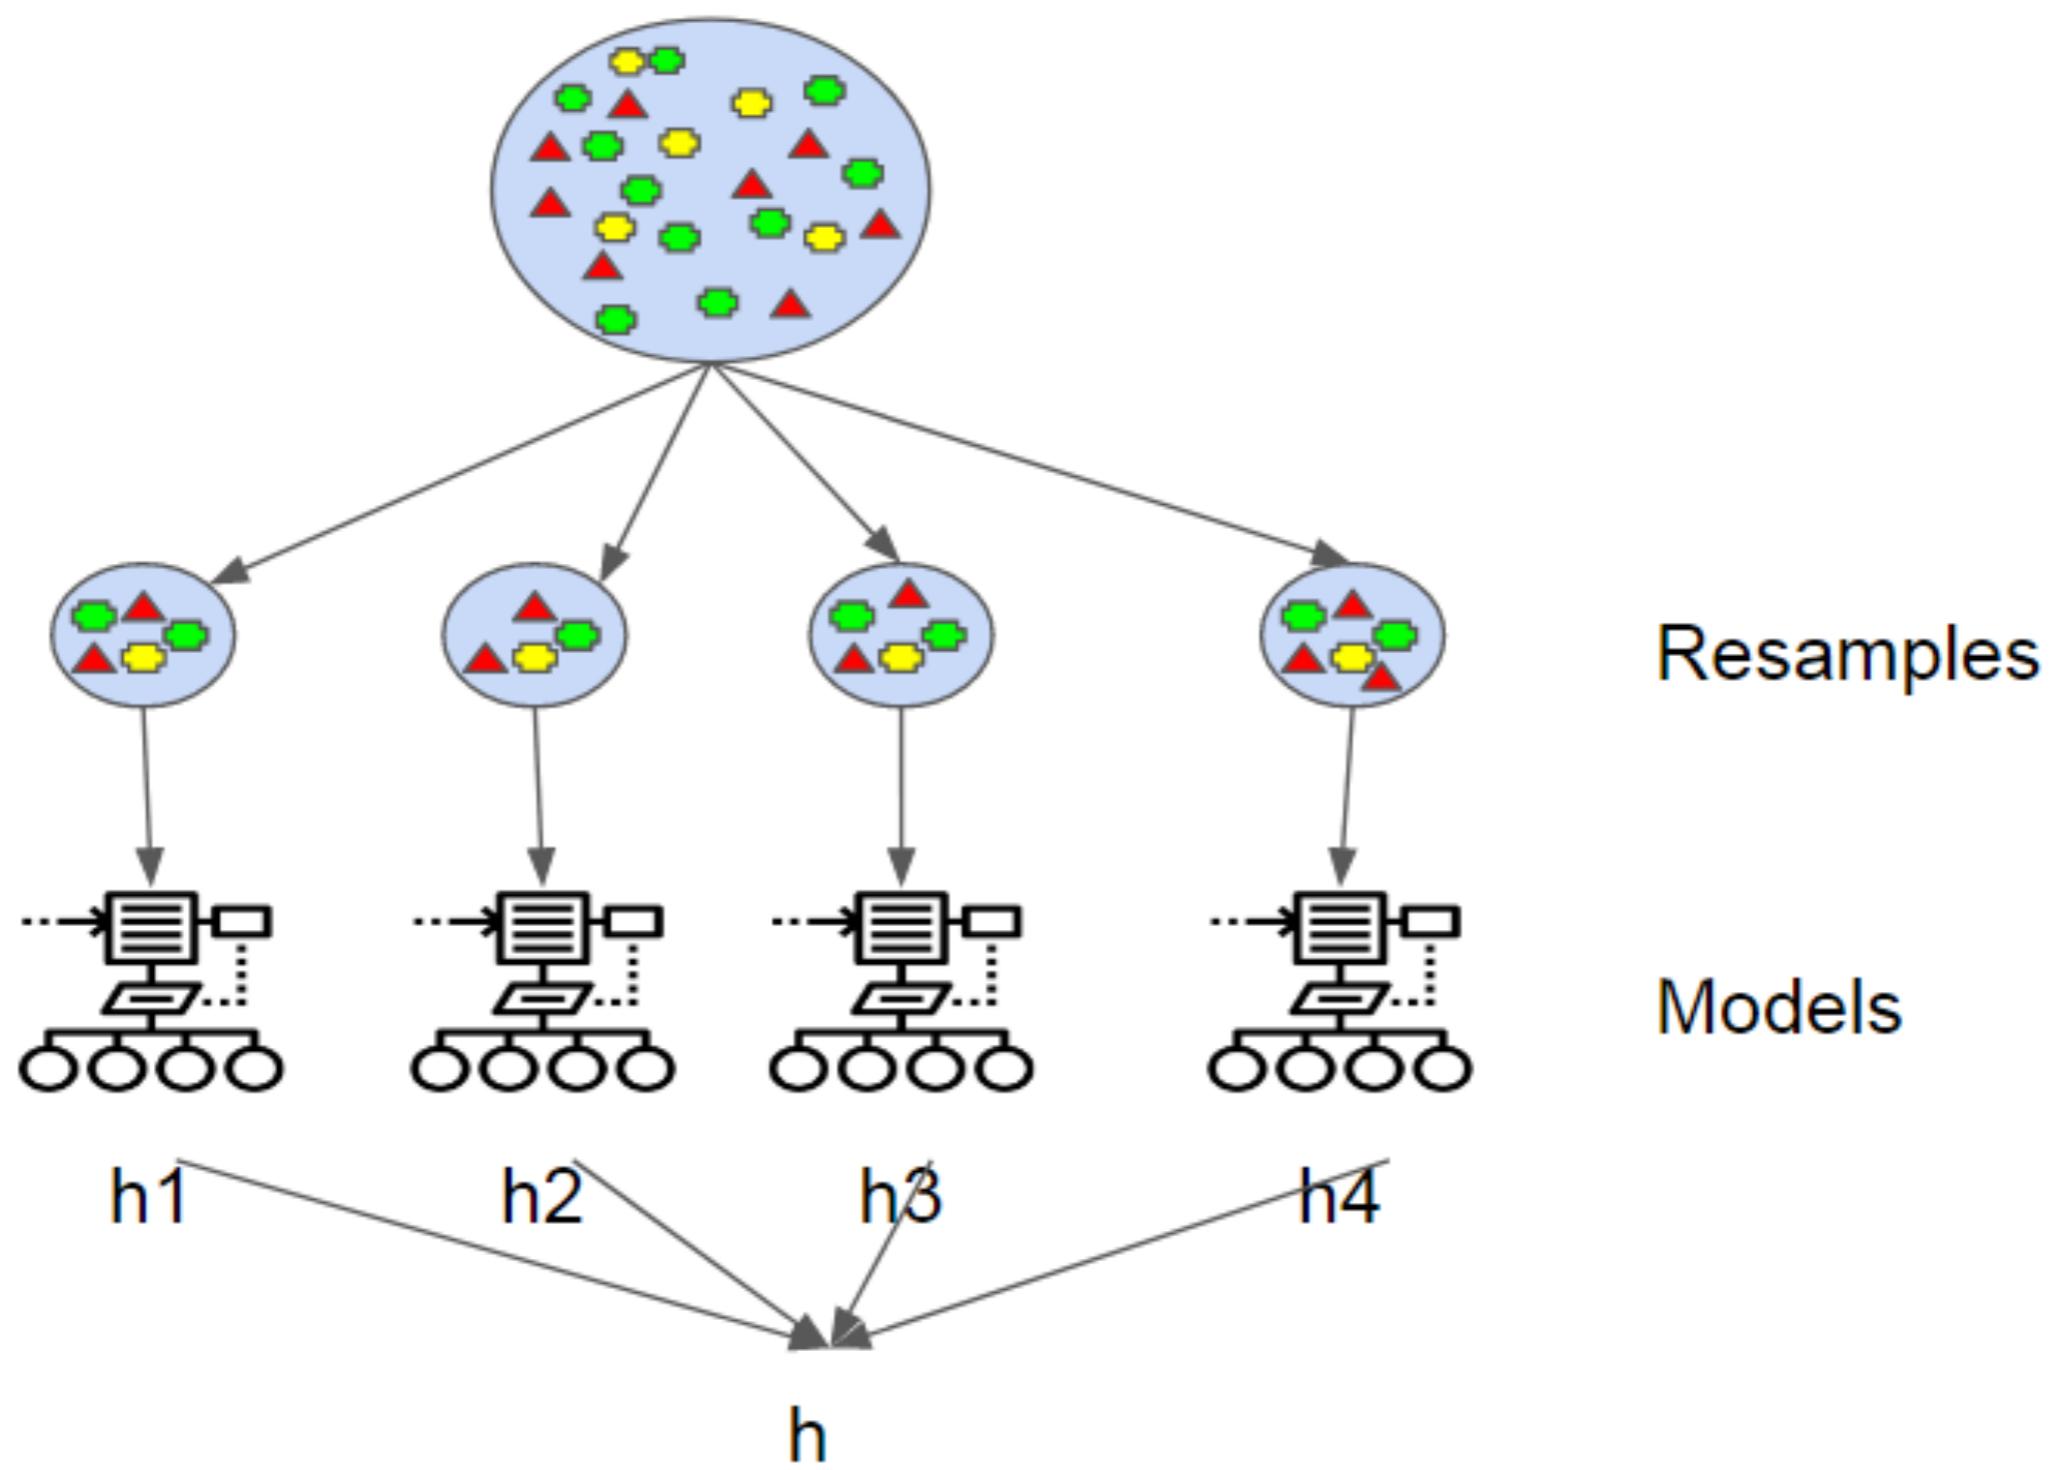
- Create a set of bins, where each bin becomes one attribute value (e.g. income <30k, 30-40k, >40k)
 - Equal depth: All the bins have the same number of instances in them
 - Equal width: All the bins have the same range
 - More sophisticated ways, including looking at where there are big discontinuities in the range of the input data etc

Can we do better?

- It's hard to train good classifiers
- Can we replace quality with quantity?
 - Train many classifiers and combine them somehow
 - Called *ensemble learning*

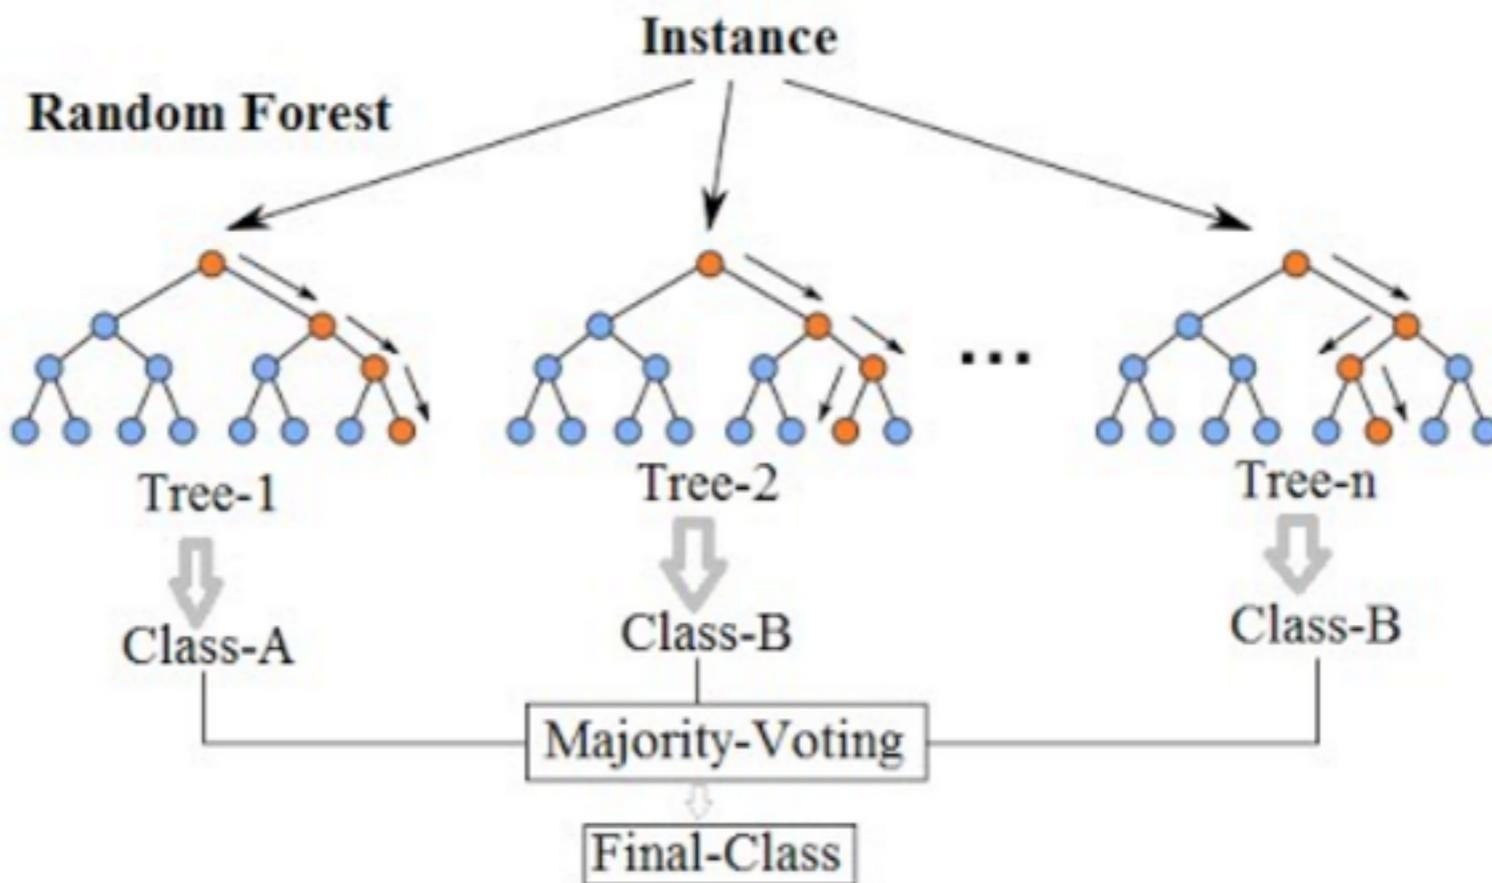
Bagging

- Short for Bootstrap Aggregation
- Train a number of classifiers, each on a randomly selected subset of the data
 - Drawn with replacement (the same instance can be part of many subsets)
- Use the majority vote of the models as the output
- For prediction, a mean or weighted mean



Random forests

Random Forest Simplified



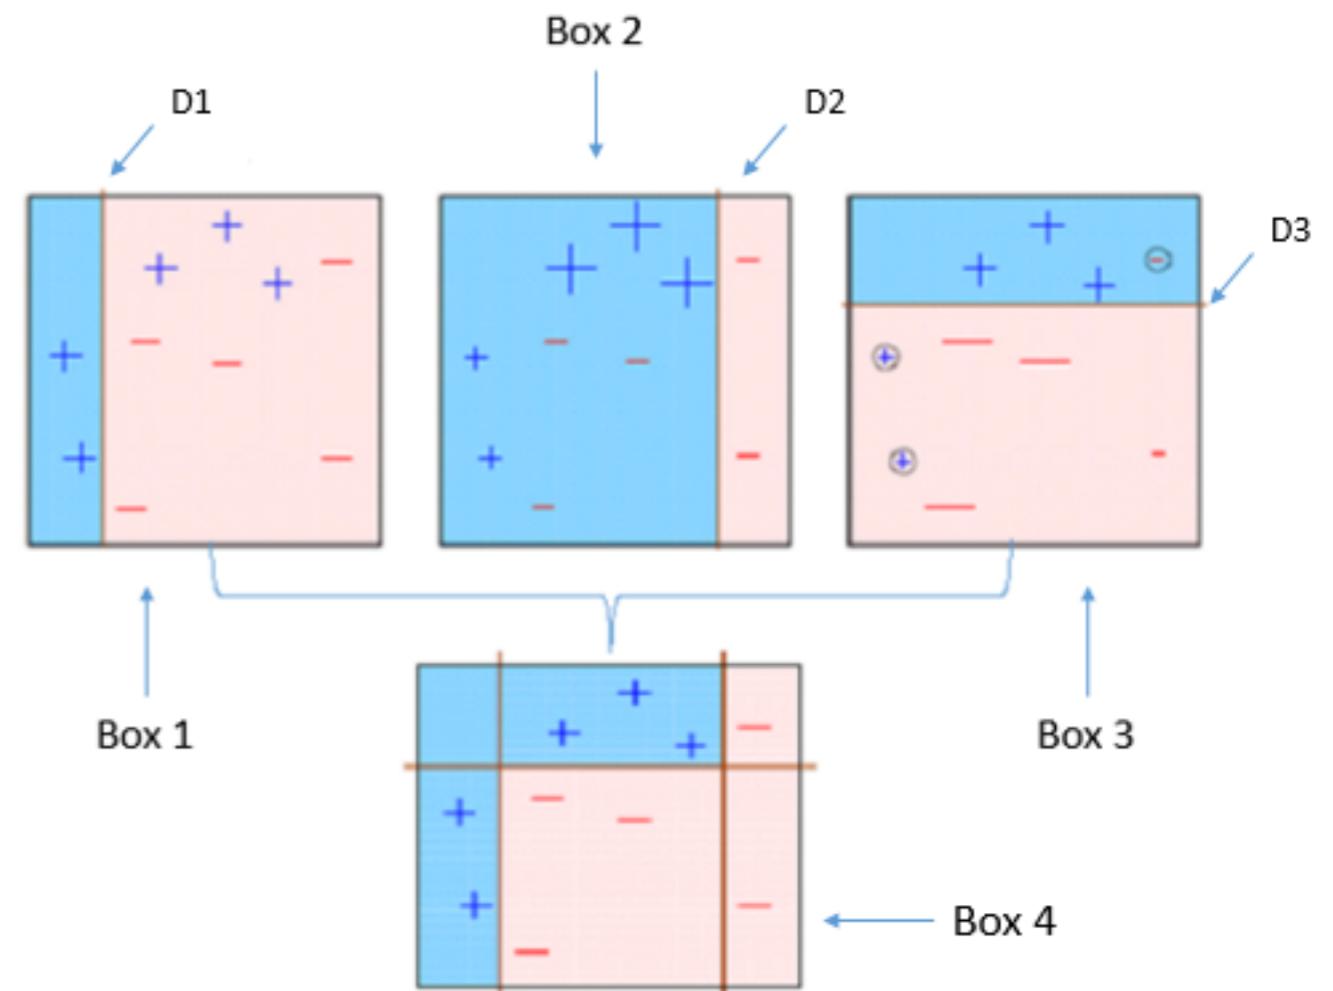
- Bagging for decision trees
 - Also *randomly select features* for each subset

Boosting

- Train many models and return the (weighted) average/majority classification, like bagging
- Unlike bagging, train each model to be good at where the previous models failed

AdaBoost

- First train a classifier on the whole data set
- Then, train again while weighing the misclassified instances higher



Gradient boosting

- See boosting as an optimization problem
- Optimize the parameters of each new classifier for minimum error using gradient descent