

Lecture 8: Multilayer Perceptrons

Artificial Intelligence

CS-GY-6613

~~Julian Togelius~~

julian.togelius@nyu.edu


Philip Bontrager
2019


Contents

- Review Perceptrons
- Multilayer Perceptron Intro
- Computational Graphs
- How to Train a Multilayer Perceptron
- Why Deep Learning now

Perceptron

The Neuron









$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i + b > 0 \\ 0 & \text{else} \end{cases}$$

Perceptron problems

- Noise: if the data isn't separable, weights might thrash
- Mediocre generalization: finds a “barely” separating solution
- Overtraining: test / held-out accuracy usually rises, then falls



Which of these separators is optimal?




With all of these problems and with so many other learning algorithms, why were people still interested in pursuing this technique?

Multilayer Perceptron

“When an axon of cell A is near enough cell B and repeatedly or persistently takes part in firing it, some growth process or metabolic change takes place in one or both cells such that A’s efficiency, as one of the cells firing B, is increased”

-Donald Hebb

Multi-layer Perceptron (MLP)



Universal Approximator

- George Cybenko in 1989 provided one of the first proofs for the Universal Approximation Theorem for Neural Networks
- States that a feedforward network with at least one hidden layer of a finite size can approximate any continuous function on compact subsets of R^n
- Says nothing about learnability
 - How do you train a hidden layer? What is its loss?

Computational Graphs

Section inspired by slides from Fei-Fei Li & Justin Johnson & Serena Yeung

Quick Calculus Refresh

- Partial Derivatives

- $f(x, y, z) = 3xz + y^2$


- $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} = 3z, 2y, 3x$

- Chain Rule

- $(f \circ g)' = (f' \circ g) \cdot g'$


Functions as Computation Graphs

$$f(x, y, z) = 3xz + y^2$$



Functions as Computation Graphs



$$f(x, y, z) = 3xz + y^2$$
$$x = 2, y = 3, z = 1$$



Functions as Computation Graphs

$$f(x, y, z) = 3xz + y^2$$
$$x = 2, y = 3, z = 1$$

Chain Rule:





Functions as Computation Graphs

$$f(x, y, z) = 3xz + y^2$$

$$x = 2, y = 3, z = 1$$

Chain Rule:





Functions as Computation Graphs

$$f(x, y, z) = 3xz + y^2$$

$$x = 2, y = 3, z = 1$$



Chain Rule:



Functions as Computation Graphs

$$f(x, y, z) = 3xz + y^2$$
$$x = 2, y = 3, z = 1$$

Chain Rule:




Functions as Computation Graphs

$$f(x, y, z) = 3xz + y^2$$

$$x = 2, y = 3, z = 1$$


Chain Rule:



$$(f \circ g)' = (f' \circ g) \cdot g'$$

**Encapsulating
Function
Gradient**

**Node
Gradient**




Functions as Computation Graphs

$$f(x, y, z) = 3xz + y^2$$

$$x = 2, y = 3, z = 1$$

Chain Rule:



$$(f \circ g)' = (f' \circ g) \cdot g'$$


Encapsulating Function Gradient

Node Gradient

$$\frac{\partial f}{\partial x} = 3$$


$$\frac{\partial f}{\partial z} = 6$$

$$\frac{\partial f}{\partial y} = 6$$



Perceptron Example


$$f(x, y) = \sigma(w_1x + w_2y + b)$$




Perceptron Example

Sigmoid:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



Perceptron Example



Perceptron Example

Sigmoid Derivative:


$$\sigma(x)' = \frac{-1}{(1 + e^{-x})^2} (1 + e^{-x})'$$

$$= \frac{-1}{(1 + e^{-x})^2} \cdot 1 \cdot (e^{-x})'$$


$$= \frac{-1}{(1 + e^{-x})^2} \cdot 1 \cdot e^{-x} \cdot (-x)'$$

$$= \frac{-1}{(1 + e^{-x})^2} \cdot 1 \cdot e^{-x} \cdot -1$$


$$w_1 = \frac{e^{-x}}{(1 + e^{-x})^2} = \left(\frac{1 + e^{-x}}{1 + e^{-x}} + \frac{-1}{1 + e^{-x}} \right) \left(\frac{1}{1 + e^{-x}} \right) = (1 - \sigma(x))\sigma(x)$$




Perceptron Example




Perceptron Example



Perceptron Example



Perceptron Example




Derivatives are Cheap

- Each derivative equals the local partial derivative times the upstream derivative, thanks to the chain rule
- $+$ → pass through gradients
- \times → multiply other inputs by upstream gradient
- σ → squish gradients

Training a Neural Network



Multi-layer Perceptron (MLP)



Backpropagation

- Forward Pass: present training input pattern to network and activate network to produce output (can also do in batch: present all patterns in succession)
- Backward Pass: calculate error gradient and update weights starting at output layer and then going back

Gradient Descent



Building a Neural Network

Weight Update Equation

$$\theta_w' = \theta_w - \eta * \frac{\partial T_C}{\partial \theta_w}$$



New weight

Learning Rate

Gradient Descent

Input : list of n training examples $(x_0, d_0), \dots, (x_n, d_n)$
where $\forall i : d_i \in \{+1, -1\}$

Output : classifying hyperplane w

Algorithm :

Randomly initialize w ;

While makes errors on training set **do**

for (x_i, d_i) **do**

 let $y_i = \text{MLP}(w, x_i)$;

$loss \leftarrow \text{Mean}((d_i - y_i)^2)$

$w' \leftarrow \text{Backprop}(w, loss)$


$w \leftarrow w - \eta w'$

end

end

*x and w are vectors;
i is the instance index*


MLP Example



MLP Example

$$w_2 = \begin{bmatrix} .1 \\ .3 \\ -.1 \end{bmatrix}$$

$$w_1 = \begin{bmatrix} .2 & .1 \\ -.2 & -.1 \\ .3 & 0 \end{bmatrix}$$



MLP Example


$$f(x) = \sigma(h \times w_2) = [\sigma(.14)] = [.54]$$

$$w_2 = \begin{bmatrix} .1 \\ .3 \\ -.1 \end{bmatrix}$$

$$h = \sigma(x \times w_1) = [\sigma(.9) \quad \sigma(.3)]$$

$$w_1 = \begin{bmatrix} .2 & .1 \\ -.2 & -.1 \\ .3 & 0 \end{bmatrix}$$

$$x = [2 \quad -1 \quad 1]$$



Calculate Gradients

$$f(x) = \sigma(h \times w_2) = [\sigma(.14)] = [.54]$$

$$w_2 = \begin{bmatrix} .1 \\ .3 \\ -.1 \end{bmatrix}$$

$$h = \sigma(x \times w_1) = [\sigma(.9) \quad \sigma(.3)]$$

$$w_1 = \begin{bmatrix} .2 & .1 \\ -.2 & -.1 \\ .3 & 0 \end{bmatrix}$$

$$x = [2 \quad -1 \quad 1]$$

Target = 1 → Loss = $(1 - .54)^2 = .21$

Loss Gradient → $-2 \cdot .46 = -.92$

Sigmoid Gradient → $-.92 \cdot (1 - \sigma(.14))\sigma(.14) = -.23$

$$w'_2 = -.23 \begin{bmatrix} \sigma(.9) \\ \sigma(.3) \\ 1 \end{bmatrix} = \begin{bmatrix} -.16 \\ -.13 \\ -.23 \end{bmatrix}$$

$$h' = -.23 [1 \quad .3] = [-.02 \quad -.07]$$

Sigmoid Gradient →

$$[-.02 \cdot (1 - \sigma(.9))\sigma(.9), -.07 \cdot (1 - \sigma(.3))\sigma(.3)]$$

$$[-.02, -.004]$$

$$w'_1 = \begin{bmatrix} -.02 \cdot 2 & -.004 \cdot 2 \\ -.02 \cdot -1 & -.004 \cdot -1 \\ -.02 \cdot 1 & -.004 \cdot 1 \end{bmatrix} = \begin{bmatrix} -.04 & -.008 \\ .02 & .004 \\ -.02 & -.004 \end{bmatrix}$$

Update Weights

let $\eta = 0.001$


$$w_1 = w_1 - \eta w'_1$$

$$w_1 = \begin{bmatrix} .2 & .1 \\ -.2 & -.1 \\ .3 & 0 \end{bmatrix} - .001 \begin{bmatrix} -.04 & -.008 \\ .02 & .004 \\ -.02 & -.004 \end{bmatrix}$$

$$w_1 = \begin{bmatrix} .20004 & .10001 \\ -.20002 & -.1000 \\ .30002 & .00000 \end{bmatrix}$$

$$w_2 = w_2 - \eta w'_2$$

$$w_2 = \begin{bmatrix} .1 \\ .3 \\ -.1 \end{bmatrix} - .001 \begin{bmatrix} -.16 \\ -.13 \\ -.23 \end{bmatrix}$$



values are rounded

Fully Connected Layer

Class FCLayer:

var $w \leftarrow$ InitializeRandom(input size + 1, output size)

var $b \leftarrow$ InitializeRandom(output size)

var $\eta \leftarrow$ Learning Rate

def forward(x):

return $x \times w + b$

def backward($grads$):

$w' \leftarrow x^T \times grads$

$x' \leftarrow grads \times w^T$


$w \leftarrow w - \eta w'$

$b \leftarrow b - \eta \cdot grads$

return x'

*Just need to define
an Activation Layer
then you can stack them*

Activation Functions




$$f(x) = \sigma\left(\sum_{i=1}^n w_i x_i + b\right)$$

$$\sigma(net) = \frac{1}{1 + e^{-net}}$$


Activation Functions

Unit step (threshold)


$$f(x) = \begin{cases} 0 & \text{if } x > 0 \\ 1 & \text{if } x \geq 0 \end{cases}$$



Activation Functions




Activation Functions




Hyperbolic Tangent

Activation Functions



**Rectified Linear Units
(ReLU)**

Elements of a Neural Network



<http://playground.tensorflow.org/>

Deep Learning

Why Deep Learning is so successful

- Deep Part (deep hierarchy improves optimization terrain)
 - Architecture, Activation, Initialization
- GPU + Internet (Data)
 - Allow more layers and weights
- Only Successful Architecture
 - CNN/RNN (only grids and sequences)

Perceptron problems

- Noise: if the data isn't separable, weights might thrash
- Multiple layers allow it to work on non-linearly separable data
- Mediocre generalization: finds a “barely” separating solution
- More data and a continuous loss function greatly improve generalization
- Overtraining: test / held-out accuracy usually rises, then falls
- Large data, many layers, and batch training help

